

Computer Vision

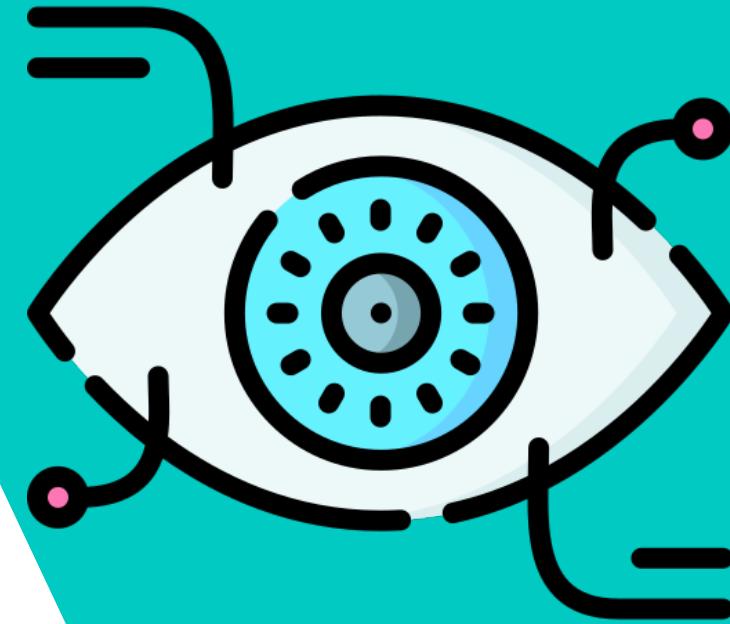
A.A. 2024-20245

Lecture 2: Image processing



SAPIENZA
UNIVERSITÀ DI ROMA

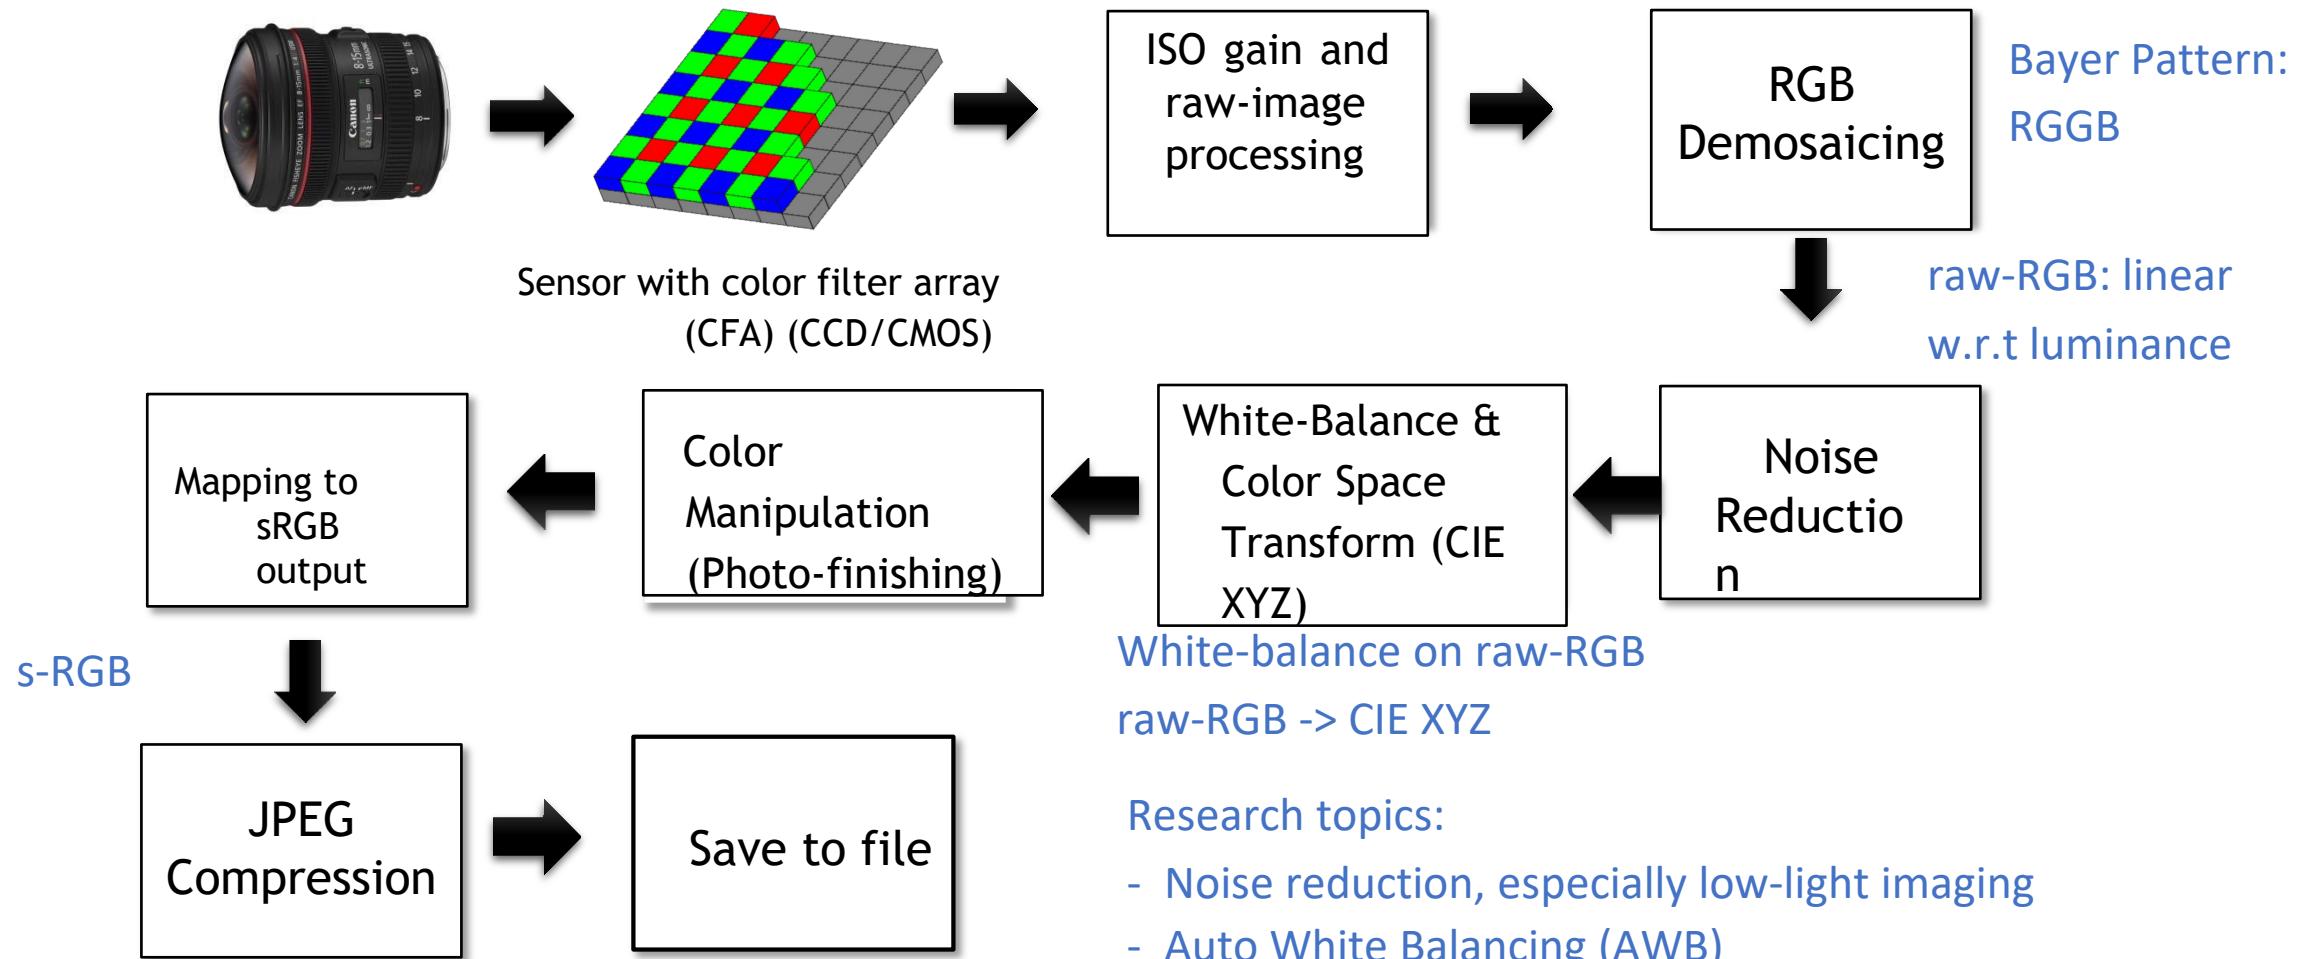
ALCOR Lab



Reading

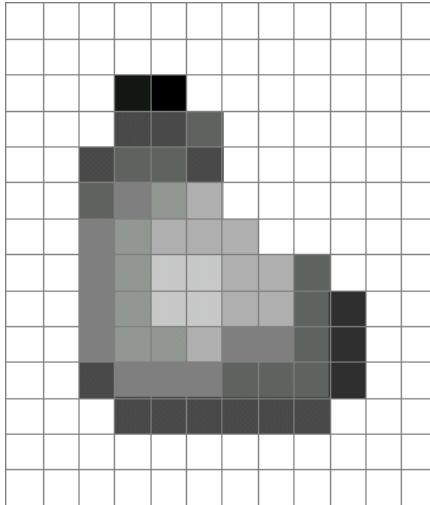
- Szeliski textbook, Chapter 3.1, 3.2, 3.3, 3.5

A typical color imaging pipeline



What is an image?

- A grid (matrix) of intensity values



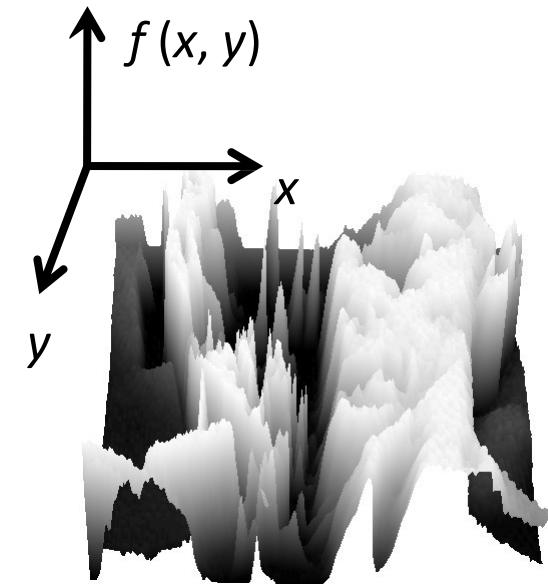
=

255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	20	0	255	255	255	255	255	255	255	255	255	255	255
255	255	255	75	75	75	255	255	255	255	255	255	255	255	255	255
255	255	75	95	95	75	255	255	255	255	255	255	255	255	255	255
255	255	96	127	145	175	255	255	255	255	255	255	255	255	255	255
255	255	127	145	175	175	175	255	255	255	255	255	255	255	255	255
255	255	127	145	200	200	175	175	175	95	255	255	255	255	255	255
255	255	127	145	200	200	175	175	175	95	47	255	255	255	255	255
255	255	127	145	145	175	127	127	95	47	255	255	255	255	255	255
255	255	74	127	127	127	95	95	95	95	47	255	255	255	255	255
255	255	255	74	74	74	74	74	74	74	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255

(common to use one byte per value: 0 = black, 255 = white)

What is an image?

- Can think of a (grayscale) image as a **function** f from \mathbb{R}^2 to \mathbb{R} :
 - $f(x, y)$ gives the **intensity** at position (x, y)



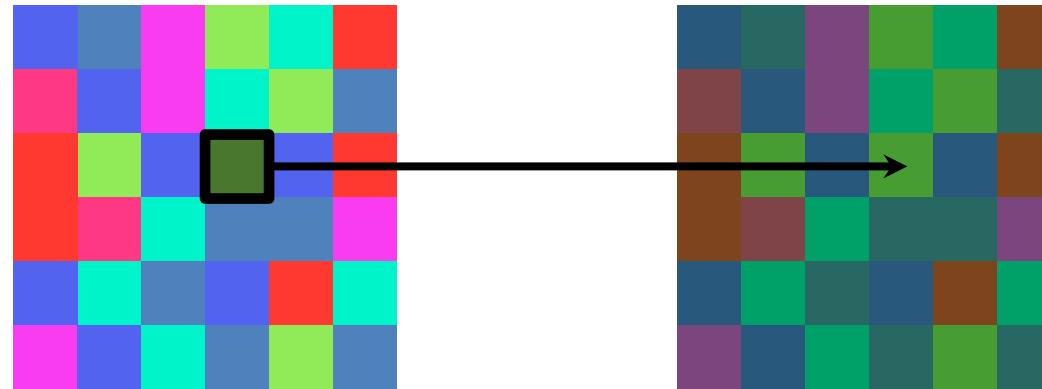
- A **digital** image is a discrete (**sampled, quantized**) version of this function

Image Transformations & Filtering

- Point Processing
- Linear Filtering
- Sampling & Aliasing
- Image Derivatives
- Edge Detection

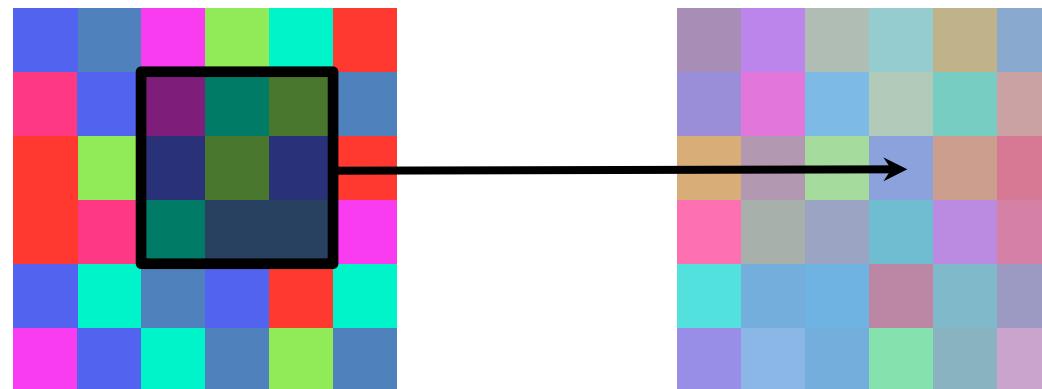
What types of image filtering can we do?

Point Operation



point processing

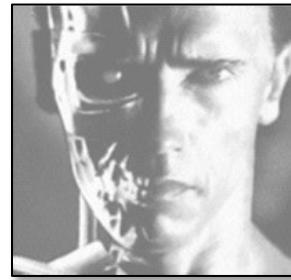
Neighborhood Operation



“filtering”

Examples of Point Processing

- As with any function, we can apply operators to an image



$$g(x,y) = f(x,y) + 20$$

brightness



$$g(x,y) = f(-x,y)$$

Horizontal flip

Examples of point processing

original



darken



lower contrast



non-linear lower contrast



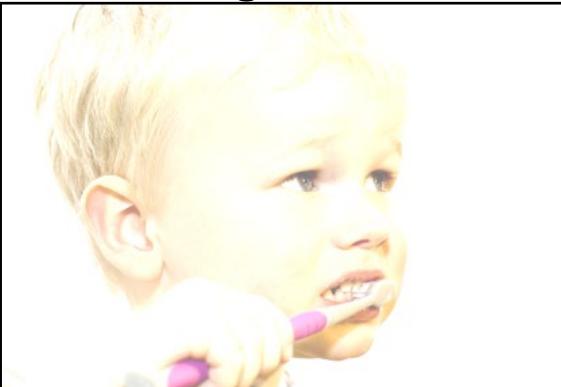
x

invert



$x - 128$

lighten



$\frac{x}{2}$

raise contrast



$\left(\frac{x}{255}\right)^{1/3} \times 255$

non-linear raise contrast



$255 - x$

$x + 128$

$x \times 2$

$\left(\frac{x}{255}\right)^2 \times 255$

Gamma Correction

- Adjust brightness and contrast in images for accurate display on different devices (monitors, TVs, etc.).
- Compensates for the non-linear relationship between input pixel values and output display brightness.

Why is it needed?

- Devices like monitors have a non-linear response to brightness.
- The human eye perceives light in a non-linear way, and gamma correction ensures images appear natural. Human vision is more sensitive to changes in darker tones than in lighter tones.

$$V_{\text{out}} = A V_{\text{in}}^{\gamma},$$

$\gamma > 1$: Reduces the brightness (darkens the image).

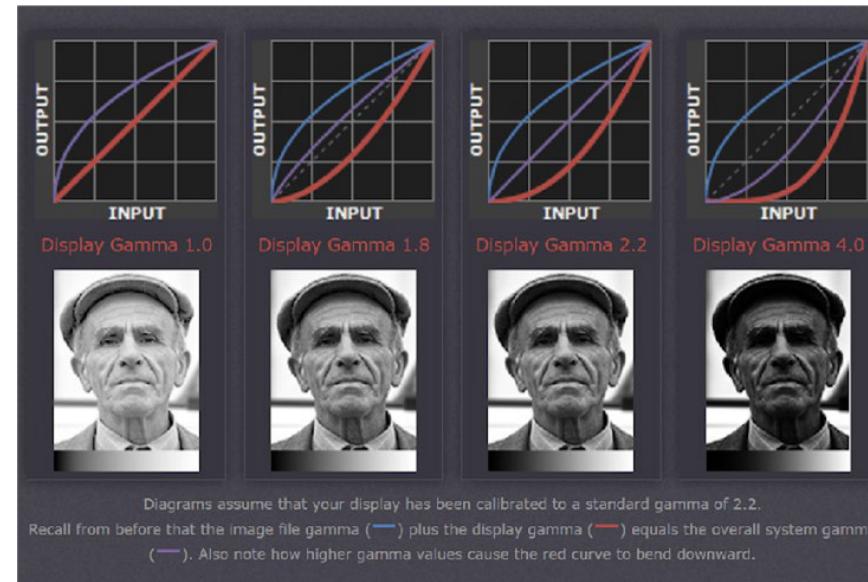
$\gamma < 1$: Increases brightness (lightens the image).

$\gamma = 1$: No change (linear).

A: A normalization constant, which might ensure that values are within a standard range (like 0-255 for an 8-bit image or 0-1 for normalized images).

Gamma Correction

- Most images we work with are in sRGB (Standard Red Green Blue) space, i.e., already tone-mapped.
- Many Vision algorithms expects the image to be in linear space.
- sRGB \rightarrow Linear space conversion requires explicit knowledge of camera processing pipeline. Requires knowledge of tone-production curve.
- In absence of it it is very common to use $\gamma = 2.2$.
- Note: This is not accurate, just a cheap approximation that works for most Vision tasks.

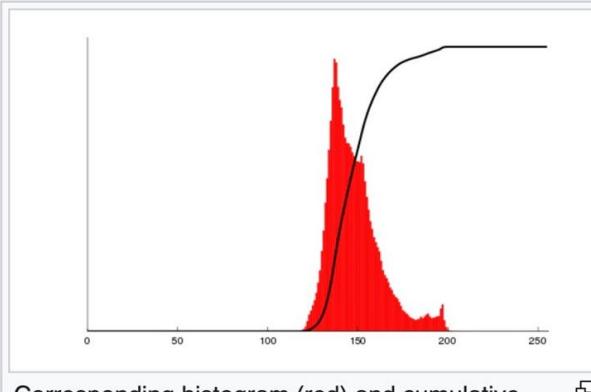


Histogram Equalization

Image processing technique that adjusts the contrast of an image by using its histogram.



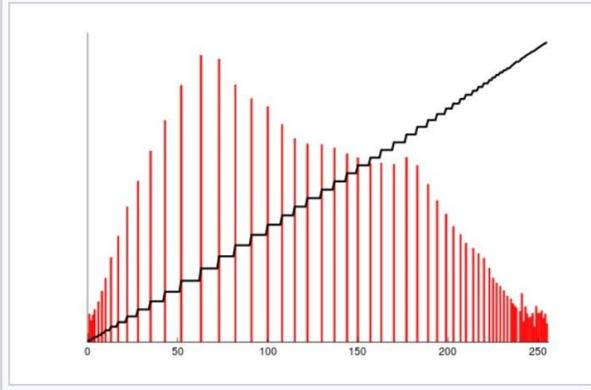
Before Histogram Equalization



Corresponding histogram (red) and cumulative histogram (black)



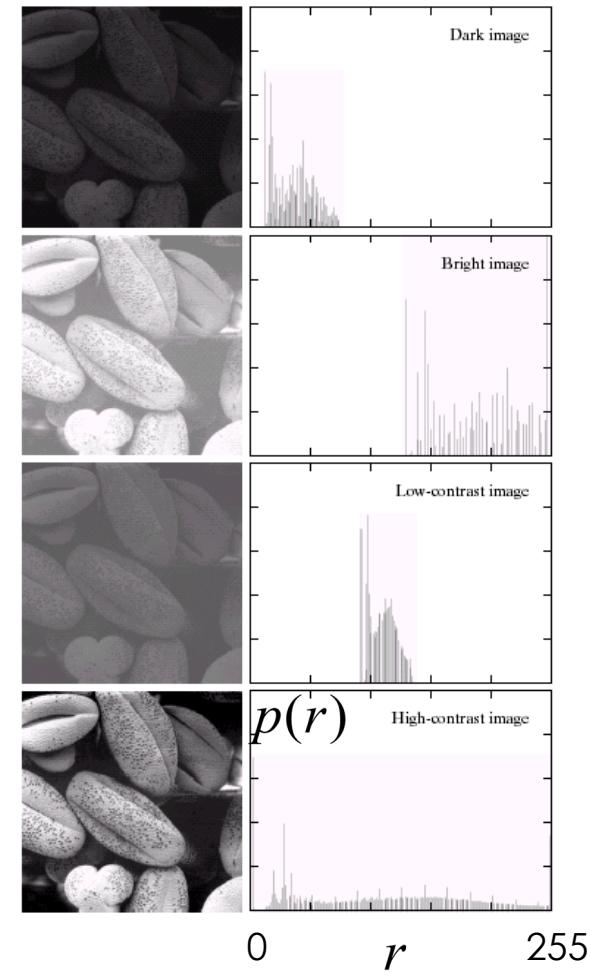
After Histogram Equalization



Corresponding histogram (red) and cumulative histogram (black)

Intensity Histogram

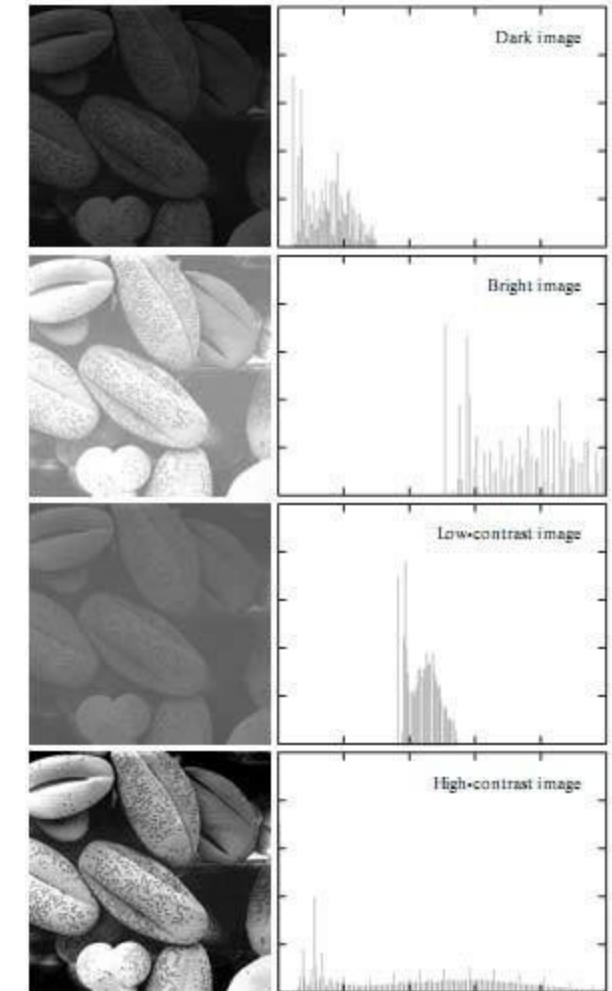
$$p(r) = \left(\frac{\text{Number of pixels with intensity } r}{\text{Total number of pixels}} \right)$$



Bahadir K. Gunturk

High contrast image

- The histogram components are distributed over all the intensity range.
- The distribution is almost uniform, with few peaks.
- If the distribution is uniform, the image tends to have a high dynamic range and the details are more easily perceived.
- This is the effect pursued by the histogram-based transformations.



How it works

- **Calculate the Histogram:** Count the number of pixels at each intensity level.

$$p(r_k) = \frac{\text{Number of pixels with intensity } r_k}{\text{Total number of pixels}}$$

- **Normalize the Histogram:** Find the cumulative distribution function (CDF) of the histogram.

$$CDF(r_k) = \sum_{i=0}^k p(r_i)$$

- **Map Intensities:** Adjust pixel values based on the CDF to achieve a more uniform distribution. The CDF is normalized in such a way the minimum and maximum values of the CDF correspond to the minimum and maximum possible intensity values (0 and 255 for an 8-bit image).

$$r'_k = \text{round} \left(\frac{(CDF(r_k) - CDF_{\min}) \cdot (L - 1)}{N - CDF_{\min}} \right)$$

Where:

- r'_k is the new intensity.
- L is the maximum pixel intensity (e.g., 255 for 8-bit images).
- N is the total number of pixels in the image.
- CDF_{\min} is the minimum non-zero value of the CDF.

$$I = \begin{bmatrix} 10 & 10 & 30 \\ 40 & 10 & 30 \\ 50 & 1 & 1 \end{bmatrix}$$

$$CDF: \left(\frac{(s-2) \cdot (255-1)}{9-2} \right) = 109$$

$$I = \begin{bmatrix} 109 & 109 & 20 \\ 40 & 109 & 20 \\ 50 & x & x \end{bmatrix}$$

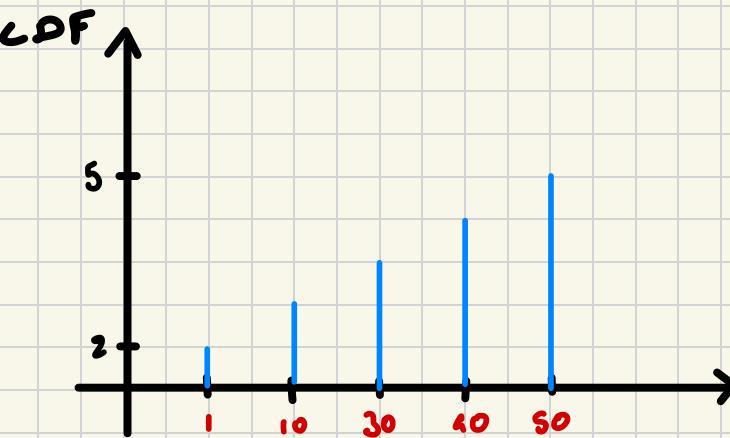
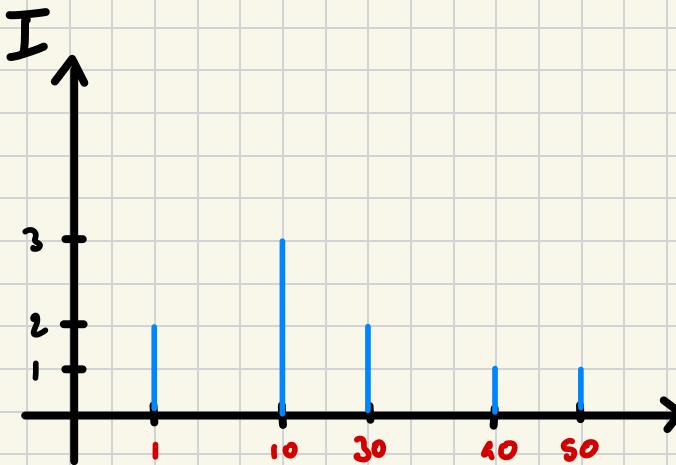
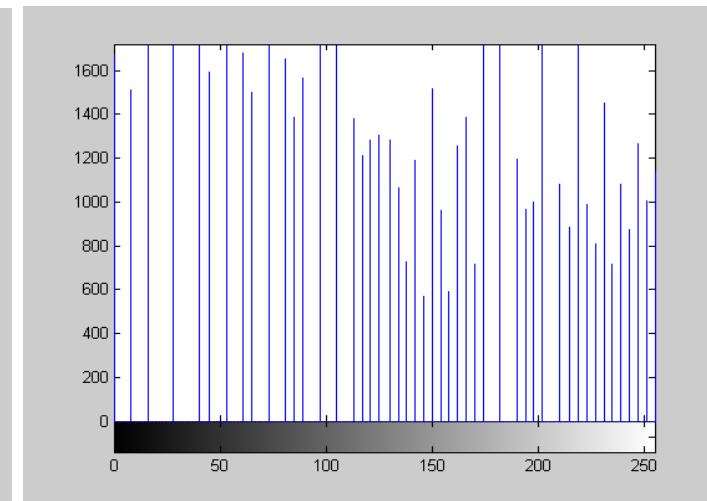
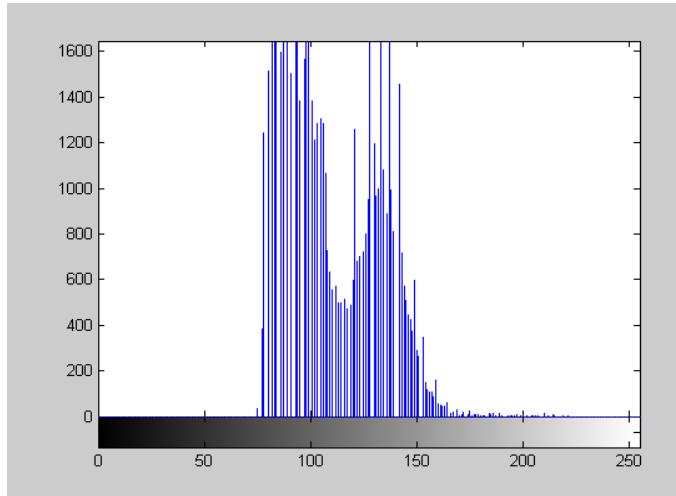


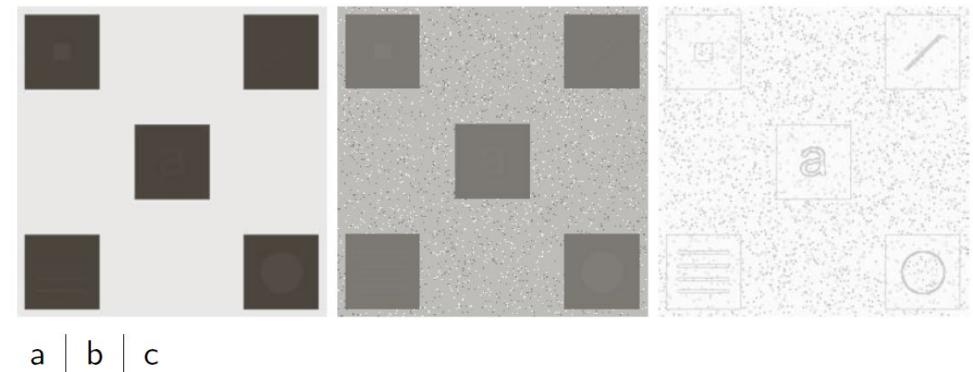
Image Enhancement by Point Processing

- Histogram Equalization



Local Histogram processing

- Histogram equalization is a global approach.
- Local histogram equalization is realized selecting, for each pixel, a suitable neighborhood on which the histogram equalization (or matching) is computed.
- More computational intensive
- Non overlapping regions may produce “blocky” effect.
- Adaptive histogram equalization (AHE), Contrast limited adaptive histogram equalization (CLAHE)



(a) original image
(b) equalized image
(c) locally equalized image (3×3 neighborhood)



Image Transformations & Filtering

- Point Processing
- Linear Filtering
- Sampling & Aliasing
- Image Derivatives
- Edge Detection

Filters

- Filtering
 - Form a new image whose pixel values are a combination of the original pixel values.
- Why?
 - To get useful information from images
 - E.g., extract edges or contours (to understand shape)
 - To enhance the image
 - E.g., to remove noise
 - E.g., to sharpen and “enhance image”
 - A key operator in Convolutional Neural Networks

Canonical Image Processing Problems

- Image Restoration
 - denoising
 - deblurring
- Image Compression
 - JPEG, HEIF, MPEG, ...
- Computing Field Properties
 - optical flow
 - disparity
- Locating Structural Features
 - corners
 - edges

Noise reduction

- Given a camera and a still scene, how can you reduce noise?



Take lots of images and average them!

What's the next best thing?

Image filtering

- Modify the pixels in an image based on some function of a local neighborhood of each pixel

10	5	3
4	5	1
1	1	7

Local image data

Some function

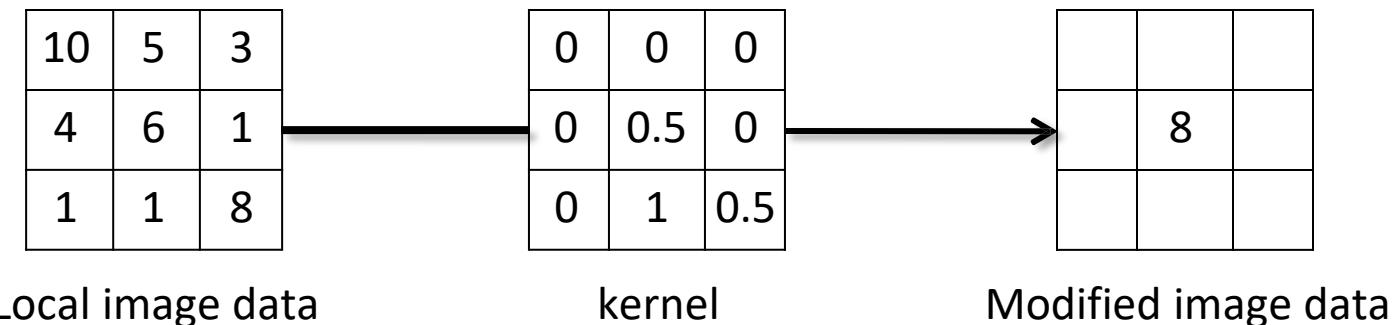


	7	

Modified image data

Linear filtering

- One simple version of filtering: linear filtering (cross-correlation, convolution)
 - Replace each pixel by a linear combination (a weighted sum) of its neighbors
- The prescription for the linear combination is called the “kernel” (or “mask”, “filter”)



Cross-correlation

Let F be the image, H be the kernel (of size $2k+1 \times 2k+1$), and the output image G

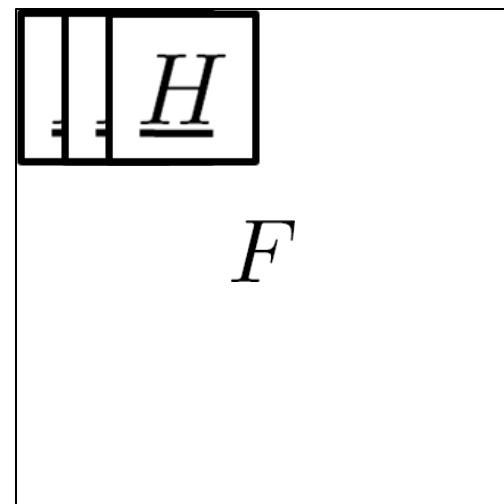
$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

This is called a **cross-correlation** operation:

$$G = H \otimes F$$

- Can think of as a “dot product” between local neighborhood and kernel for each pixel

Cross-correlation



- Sometimes called just **correlation**
- Neither associative nor commutative
- In the literature, people often just call “convolution”
- Filters often symmetric, so won’t matter

Convolution

- Same as cross-correlation, except that the kernel is “flipped” (horizontally and vertically)

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

Cross-correlation

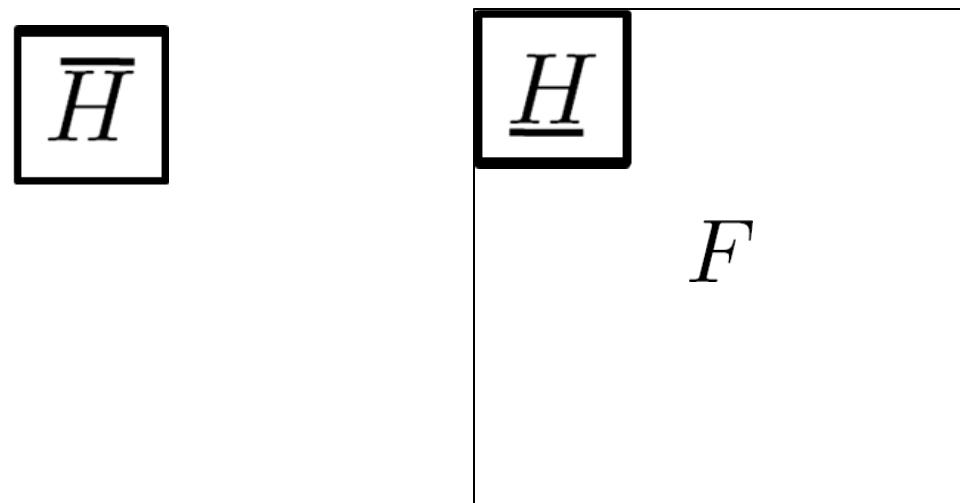
- Convolution is **commutative** and **associative**

Convolution is nice!

- Notation: $b = c \star a$
- Convolution is a multiplication-like operation
 - commutative
 - associative $a \star b = b \star a$
 - distributes over addition $a \star (b \star c) = (a \star b) \star c$
 - scalar factor out $\alpha a \star b = a \star \alpha b = \alpha(a \star b)$
 - identity: unit impulse $e = [..., 0, 0, 1, 0, 0, ...]$ $a \star e = a$
- Usefulness of associativity
 - often apply several filters one after another: $((a * b_1) * b_2) * b_3$
 - this is equivalent to applying one filter: $a * (b_1 * b_2 * b_3)$

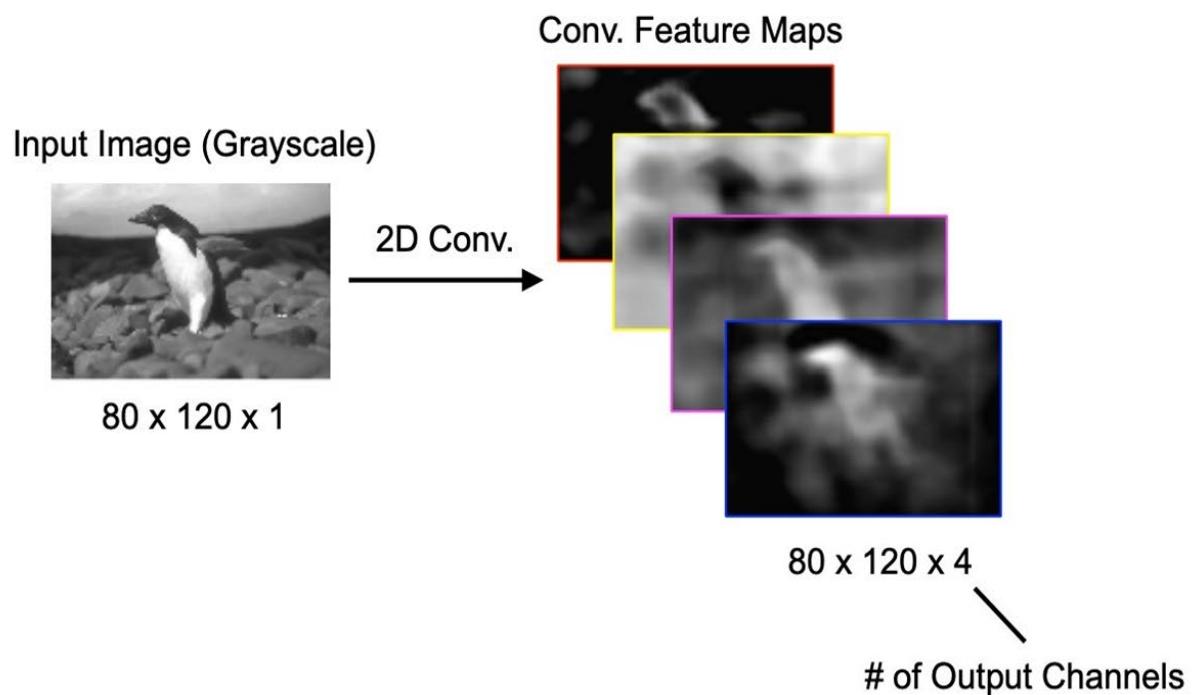
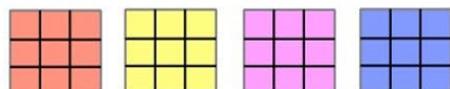
Convolution

The kernel is “flipped” (horizontally and vertically)



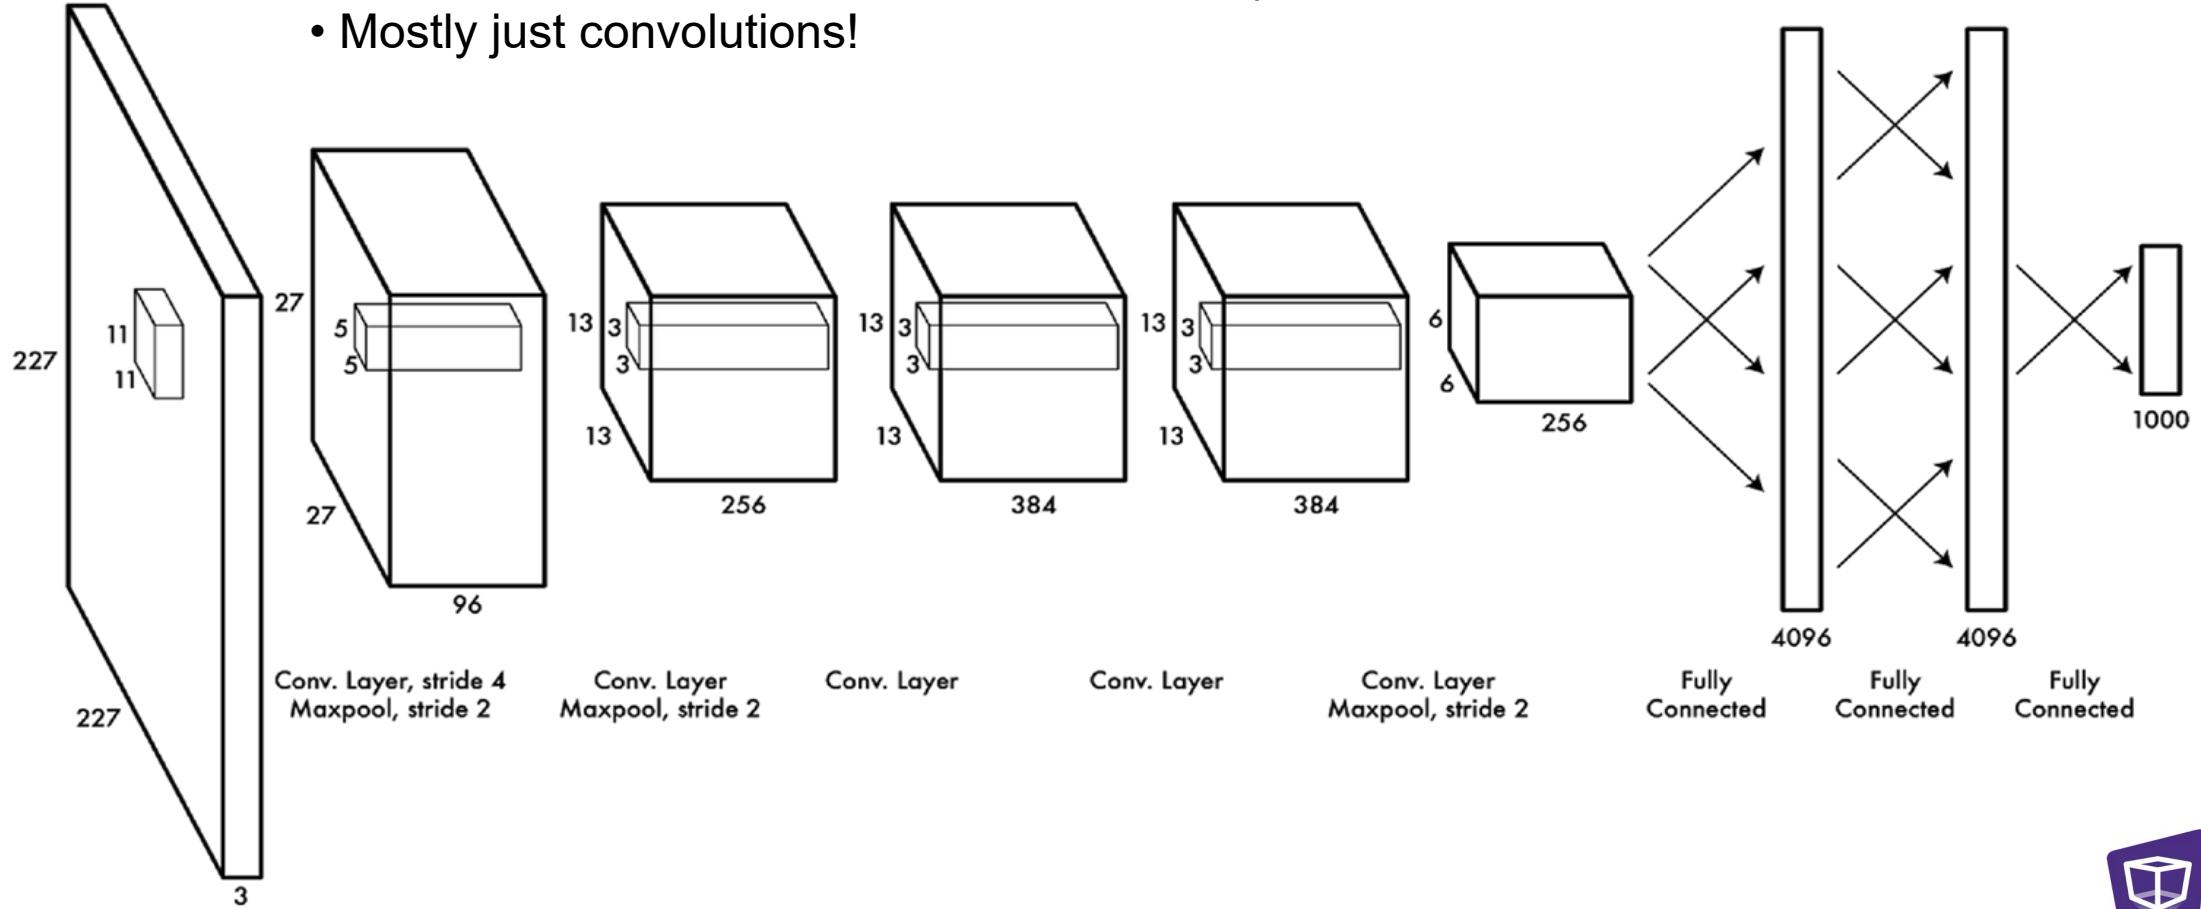
Convolutional Networks

Learnable 3x3 Convolutional Kernels



AlexNet: An Early Example

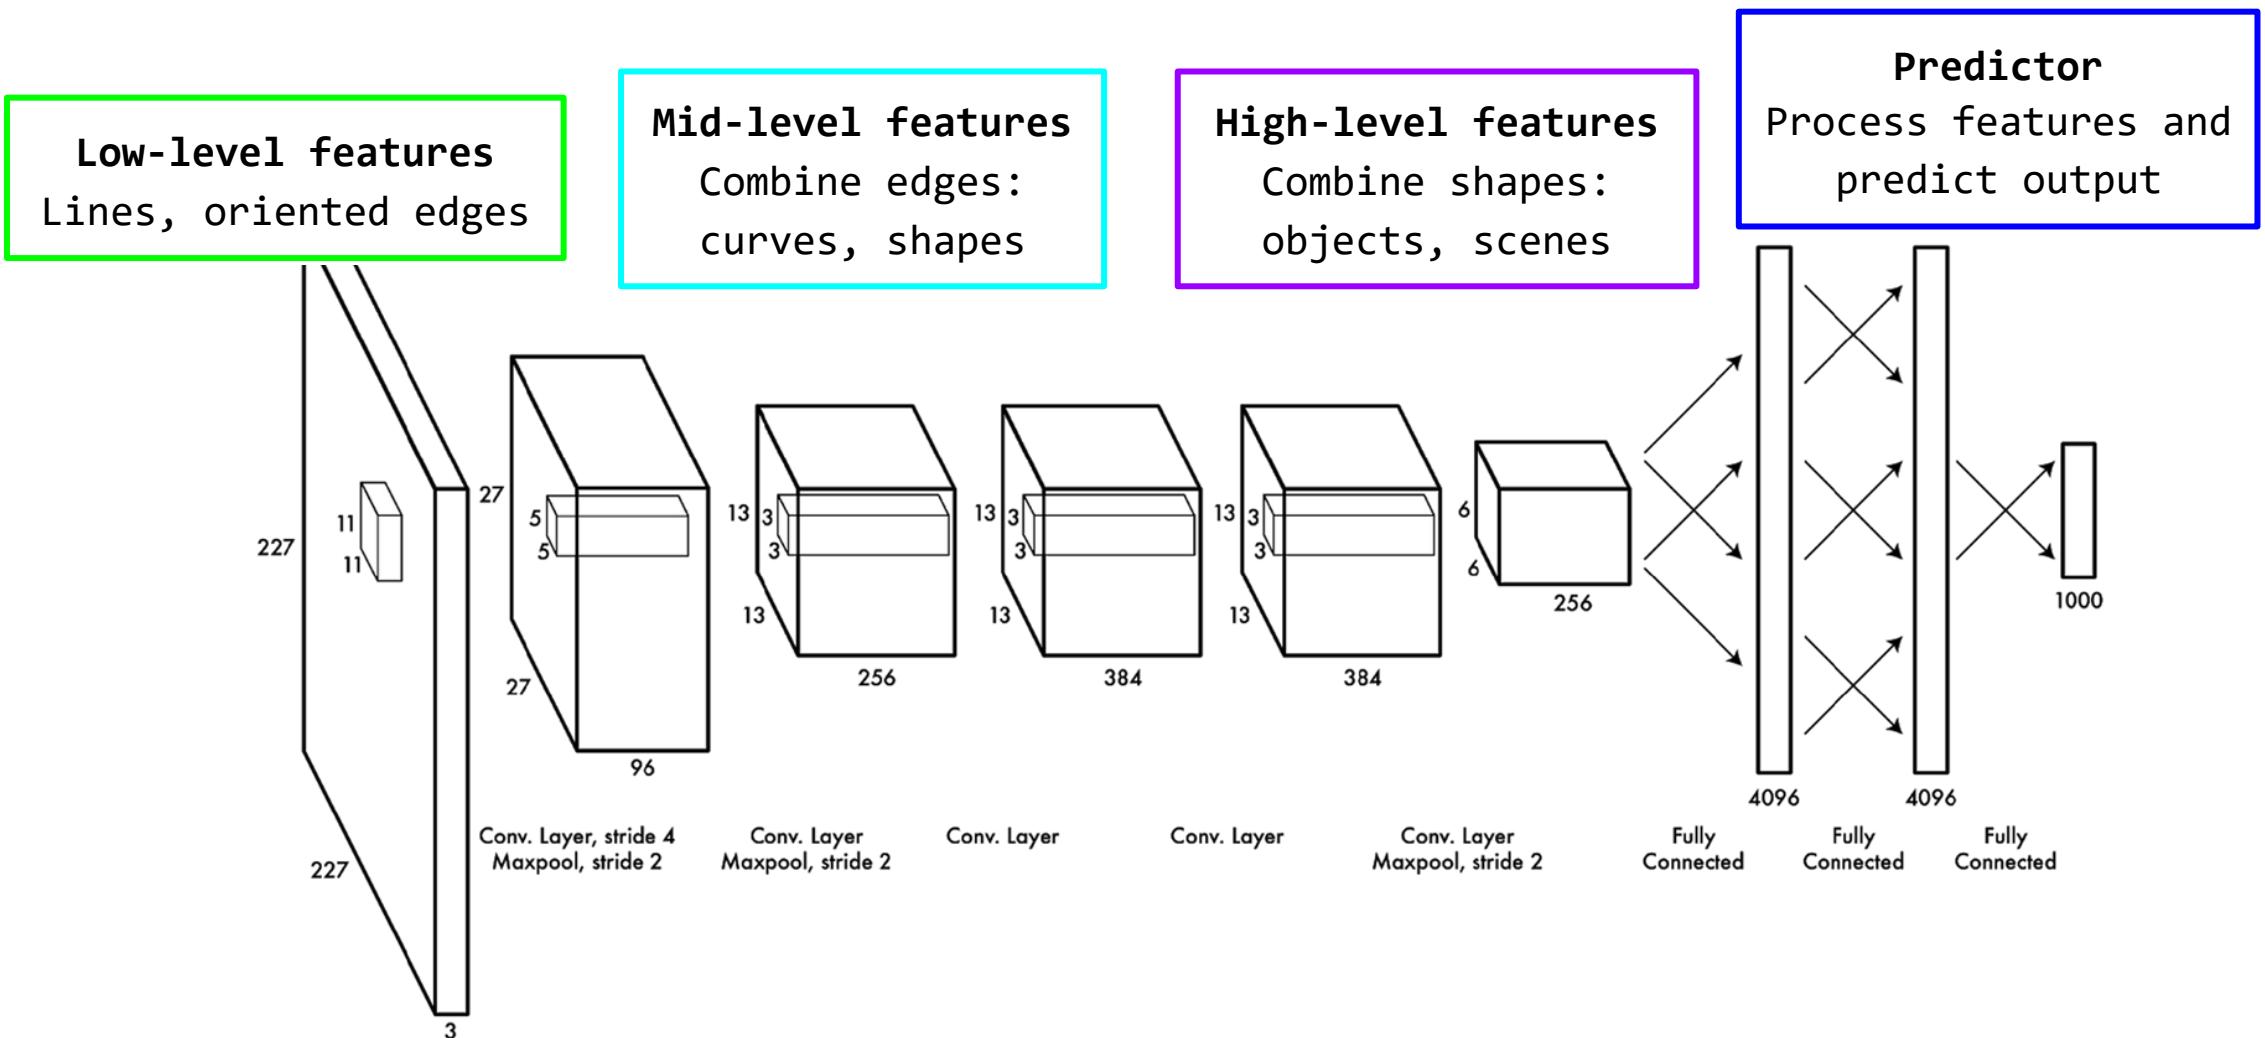
- Neural network with specialized connectivity structure
- Mostly just convolutions!



<https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks.pdf>



Where Models Learn Features of an Image

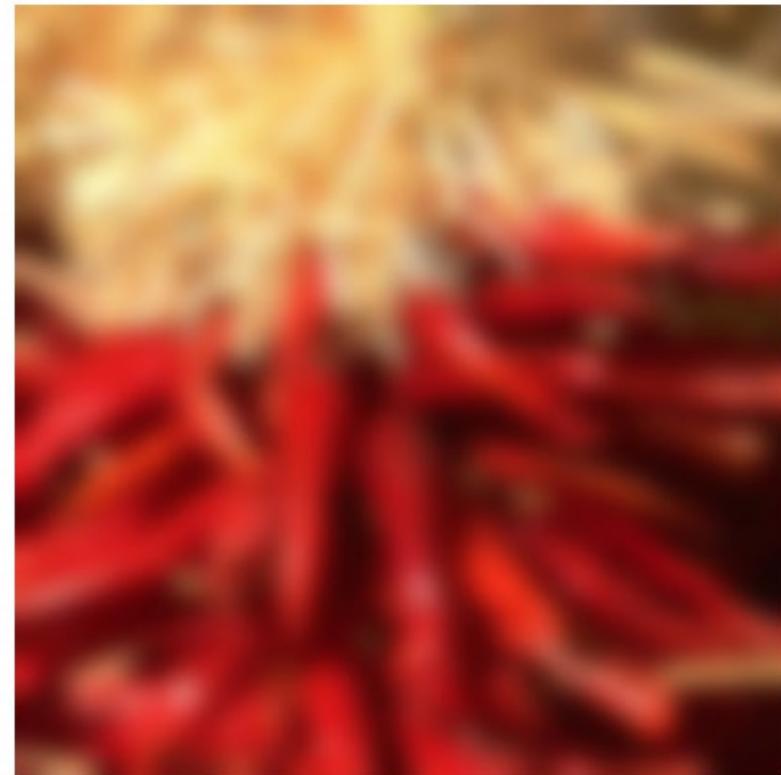


What are Models Actually Looking For?



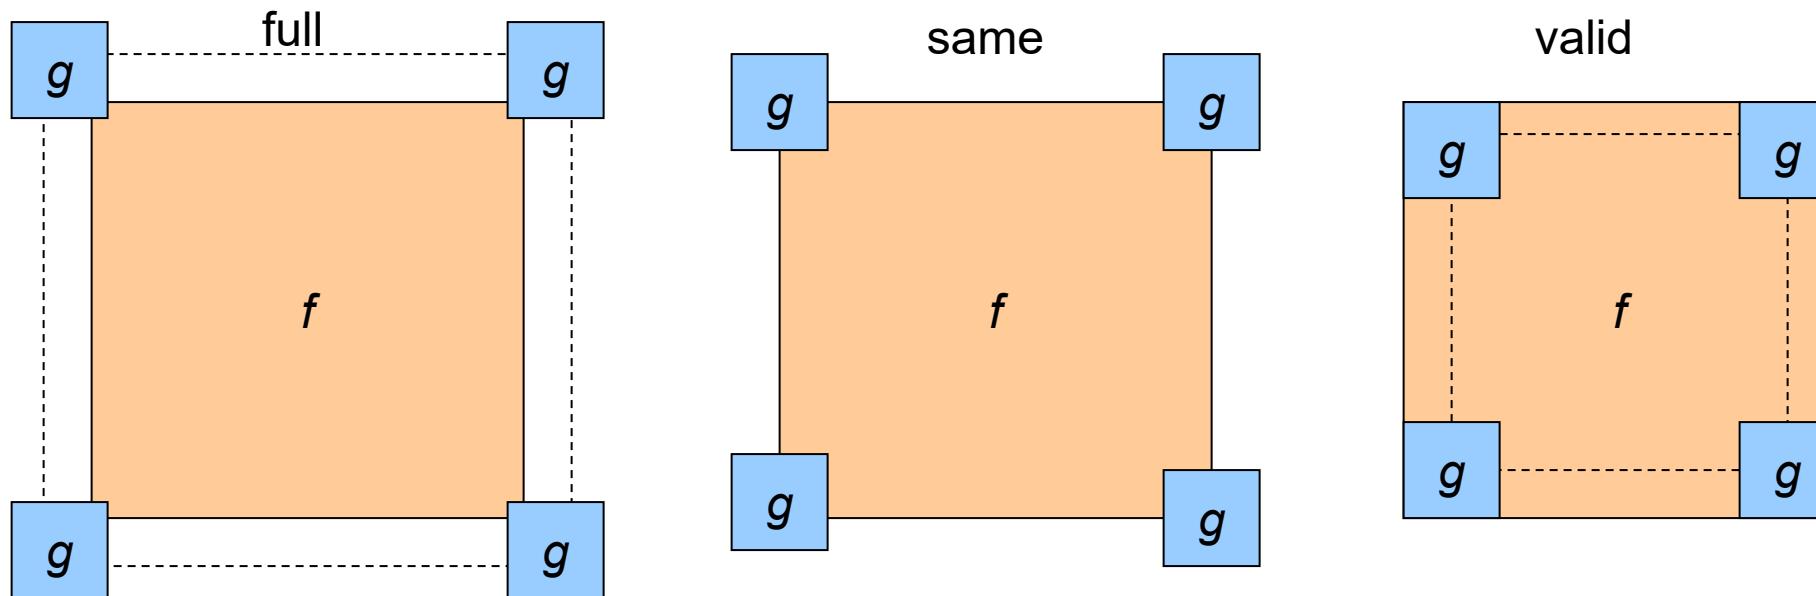
Practical matters

- What about near the edge?
 - the filter window falls off the edge of the image
 - need to extrapolate
 - Methods (padding):
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge

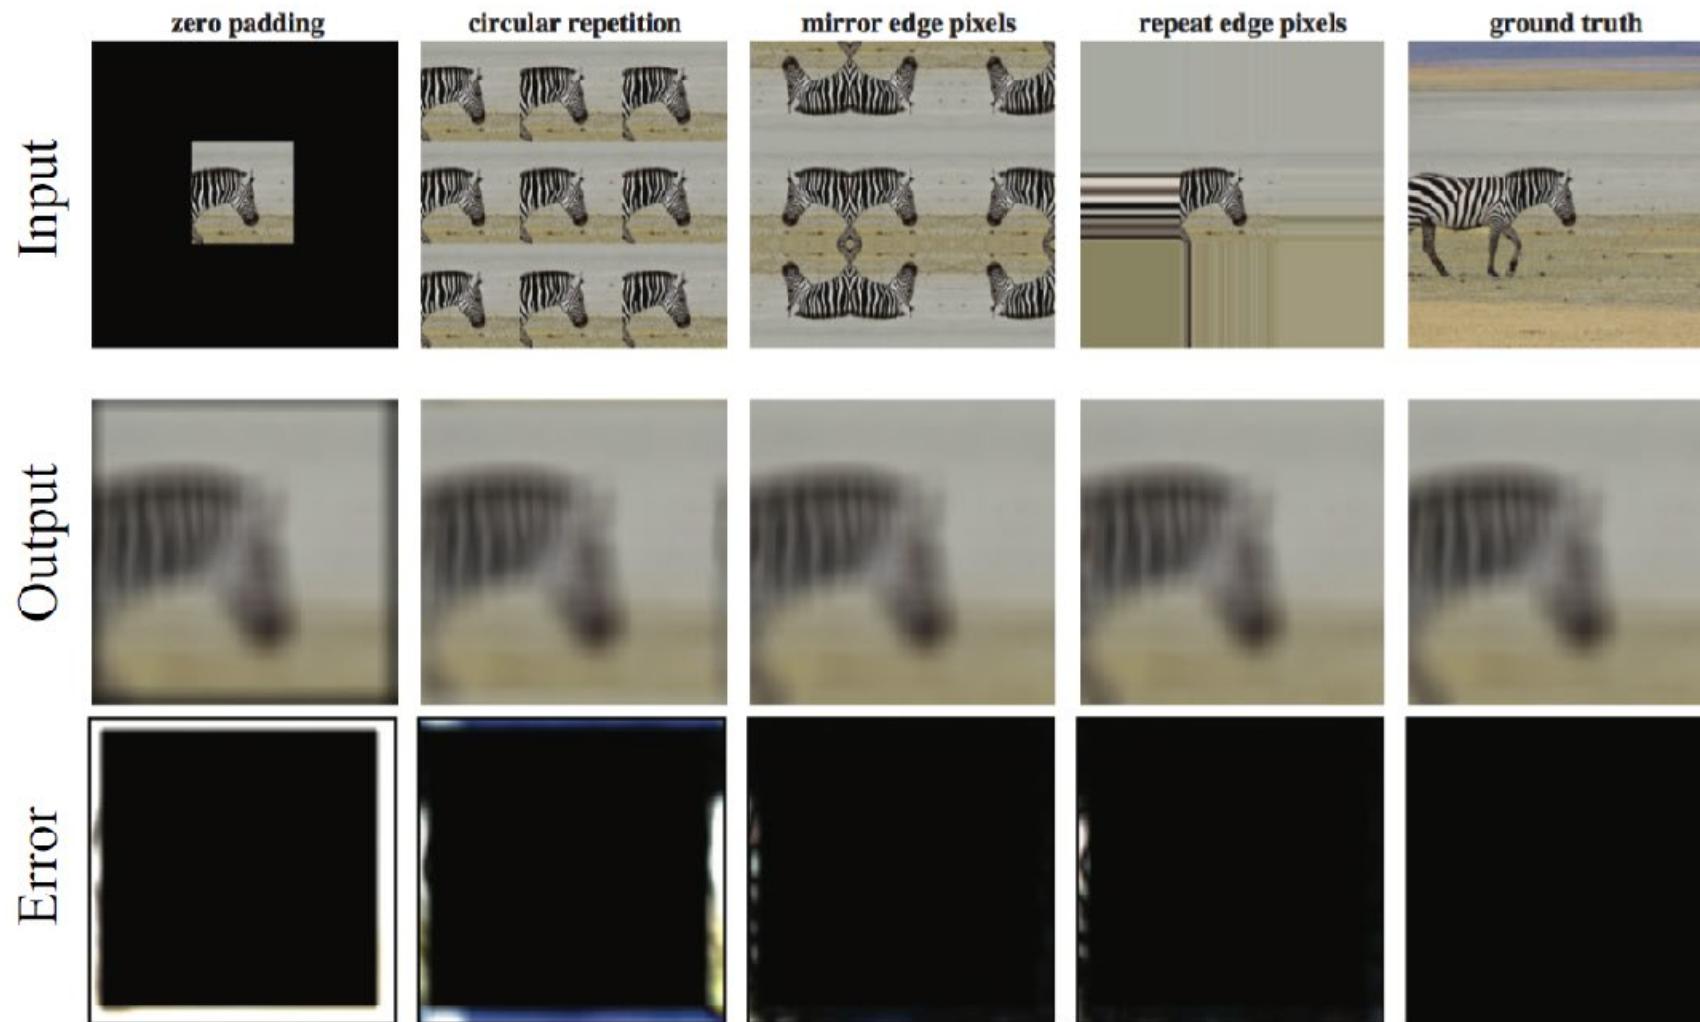


Practical matters

- What is the size of the output?
 - `shape = 'full'`: output size is sum of sizes of f and g
 - `shape = 'same'`: output size is same as f
 - `shape = 'valid'`: output size is difference of sizes of f and g



Handling boundaries



Source: Torralba, Freeman, Isola

Padding & Stride in CNN

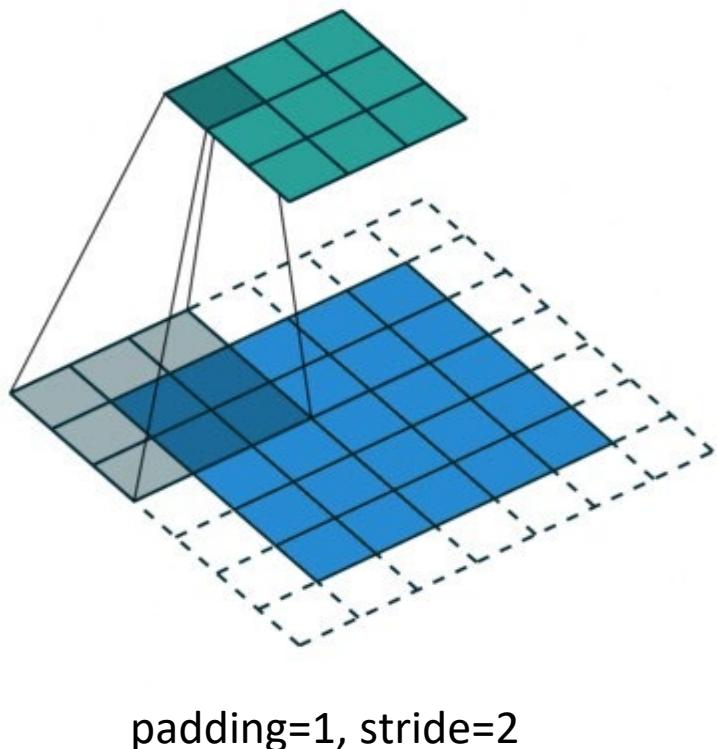
```
CLASS torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0,  
dilation=1, groups=1, bias=True, padding_mode='zeros', device=None, dtype=None) [SOURCE]
```

Applies a 2D convolution over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size (N, C_{in}, H, W) and output $(N, C_{\text{out}}, H_{\text{out}}, W_{\text{out}})$ can be precisely described as:

$$\text{out}(N_i, C_{\text{out}_j}) = \text{bias}(C_{\text{out}_j}) + \sum_{k=0}^{C_{\text{in}}-1} \text{weight}(C_{\text{out}_j}, k) \star \text{input}(N_i, k)$$

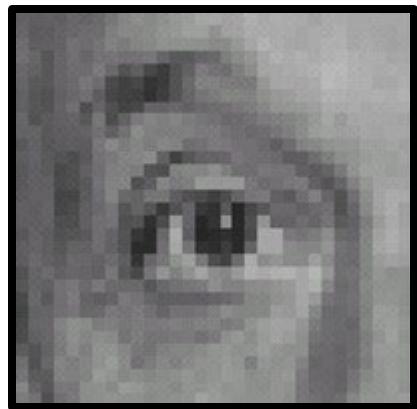
where \star is the valid 2D **cross-correlation** operator, N is a batch size, C denotes a number of channels, H is a height of input planes in pixels, and W is width in pixels.



$$H_{\text{out}} = \left\lfloor \frac{H_{\text{in}} + 2 \times \text{padding}[0] - \text{dilation}[0] \times (\text{kernel_size}[0] - 1) - 1}{\text{stride}[0]} + 1 \right\rfloor$$

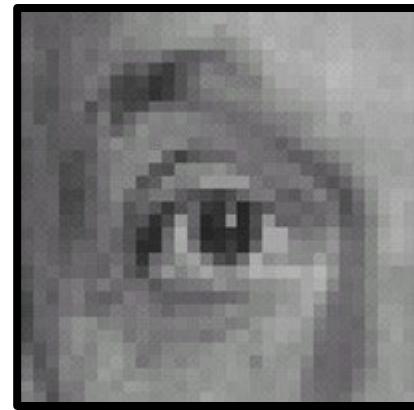
$$W_{\text{out}} = \left\lfloor \frac{W_{\text{in}} + 2 \times \text{padding}[1] - \text{dilation}[1] \times (\text{kernel_size}[1] - 1) - 1}{\text{stride}[1]} + 1 \right\rfloor$$

Linear filters: examples



*

0	0	0
0	1	0
0	0	0



Original

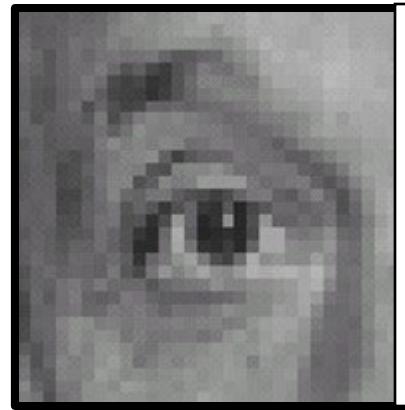
Linear filters: examples



Original

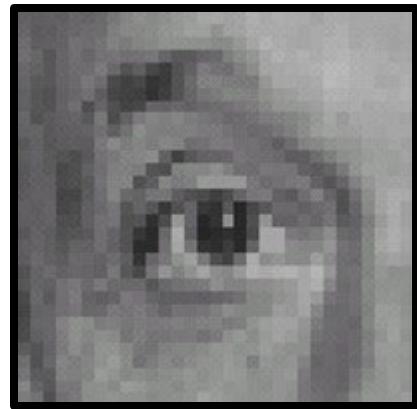
*

0	0	0
1	0	0
0	0	0

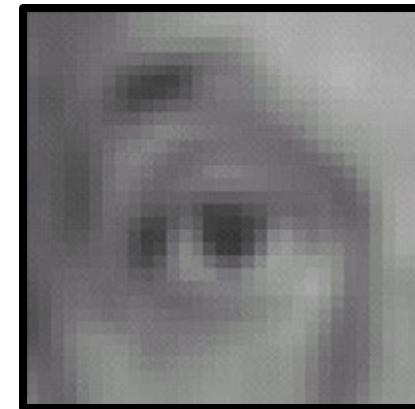


Shifted left by 1 pixel

Linear filters: examples

 $*$

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

 $=$ 

Original

Blur (with a mean filter)

Mean filtering



H

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

=

0	10	20	30	30	30	20	10		
0	20	40	60	60	60	40	20		
0	30	60	90	90	90	60	30		
0	30	50	80	80	90	60	30		
0	30	50	80	80	90	60	30		
0	20	30	50	50	60	40	20		
10	20	30	30	30	30	20	10		
10	10	10	0	0	0	0	0		

F

G

Mean filtering/Moving average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

Mean filtering/Moving average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

0	10									

Mean filtering/Moving average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

0	10	20								

Mean filtering/Moving average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

0	10	20	30						

Mean filtering/Moving average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

Mean filtering/Moving average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0	10	20	30	30	30	20	10		
	0	20	40	60	60	60	40	20		
	0	30	60	90	90	90	60	30		
	0	30	50	80	80	90	60	30		
	0	30	50	80	80	90	60	30		
	0	20	30	50	50	60	40	20		
	10	20	30	30	30	30	20	10		
	10	10	10	0	0	0	0	0		

Separable filters

A 2D filter is separable if it can be written as the product of a “column” and a “row”.

example:
box filter

1	1	1
1	1	1
1	1	1

=

1
1
1

*

1	1	1
---	---	---

row

column

2D convolution with a separable filter is equivalent to two 1D convolutions (with the “column” and “row” filters).

Separable filters

A 2D filter is separable if it can be written as the product of a “column” and a “row”.

example:
box filter

1	1	1
1	1	1
1	1	1

=

1
1
1

*

1	1	1
---	---	---

row

column

2D convolution with a separable filter is equivalent to two 1D convolutions (with the “column” and “row” filters).

If the image has $M \times N$ pixels and the filter kernel has size $L \times L$:

- What is the cost of convolution with a non-separable filter? $\longrightarrow L^2 \times M \times N$
- What is the cost of convolution with a separable filter? $\longrightarrow 2 \times (L \times M \times N)$

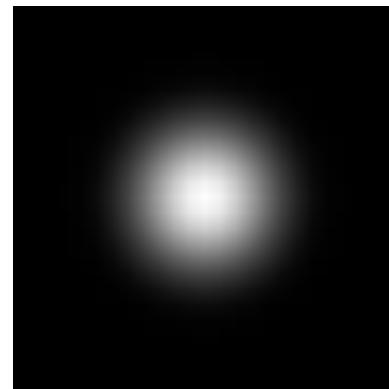
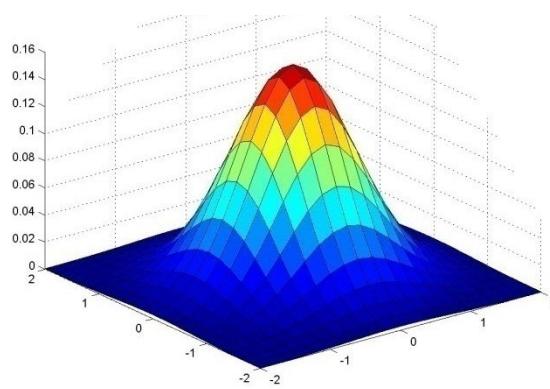
How to verify it?

- If a matrix has a rank of 1 (all of its row or columns are scalar multiples of each others). In this case the matrix can be factored as the outer product of two vectors.

Outer product: takes two vector and produces a matrix whose elements are the pairwise products of the elements of the two vectors.

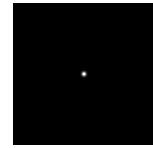
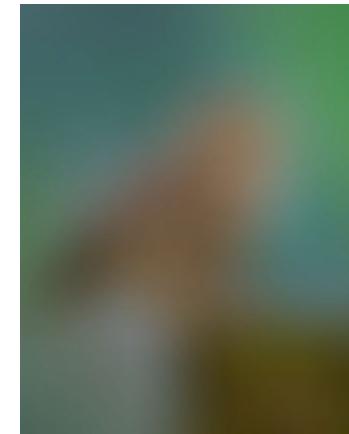
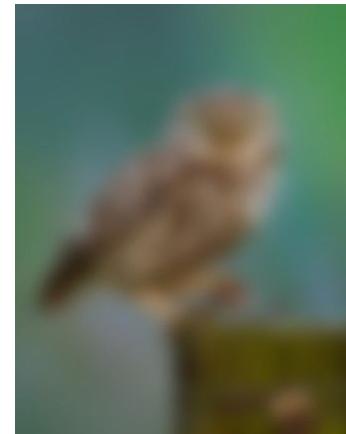
- The outer product of two vectors always results in a matrix of rank 1 (unless one of the vector is the zero vector).

Gaussian kernel

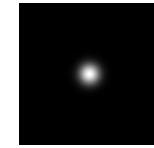


$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

Gaussian kernel



$\sigma = 1$ pixel



$\sigma = 5$ pixels

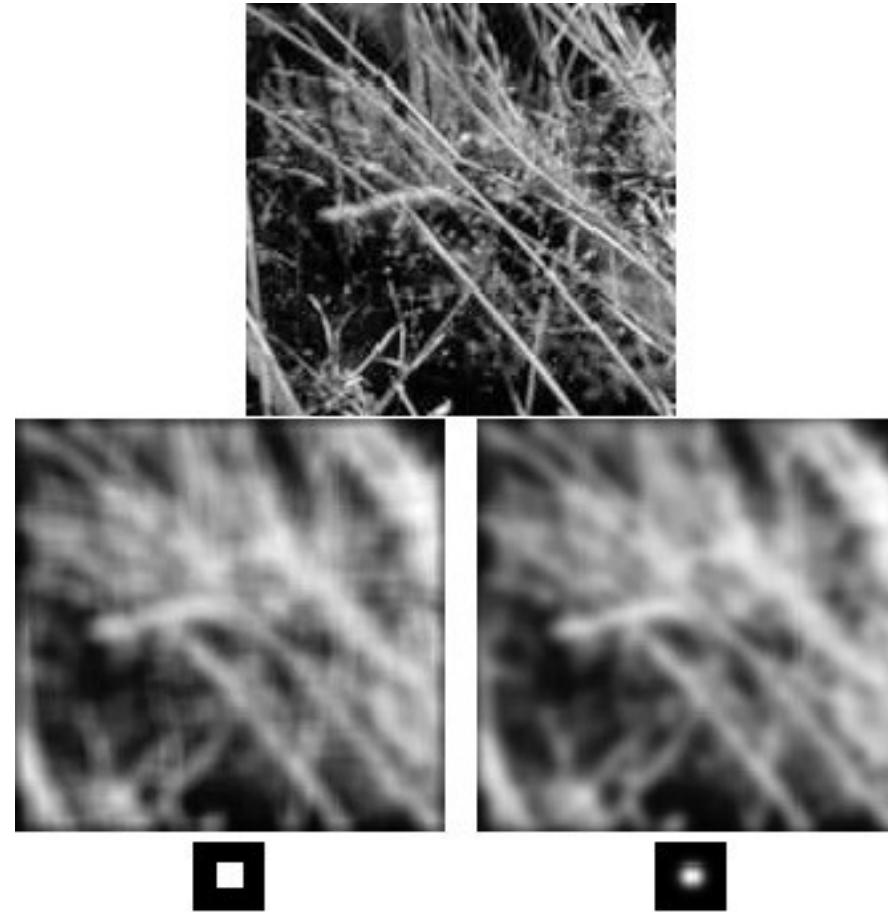


$\sigma = 10$ pixels

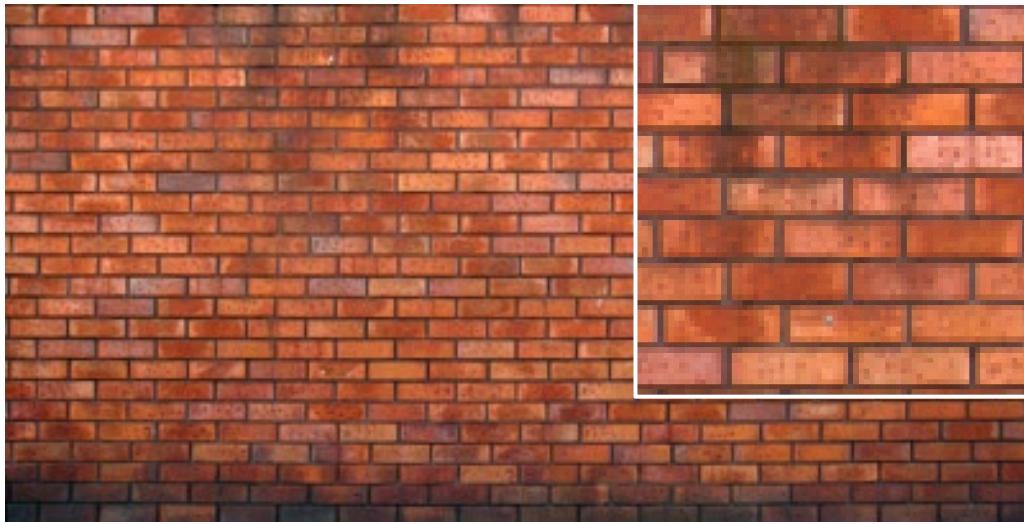


$\sigma = 30$ pixels

Mean (or box filtering) vs Gaussian filtering

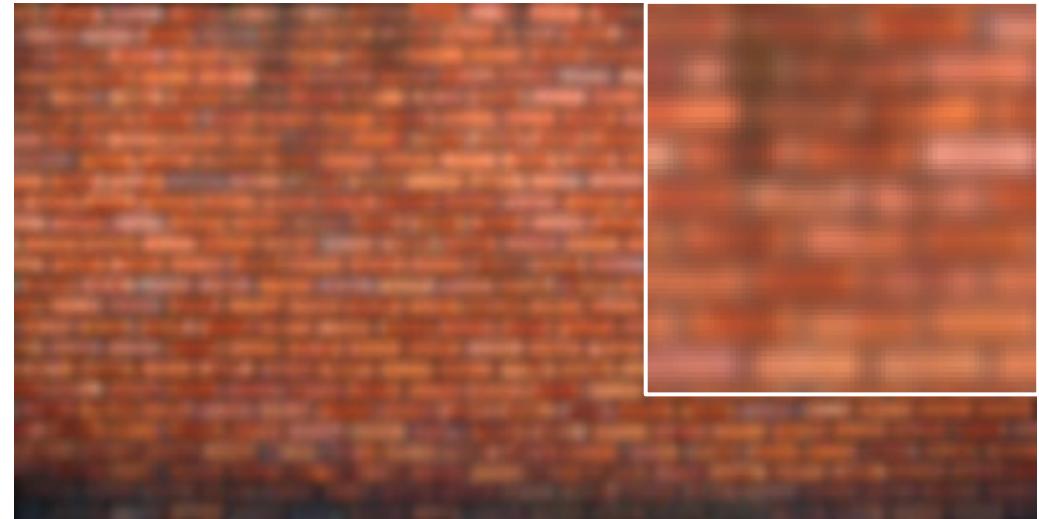


Another example

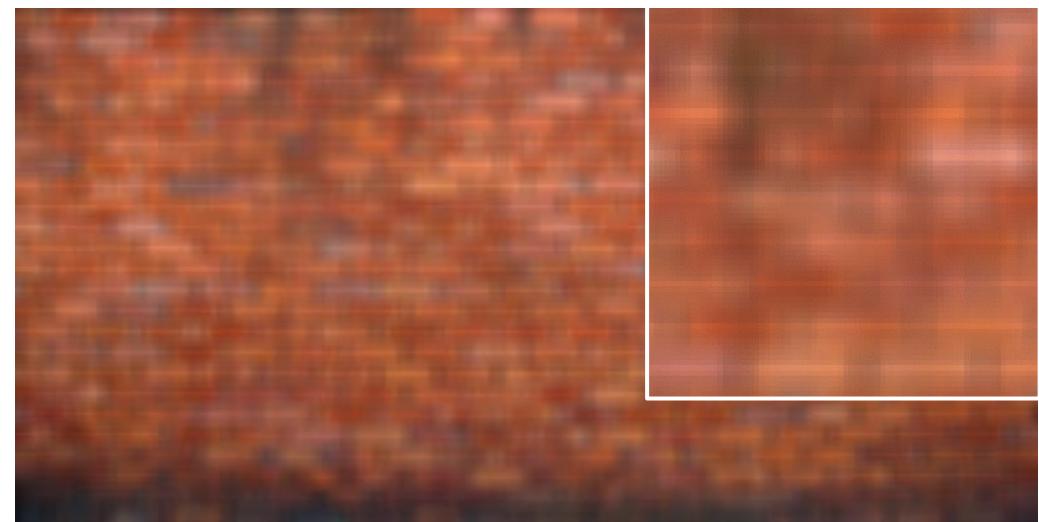


original

Which blur do you like better?



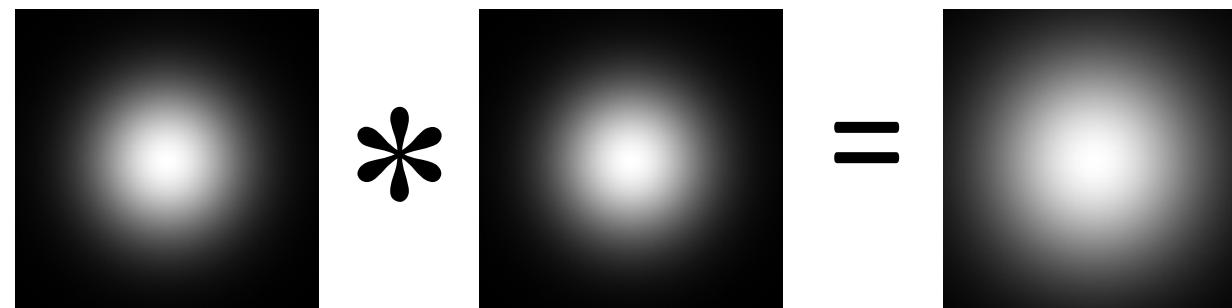
7x7 Gaussian



7x7 box

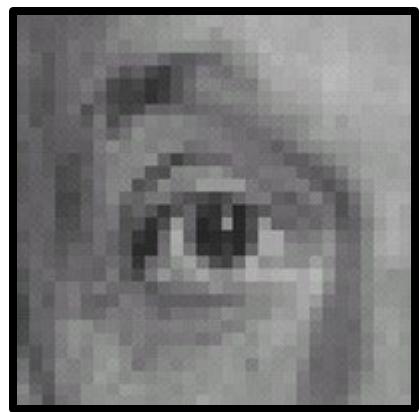
Gaussian filter

- Removes “high-frequency” components from the image (low-pass filter)
- Convolution with self is another Gaussian



- Convolving twice with Gaussian kernel of σ_{dth} = Convolving once with kernel $\sigma\sqrt{2}\text{dth}$

Linear filters: examples



$$\text{Original} * \left(\begin{array}{ccc} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{array} \right) - \frac{1}{9} \left(\begin{array}{ccc} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{array} \right) = \text{Sharpening filter output}$$



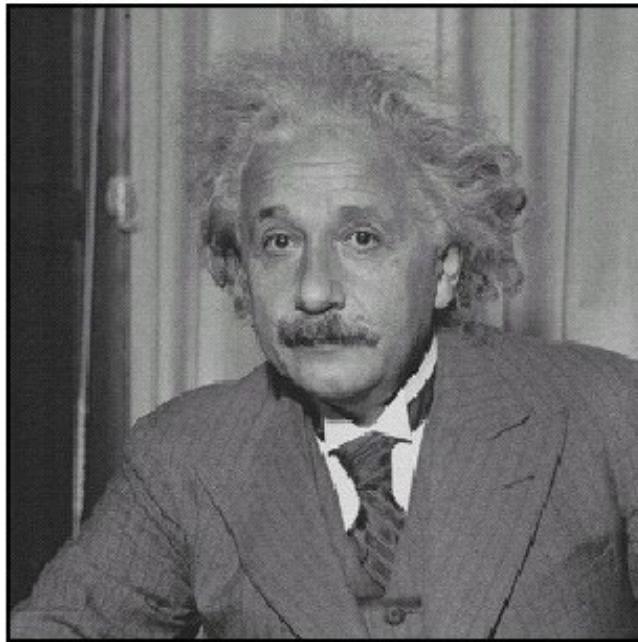
Sharpening filter
(accentuates edges)

Sharpening (derivation)

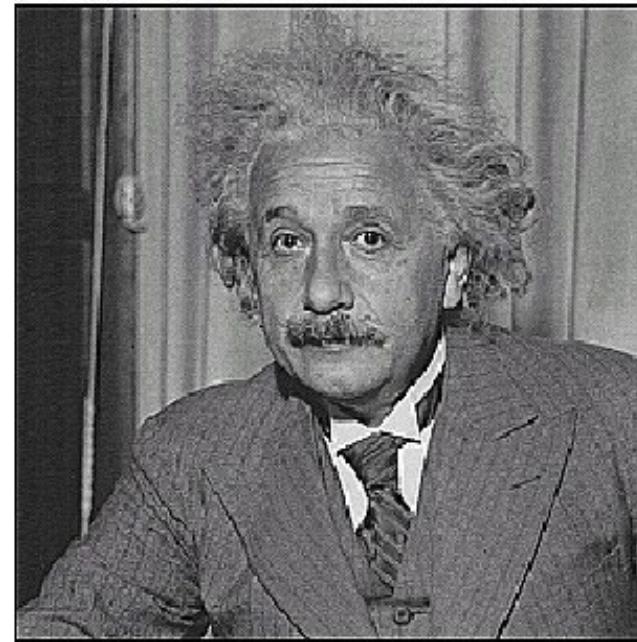


$$\begin{array}{|c|c|c|} \hline \bullet & 0 & 0 \\ \hline 0 & \bullet & 1 \\ \hline 0 & 0 & \bullet \\ \hline \end{array} + \begin{array}{|c|c|c|} \hline \bullet & 0 & 0 \\ \hline 0 & \bullet & 1 \\ \hline 0 & 0 & \bullet \\ \hline \end{array} - \frac{1}{9} \begin{array}{|c|c|c|} \hline \bullet & 1 & 1 \\ \hline 1 & \bullet & 1 \\ \hline 1 & 1 & \bullet \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline \bullet & 0 & 0 \\ \hline 0 & \bullet & 2 \\ \hline 0 & 0 & \bullet \\ \hline \end{array} - \frac{1}{9} \begin{array}{|c|c|c|} \hline \bullet & 1 & 1 \\ \hline 1 & \bullet & 1 \\ \hline 1 & 1 & \bullet \\ \hline \end{array}$$

Sharpening



before



after

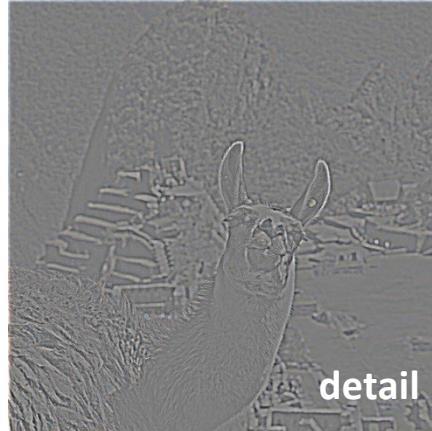
Sharpening



original



smoothed (5x5)



detail

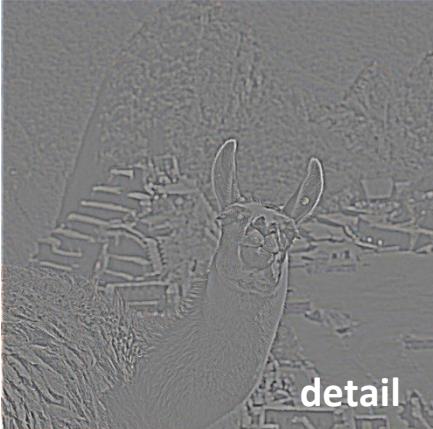
(This “detail extraction” operation is also called a ***high-pass filter***)

Let's add it back:



original

$+\alpha$



detail

$=$

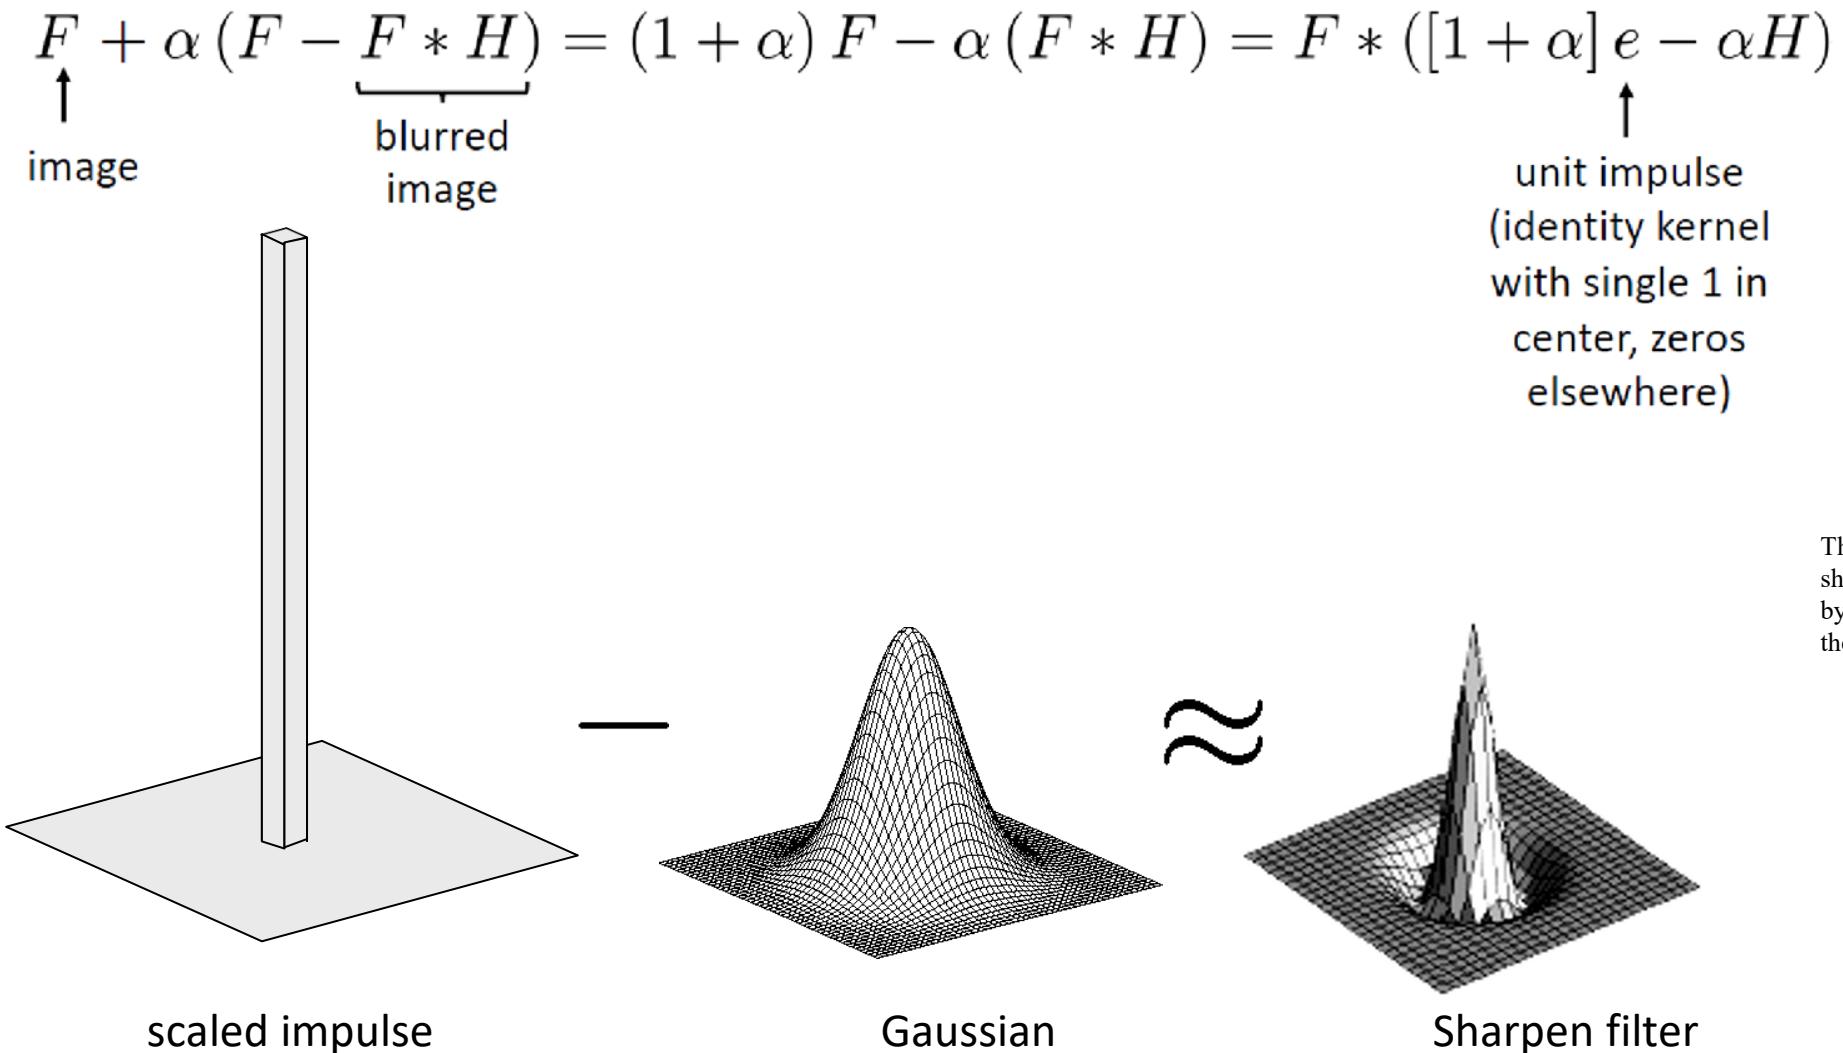


sharpened

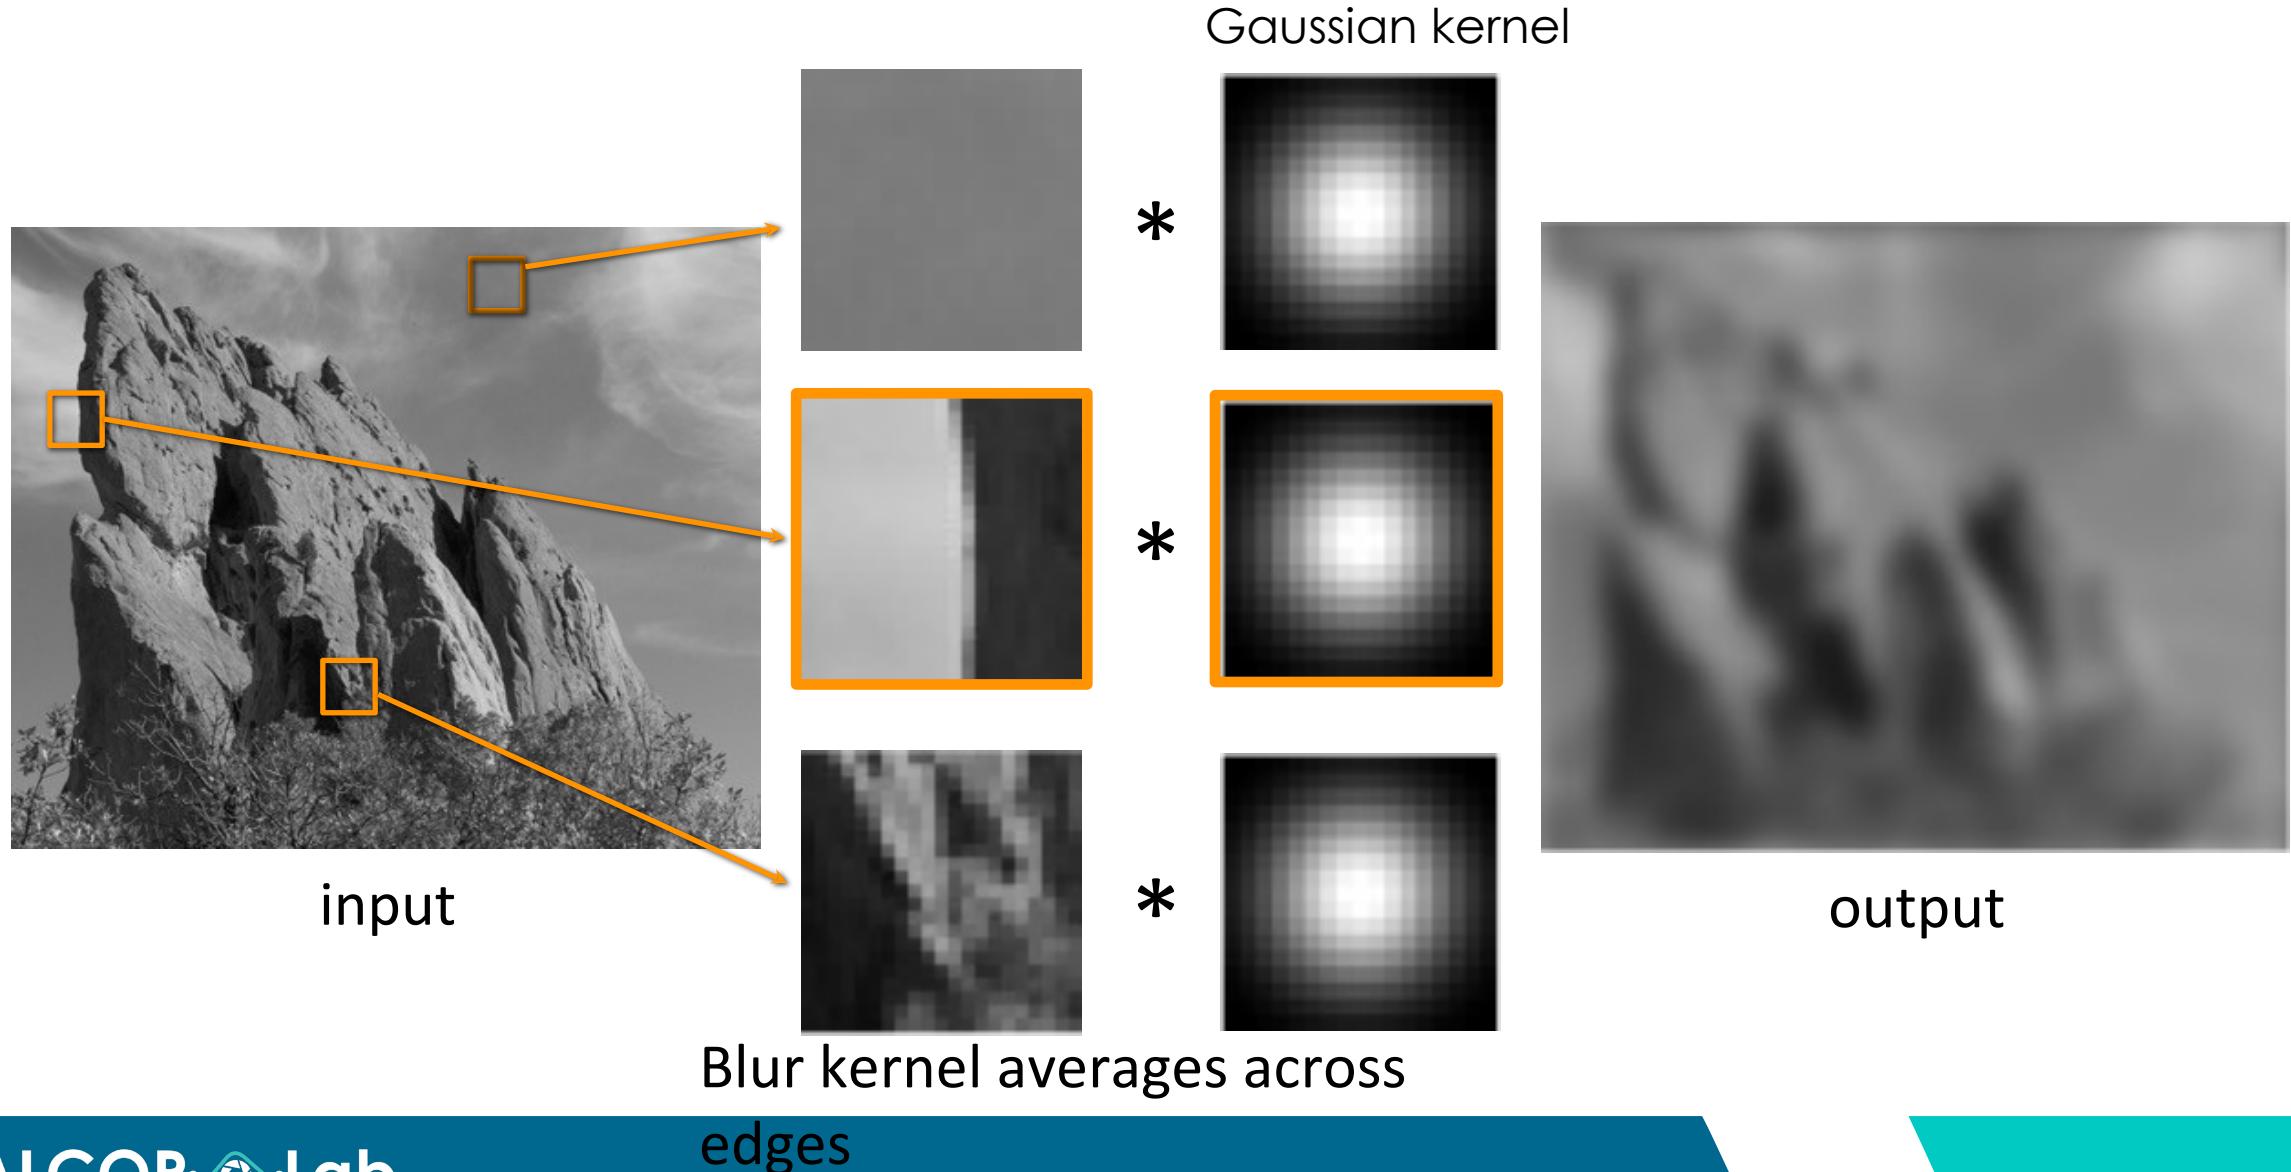
Photo credit:

<https://www.flickr.com/photos/geezaweezer/16089096376/>

Sharpen filter

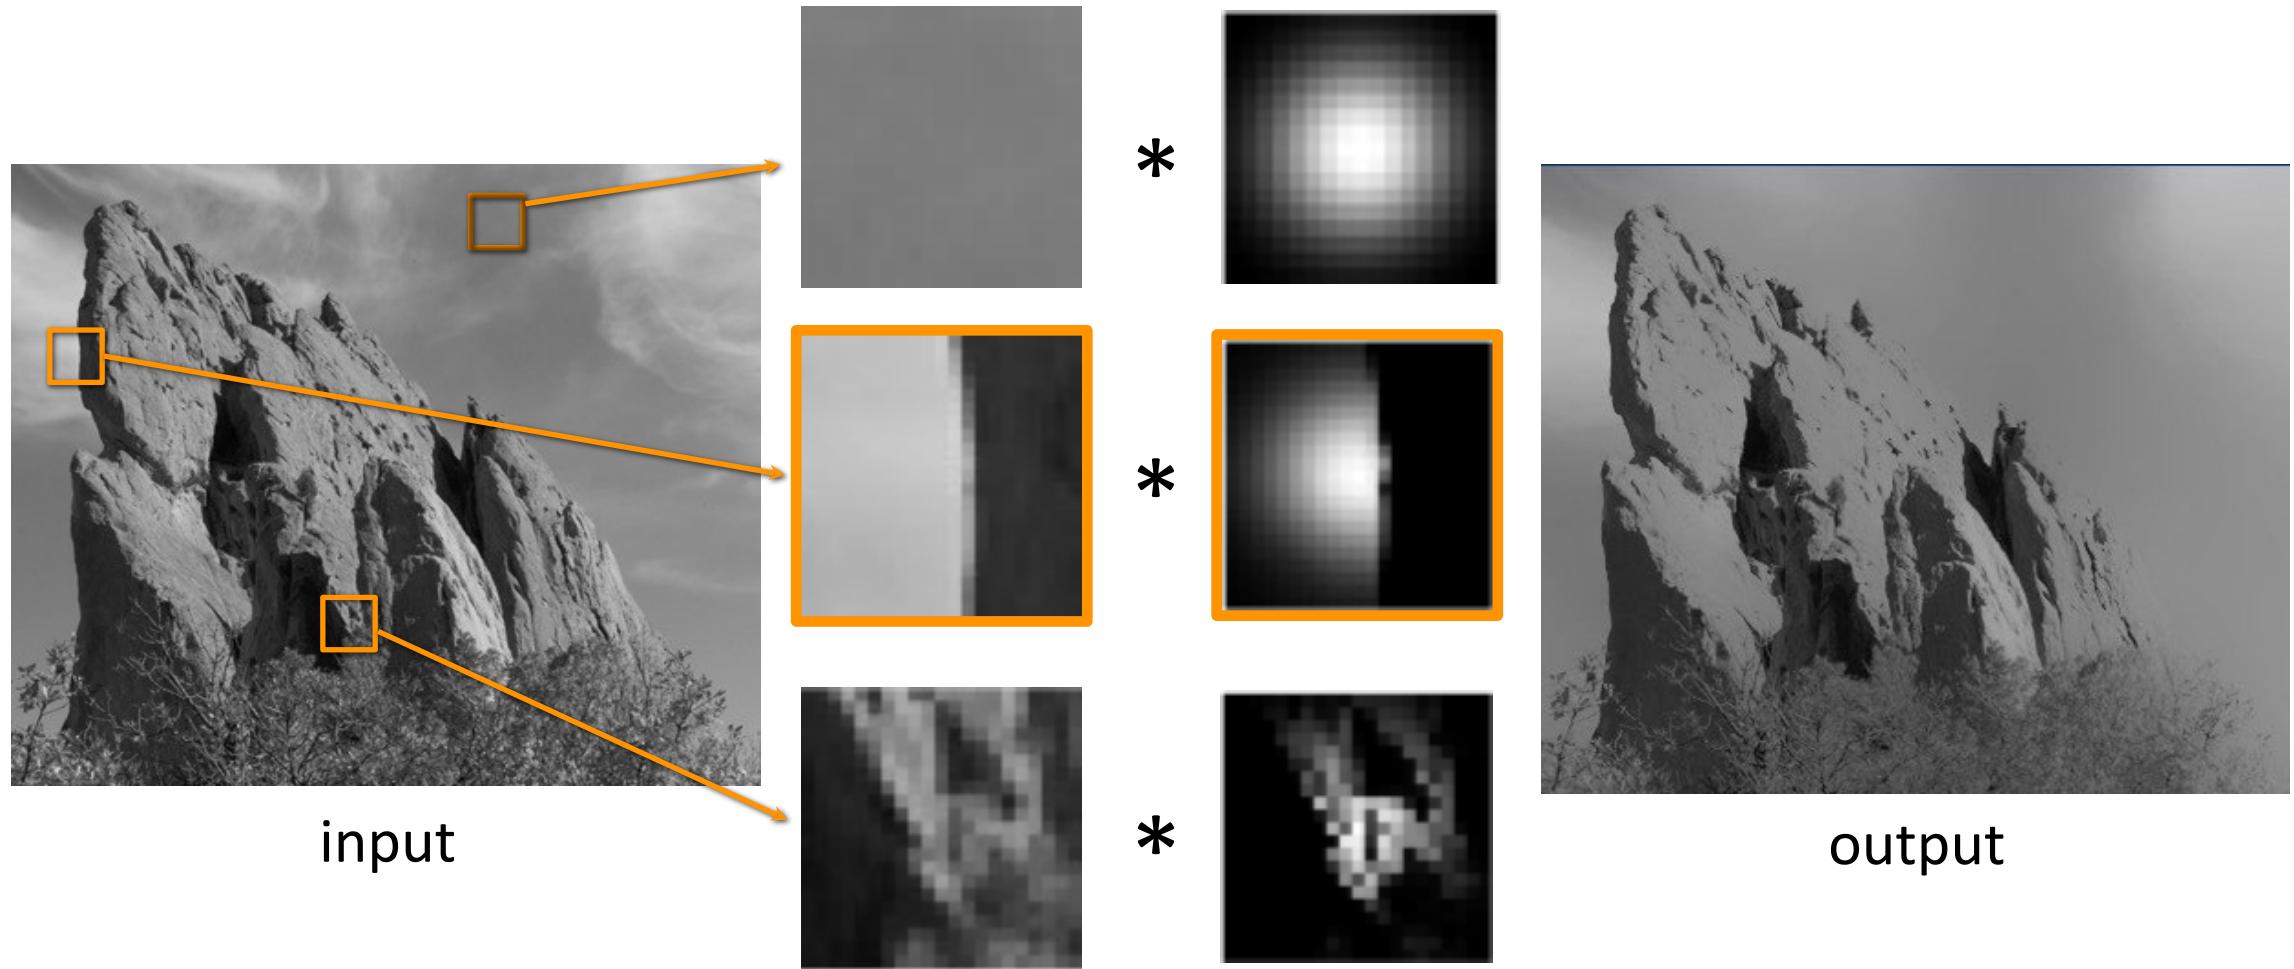


The problem with Gaussian filtering



The bilateral filtering solution

Bilateral filter kernel



Do not blur if there is an edge! How does it do that?

Bilateral filtering vs Gaussian filtering

Gaussian
filtering

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

Bilateral
filtering

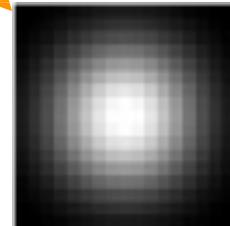
$$h[m, n] = \frac{1}{W_{mn}} \sum_{k,l} g[k, l] r_{mn}[k, l] f[m + k, n + l]$$

Bilateral filtering vs Gaussian filtering

Gaussian
filtering

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

σ_s



Spatial weighting:
favor *nearby*
pixels

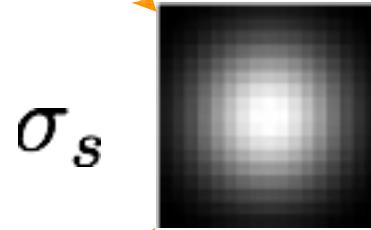
Bilateral
filtering

$$h[m, n] = \frac{1}{W_{mn}} \sum_{k,l} g[k, l] r_{mn}[k, l] f[m + k, n + l]$$

Bilateral filtering vs Gaussian filtering

Gaussian
filtering

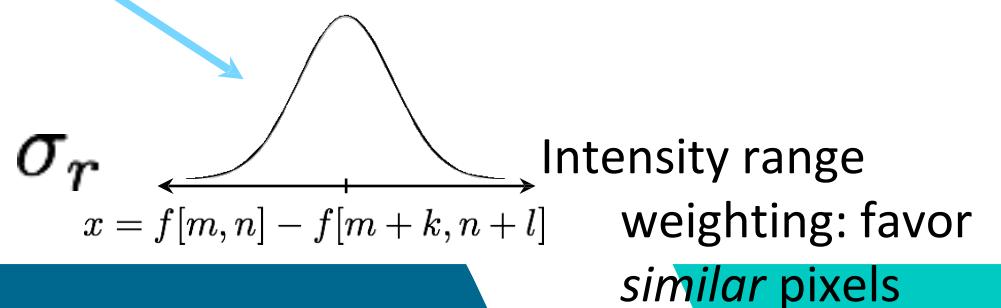
$$h[m, n] = \sum_{k, l} g[k, l] f[m + k, n + l]$$



Spatial weighting:
favor *nearby*
pixels

Bilateral
filtering

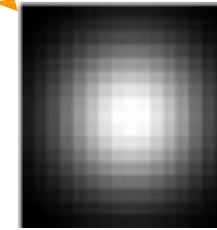
$$h[m, n] = \frac{1}{W_{mn}} \sum_{k, l} g[k, l] r_{mn}[k, l] f[m + k, n + l]$$



Bilateral filtering vs Gaussian filtering

Gaussian
filtering

$$h[m, n] = \sum_{k, l} g[k, l] f[m + k, n + l]$$

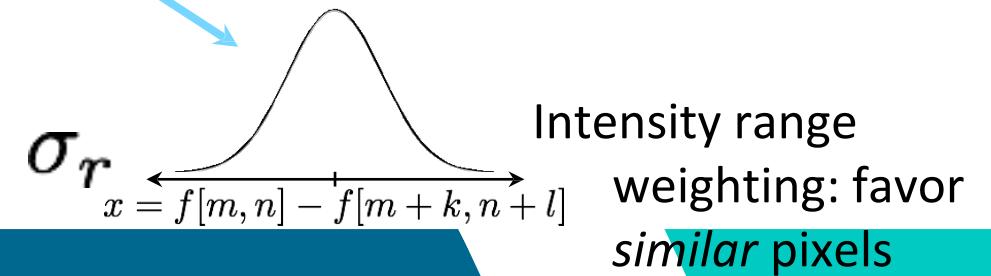


Spatial weighting:
favor *nearby*
pixels

Bilateral
filtering

$$h[m, n] = \frac{1}{W_{mn}} \sum_{k, l} g[k, l] r_{mn}[k, l] f[m + k, n + l]$$

Normalization
factor



Bilateral filtering vs Gaussian filtering

Gaussian filtering

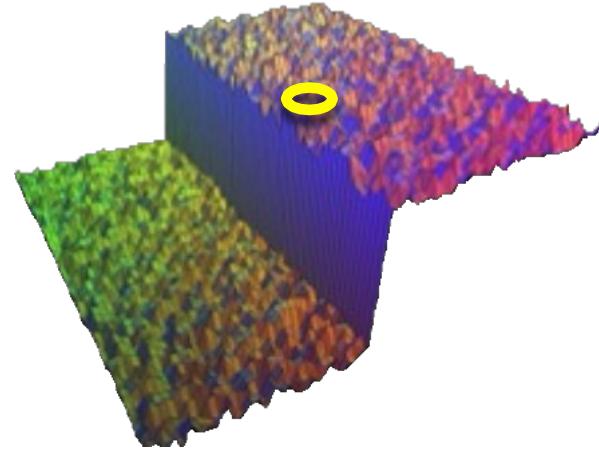
Smooths everything nearby (even edges) Only depends on *spatial* distance

Bilateral filtering

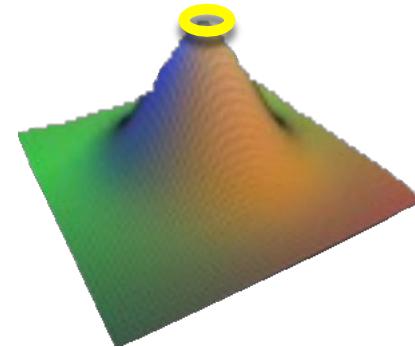
Smooths ‘close’ pixels in space and intensity Depends on *spatial* and *intensity* distance

Gaussian filtering visualization

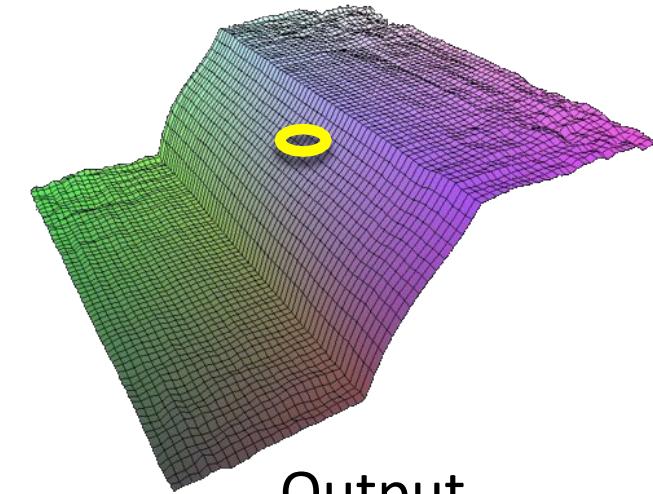
$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$



Input



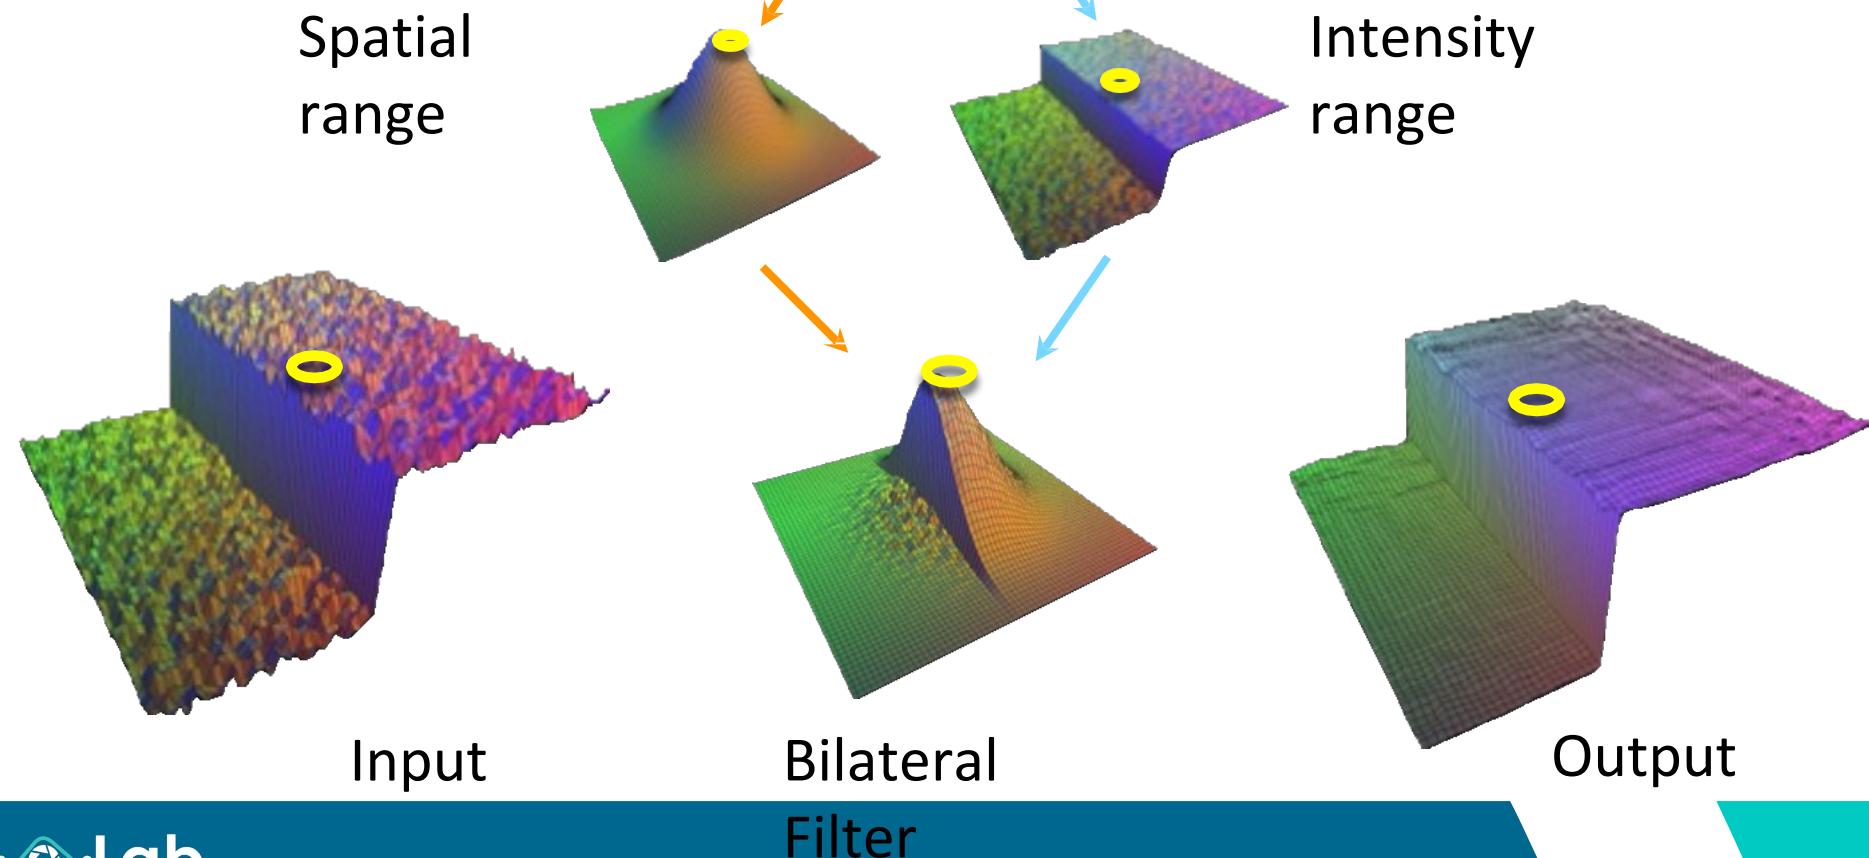
Gaussian
Filter



Output

Bilateral filtering visualization

$$h[m, n] = \frac{1}{W_{mn}} \sum_{k,l} g[k, l] r_{mn}[k, l] f[m + k, n + l]$$



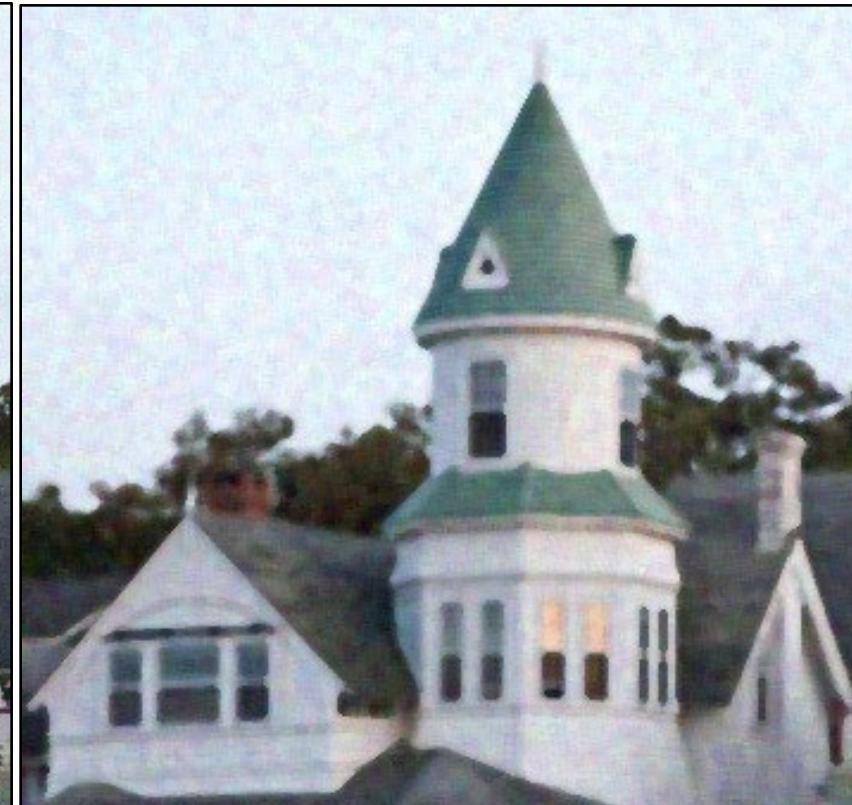
Denoising



noisy
input



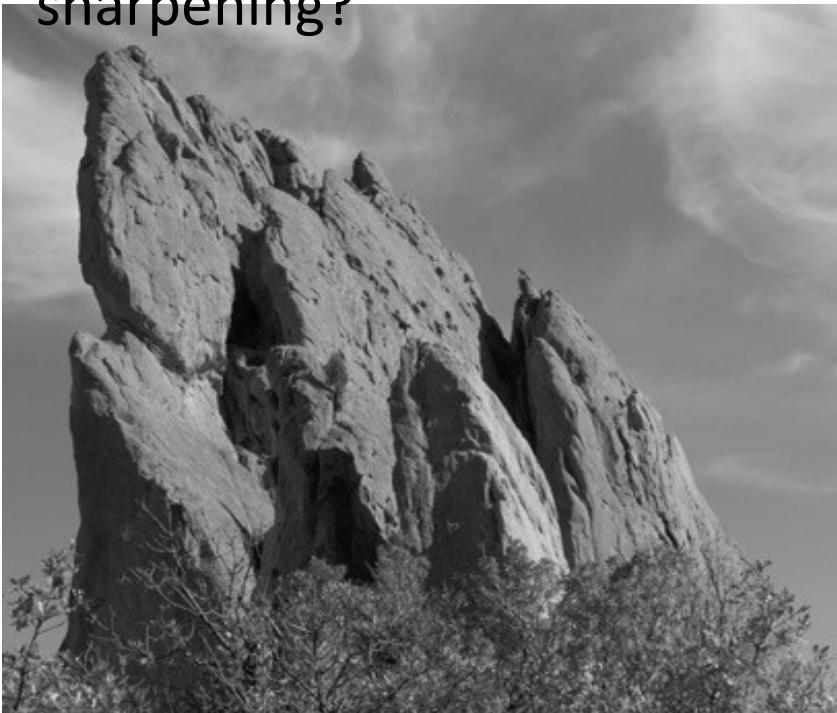
bilateral
filtering



median
filtering

Contrast enhancement

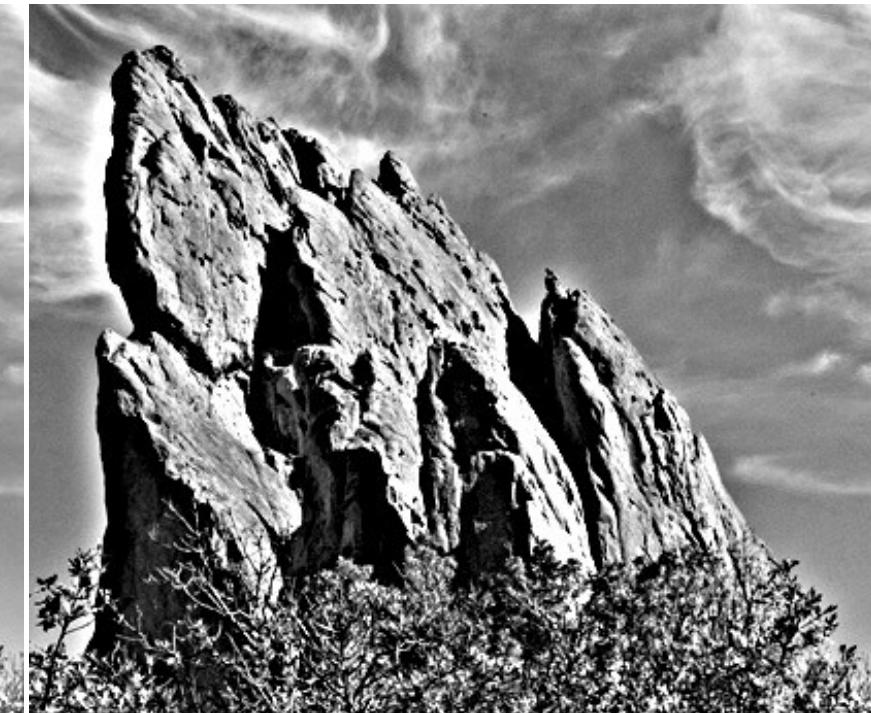
How would you use Gaussian or bilateral filtering for sharpening?



input



sharpening based
on bilateral
filtering



sharpening based
on Gaussian
filtering

Photo retouching



Photo retouching



original
|



digital pore removal (aka bilateral
filtering)

Close-up comparison



original
|



digital pore removal (aka bilateral
filtering)

Cartoonization



input

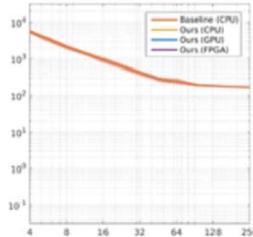
The bilateral filter will reduce the color palette, or the numbers of colors that are used in the image.

Then we can apply edge detection to the resulting image to generate bold silhouettes.



cartoon
rendition

Actively used in various research problems



A Hardware-Friendly Bilateral Solver for Real-Time Virtual Reality Video

Amrita Mazumdar, Armin Alaghi, **Jonathan T. Barron**, David Gallup, Luis Ceze, Mark Oskin, Steven M. Seitz

High-Performance Graphics (HPG), 2017
[project page](#)

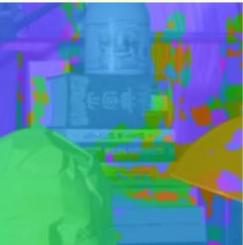
A reformulation of the **bilateral** solver can be implemented efficiently on GPUs and FPGAs.



Deep Bilateral Learning for Real-Time Image Enhancement

Michaël Gharbi, Jiawen Chen, **Jonathan T. Barron**, Samuel W. Hasinoff, Frédo Durand
SIGGRAPH, 2017
[project page](#) / [video](#) / [bibtex](#) / [press](#)

By training a deep network in **bilateral** space we can learn a model for high-resolution and real-time image enhancement.



The Fast Bilateral Solver

Jonathan T. Barron, Ben Poole
ECCV, 2016 ([Oral Presentation](#), [Best Paper Honorable Mention](#))
[arXiv](#) / [bibtex](#) / [video](#) (they messed up my slides, use →) / [keynote](#) (or [PDF](#)) / [code](#) / [depth super-res results](#) / [reviews](#)

Our solver smooths things better than other filters and faster than other optimization algorithms, and you can backprop through it.

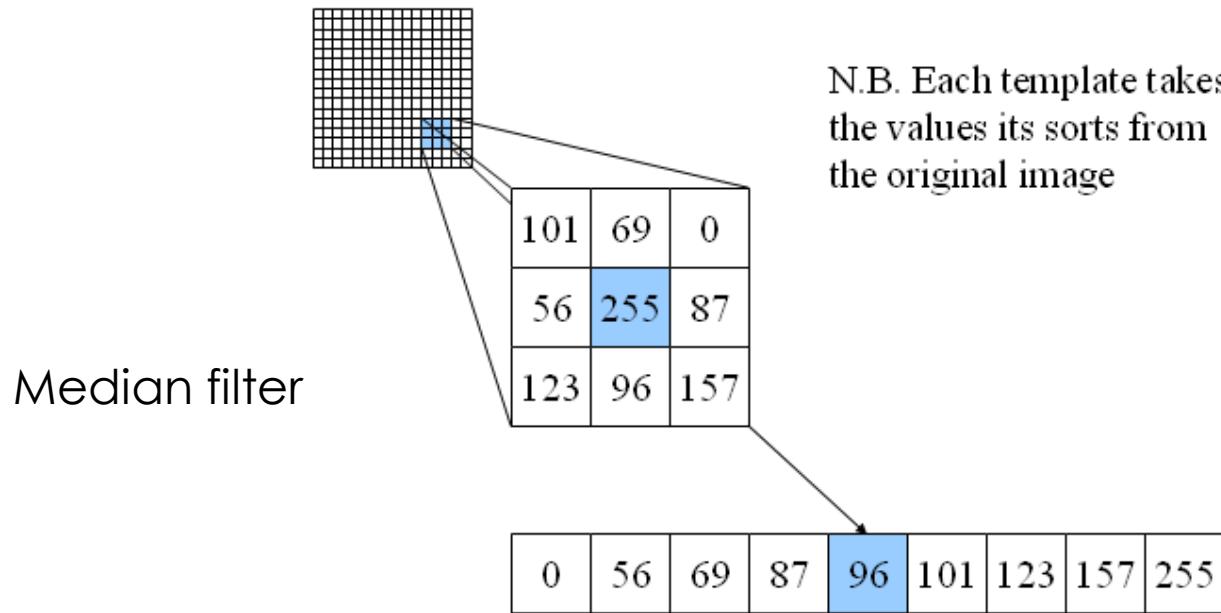
Filters: Thresholding



$$g(m, n) = \begin{cases} 255, & f(m, n) > A \\ 0 & otherwise \end{cases}$$

Non-Linear filters

- Is thresholding a linear filter?
- No is a non-linear filter together with image equalization or median filters



Salt and pepper noise



Image Transformations & Filtering

- Point Processing
- Linear Filtering
- Sampling & Aliasing
- Image Derivatives
- Edge Detection

Image resizing

This image is too big to fit on the screen.
How can we reduce it?

How to generate a half-sized version?

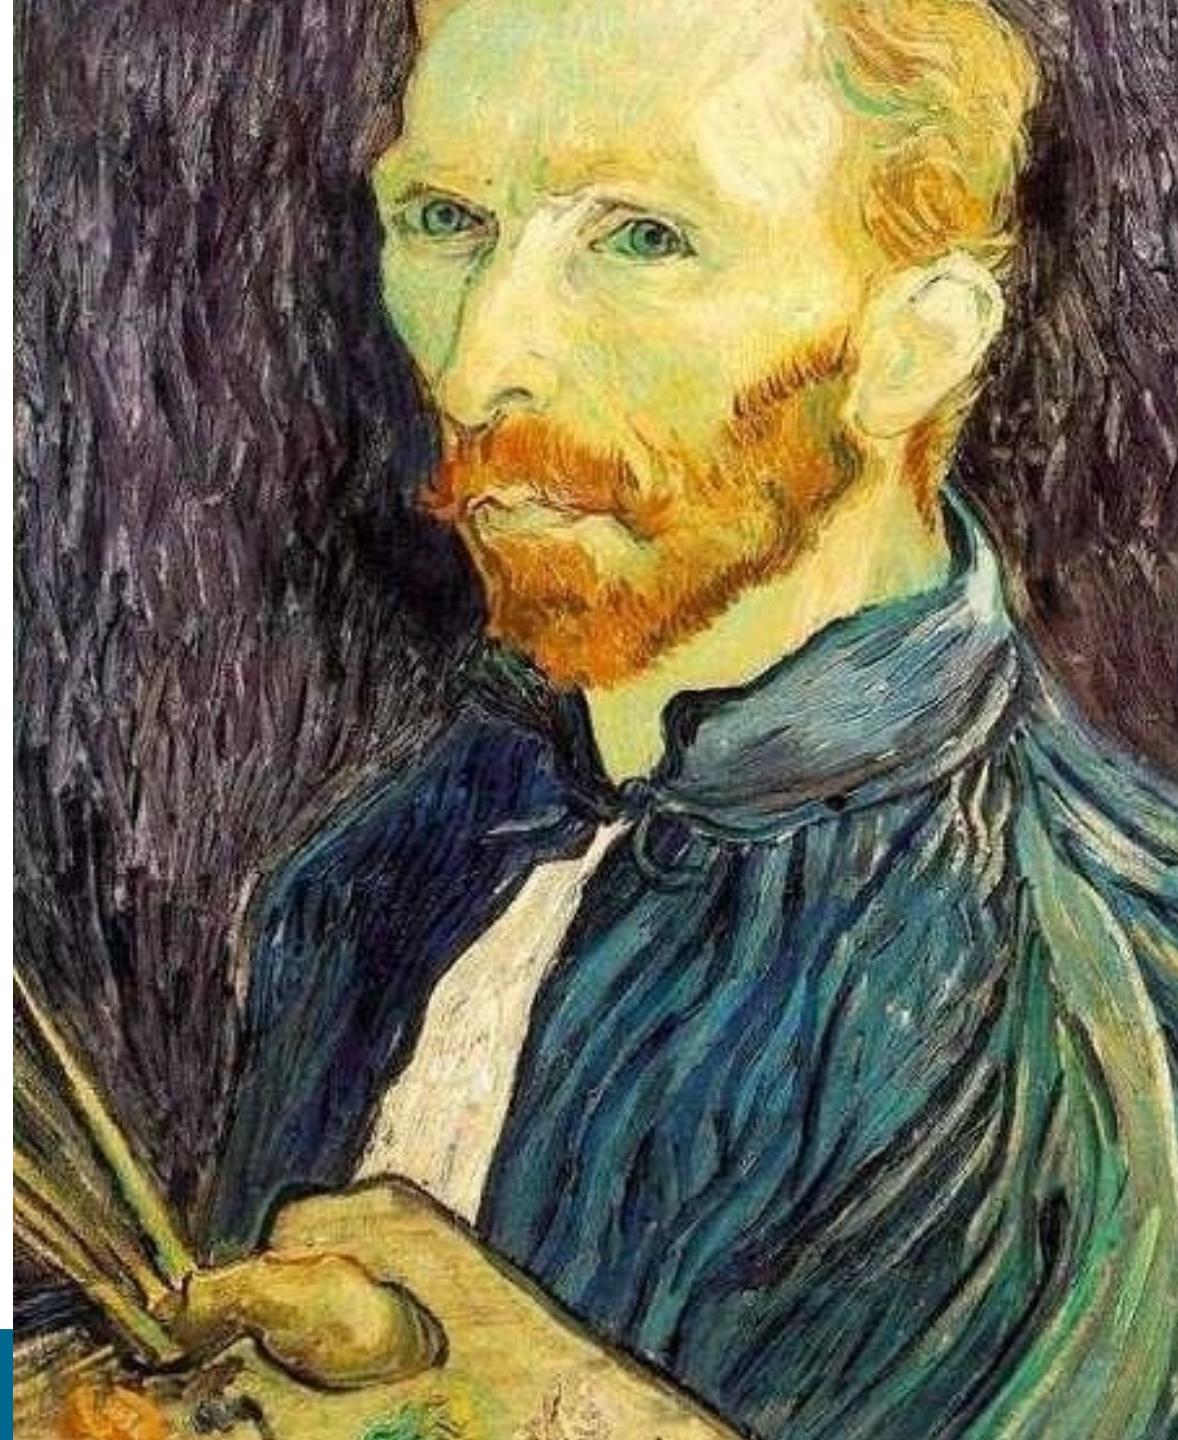
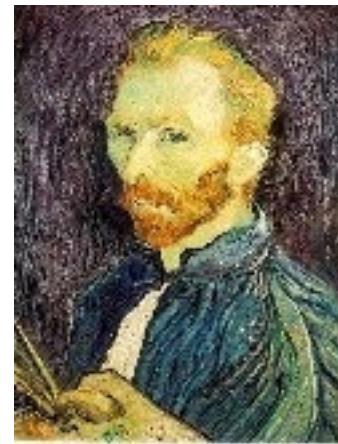
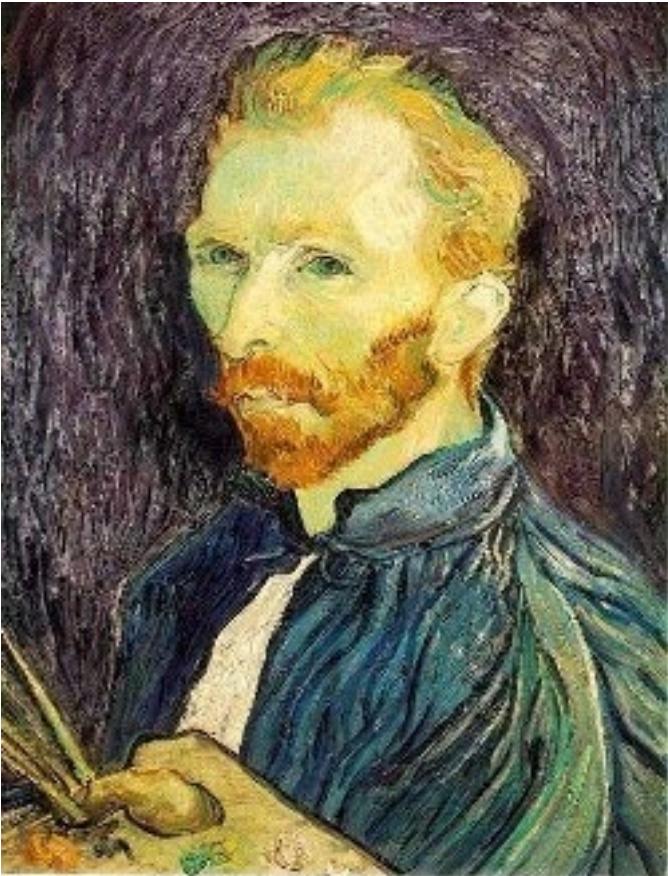


Image sub-sampling

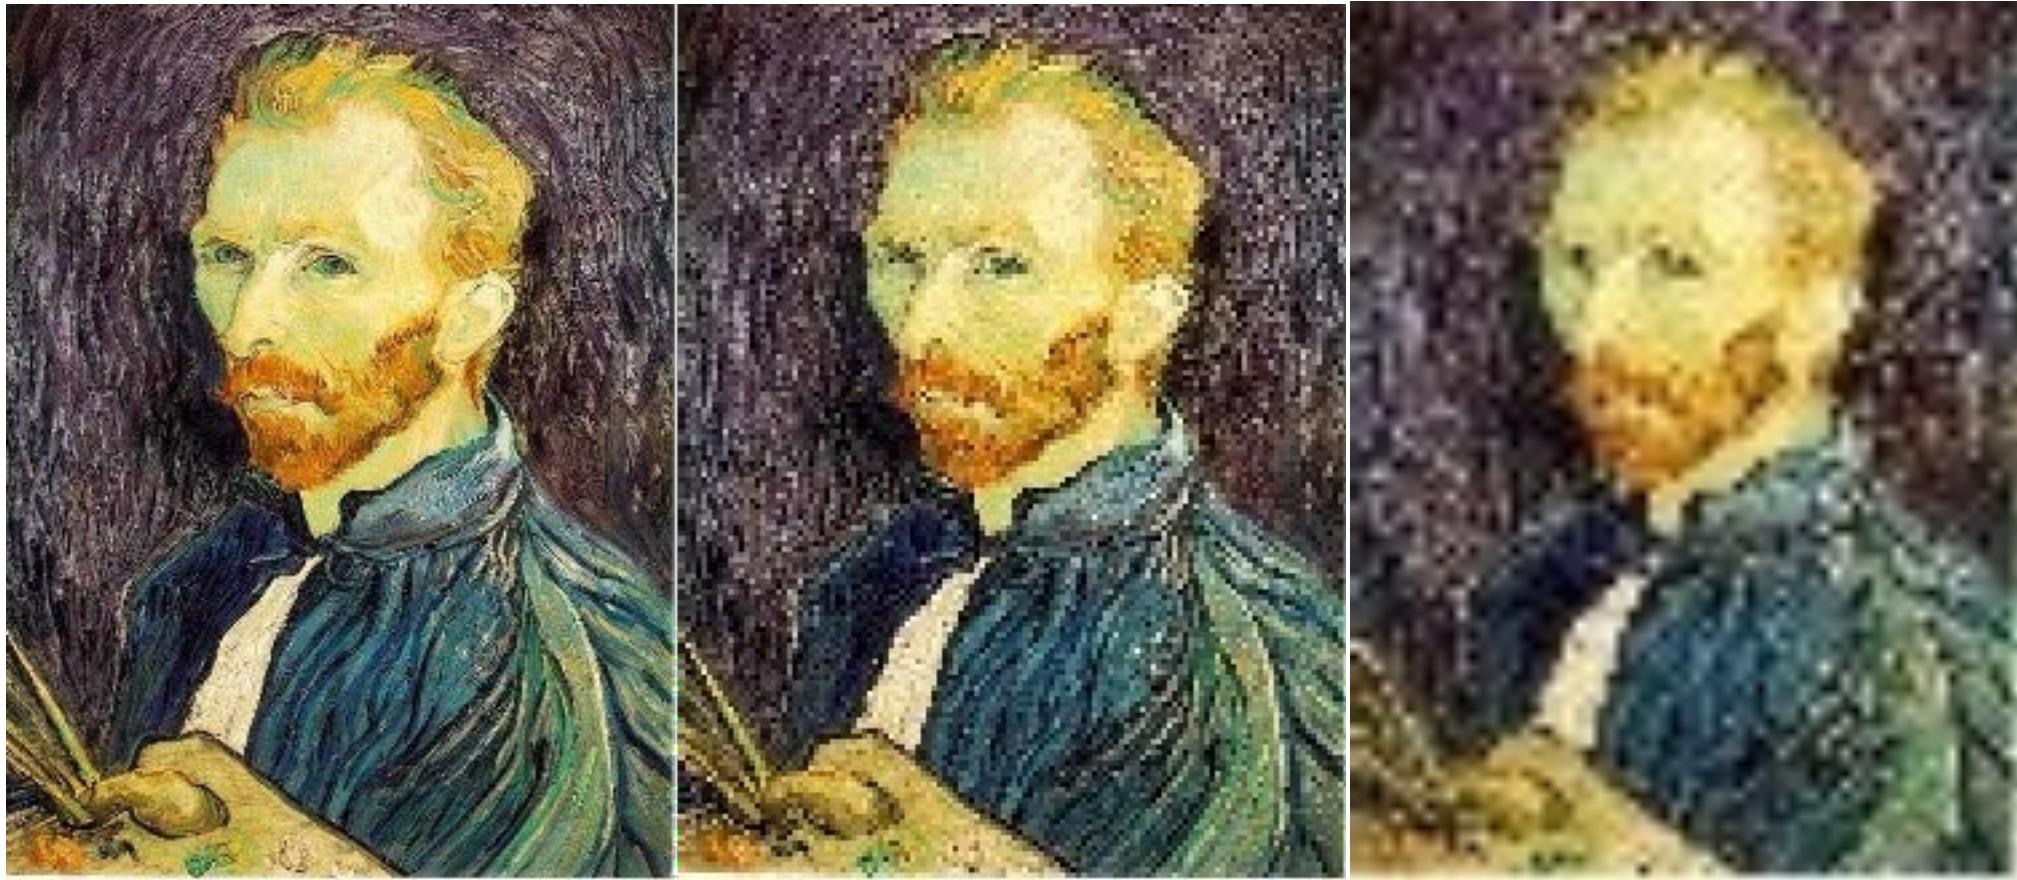


1/4



1/8

Throw away every other row and column to create a $1/2$ size image
- called *image sub-sampling*



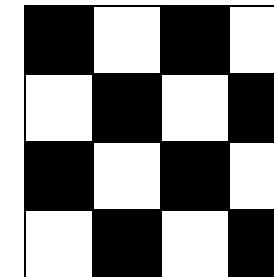
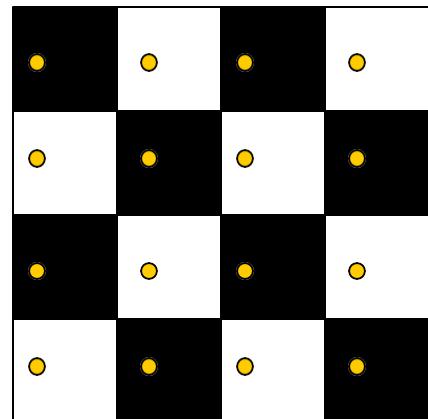
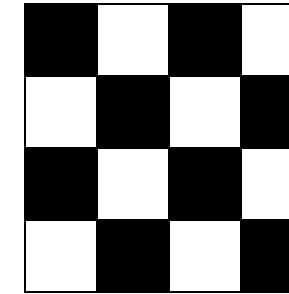
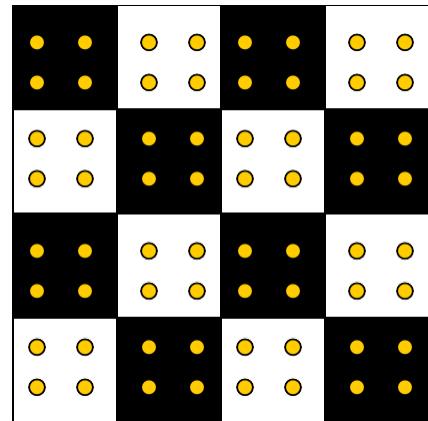
1/2

1/4
(2x
zoom)

1/8
(4x
zoom)

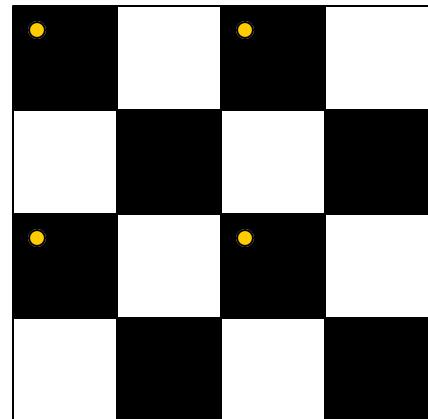
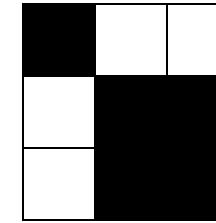
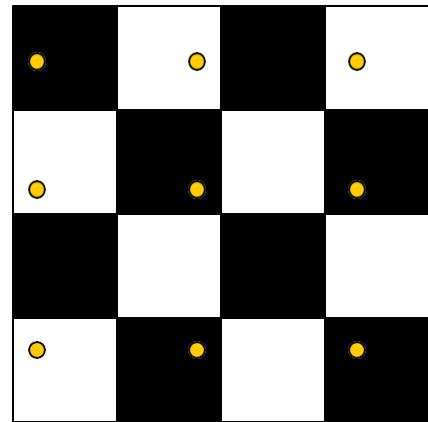
Aliasing!
What do we
do?

Sampling an image



Examples of GOOD sampling

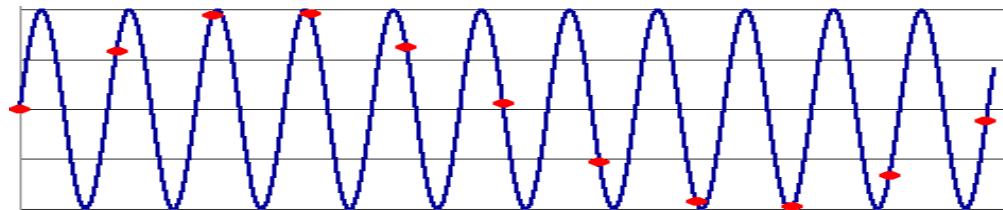
Undersampling



Examples of BAD sampling ->
Aliasing

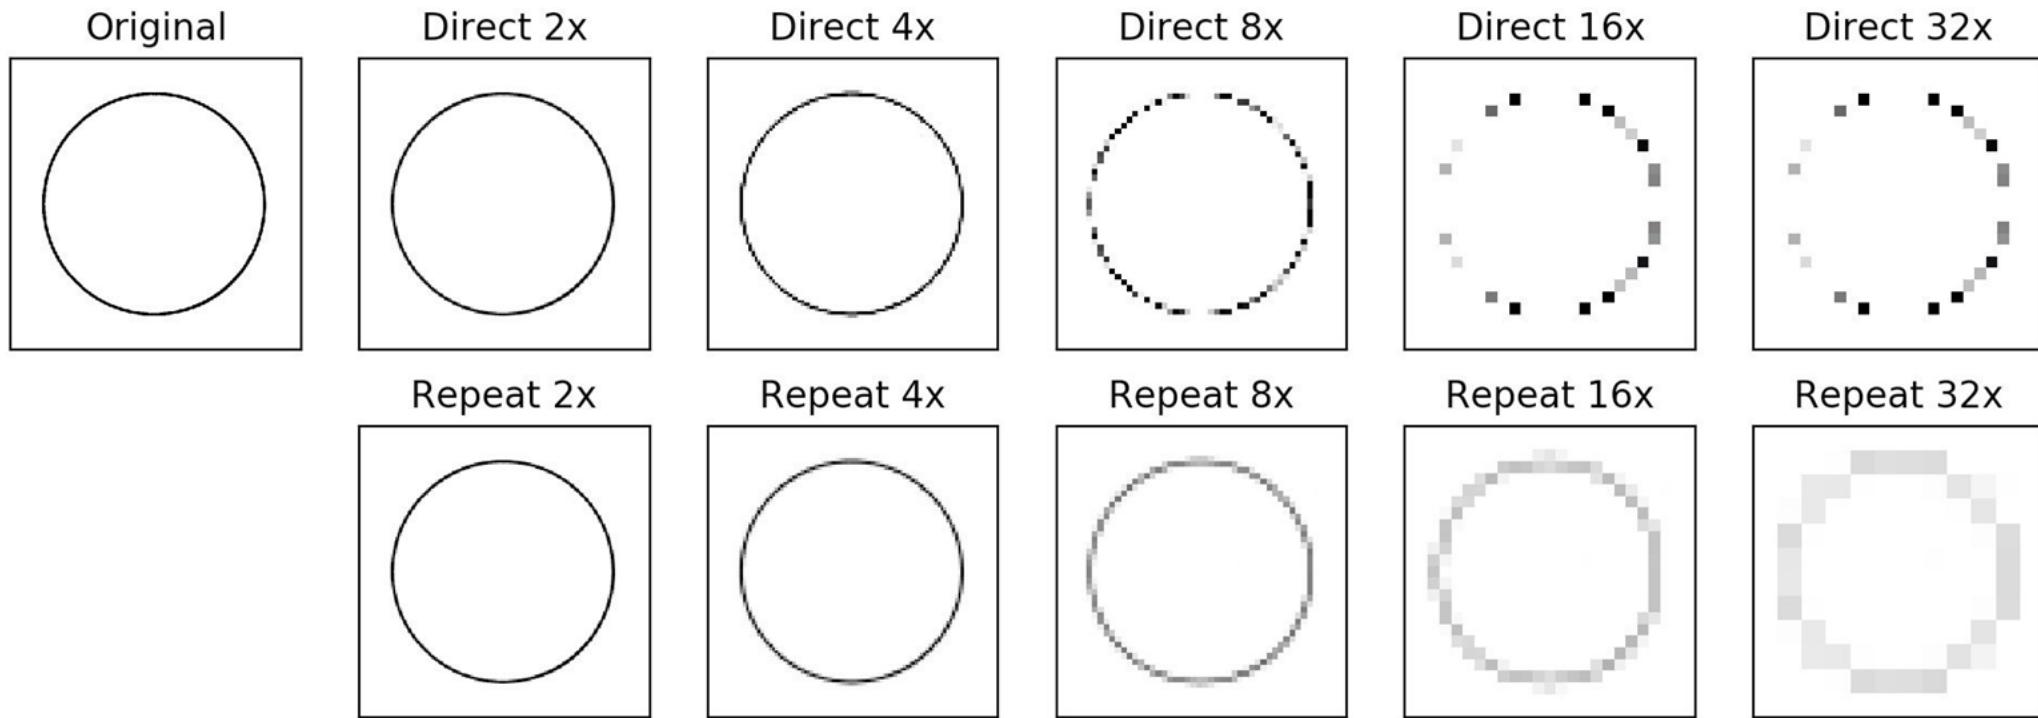
Aliasing

- Occurs when your sampling rate is not high enough to capture the amount of detail in your image
- Can give you the wrong signal/image—an *alias*
- To do sampling right, need to understand the structure of your signal/image
- Fourier Transform (we will talk about it in the next class)
- To avoid aliasing:
 - sampling rate $\geq 2 * \text{max frequency in the image}$
 - said another way: $\geq \text{two samples per cycle}$
 - this minimum sampling rate is called the **Nyquist rate**



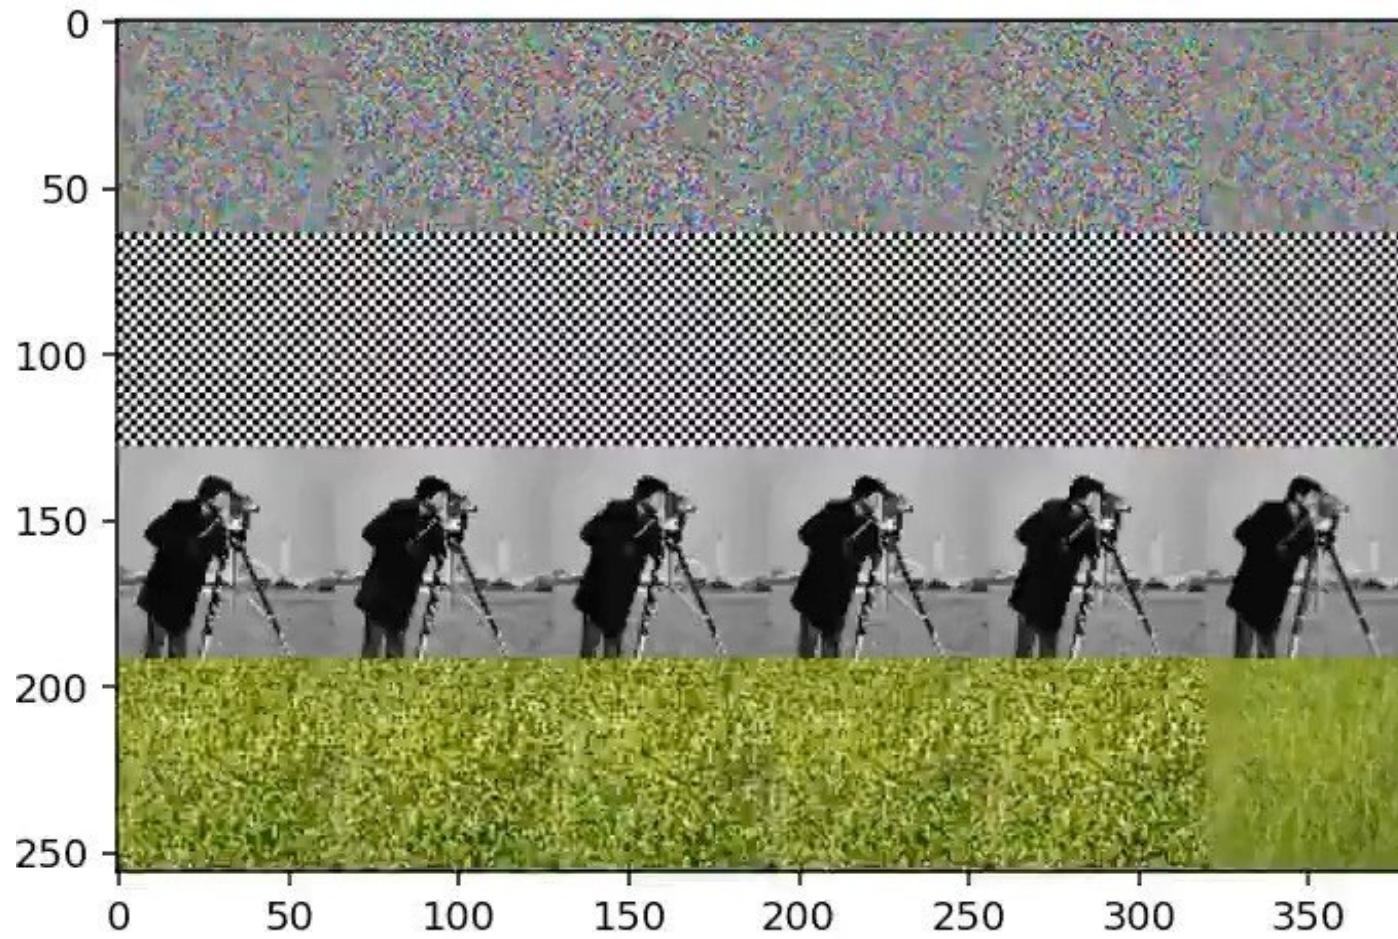
A real problem!

Bilinear downsampling with F.interpolate() in PyTorch



A real problem!

128 x 128 □ 64x64



Open-CV:
default, bicubic, Lanczos4

Pytorch:
bilinear,
bicubic

PIL: Lanczos
credit: @jaakkolehtinen

Problems in NN too

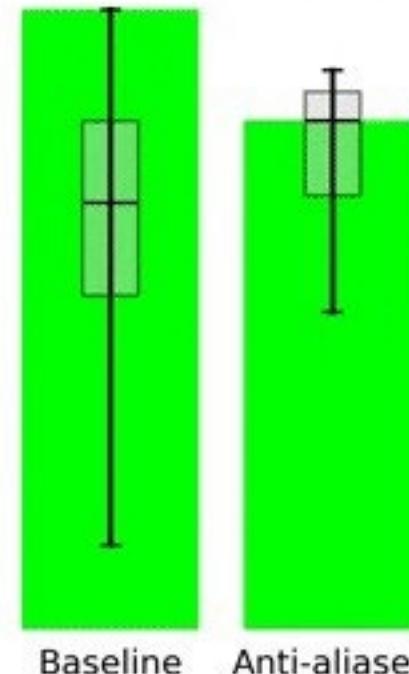


Anti-aliasing in CNNs

```
pip install antialiased-cnns
```

84.5
Baseline

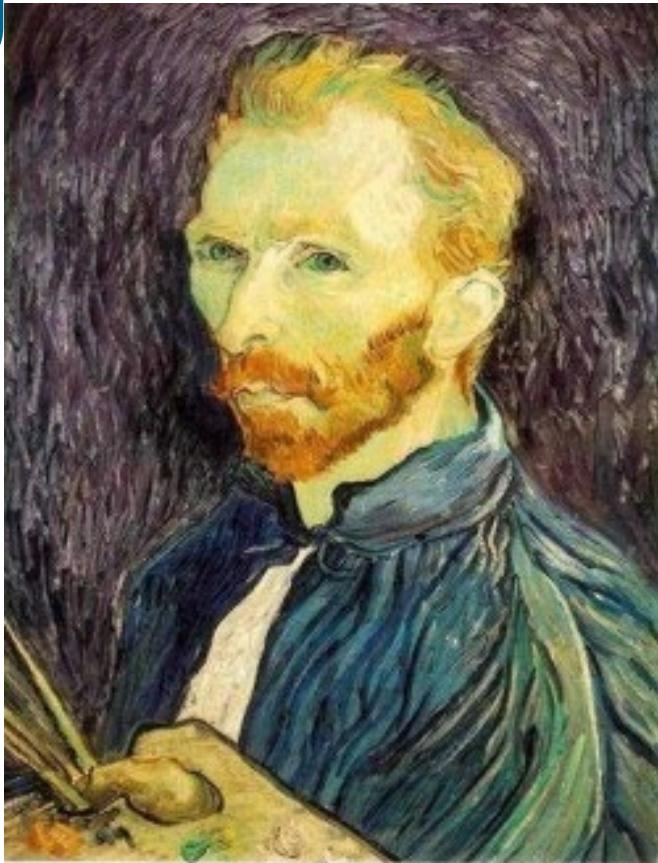
69.3
Anti-aliased



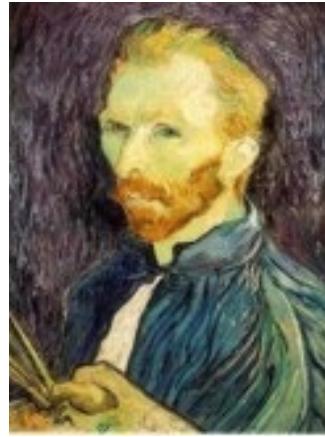
Classification stability:
predicted probability of
the correct class changes
when shifting the image.
Anti-aliased max-pooling.
Pooling does not preserve
shift-equivariance

Making Convolutional Networks Shift-Invariant Again, Richard Zhang ICML 2019

Gaussian (lowpass) pre-filtering



Gaussian 1/2



G 1/4

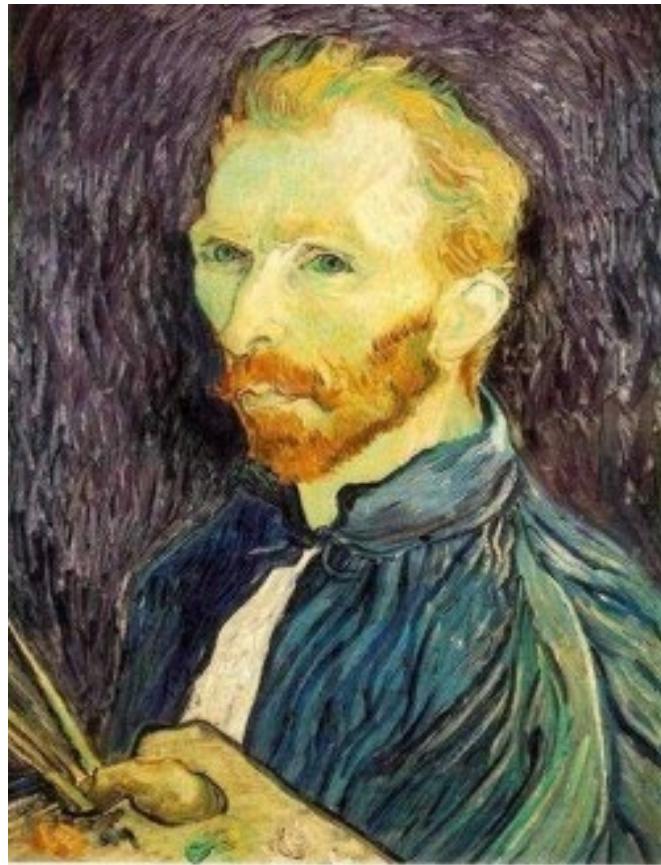


G 1/8

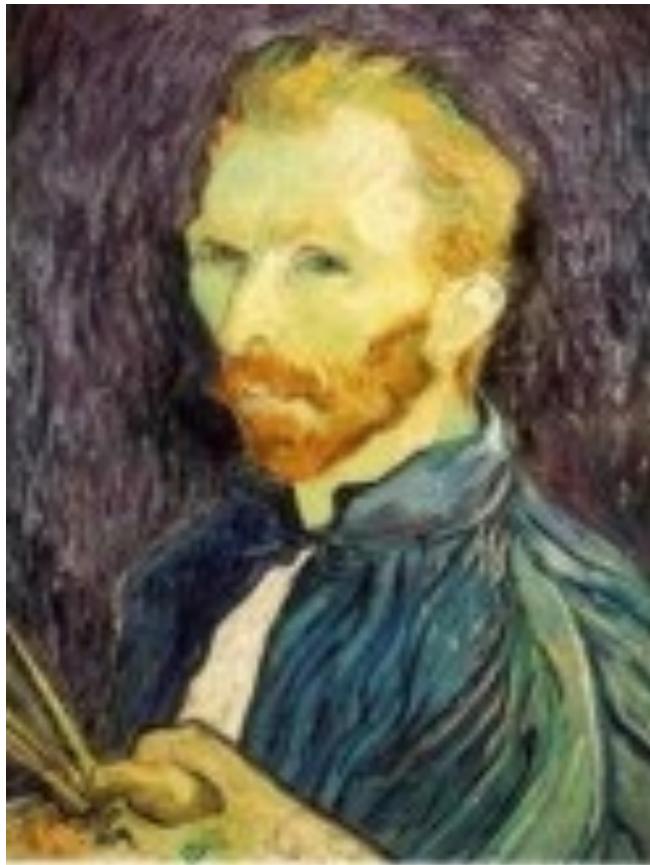
Solution: filter the image, *then* subsample

- Filter size should double for each $\frac{1}{2}$ size reduction.

Subsampling with Gaussian pre-filtering



Gaussian 1/2

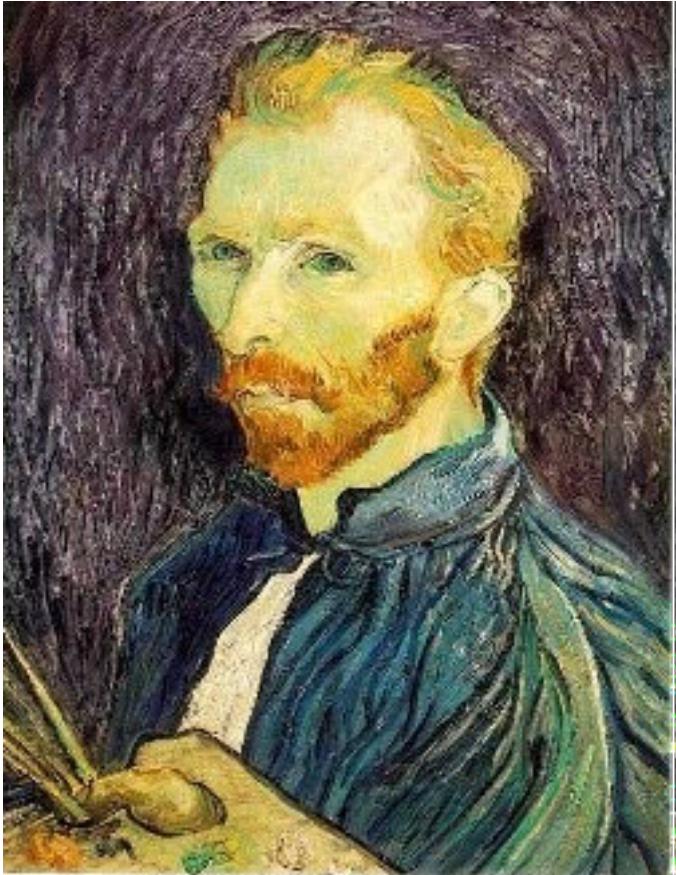


$G \frac{1}{4}$ (2x zoom)



$G \frac{1}{8}$ (4x zoom)

Compare with...



1/2



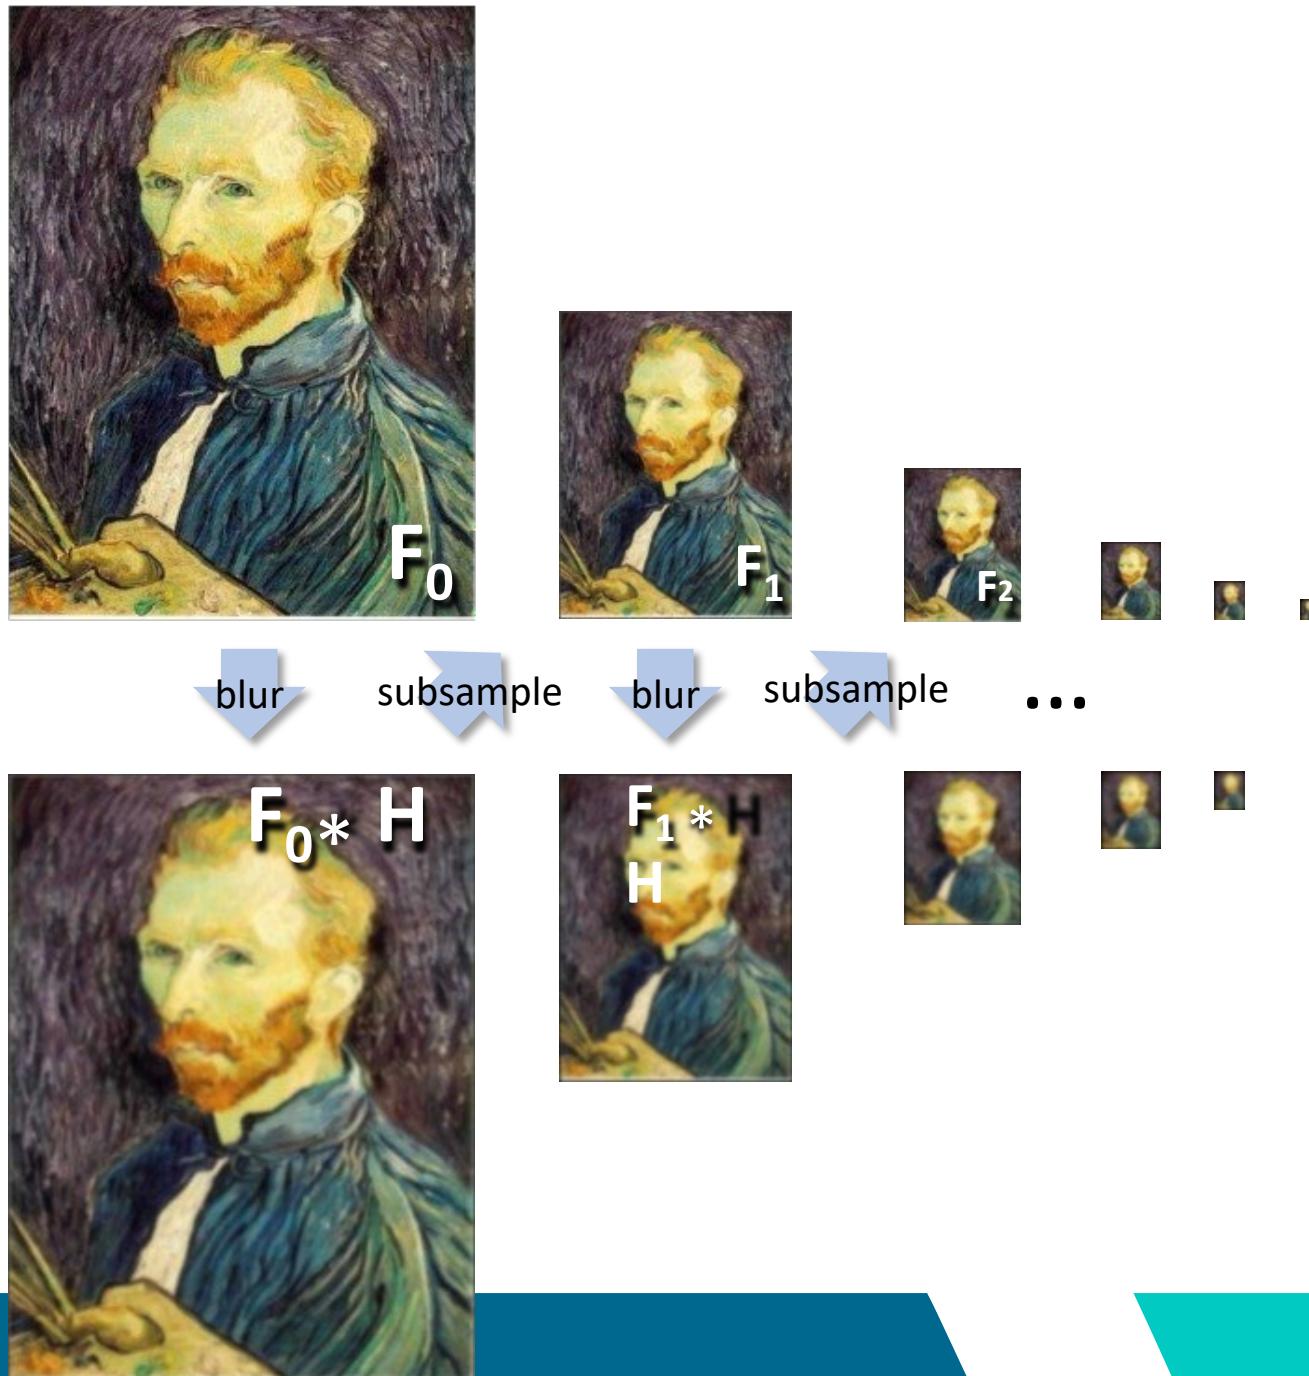
1/4
(2x
zoom)



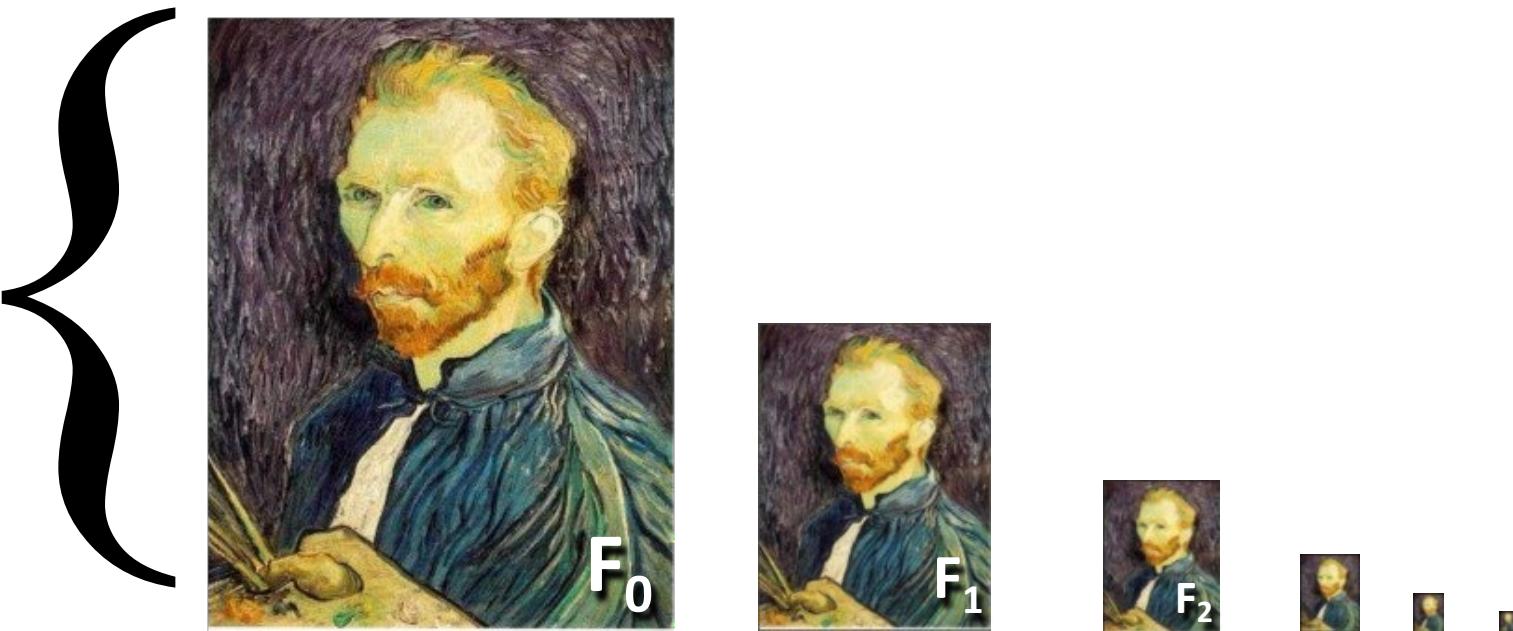
1/8
(4x
zoom)

Gaussian pre-filtering

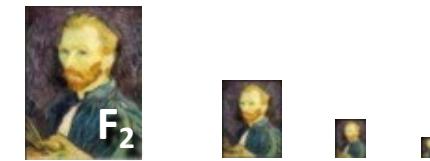
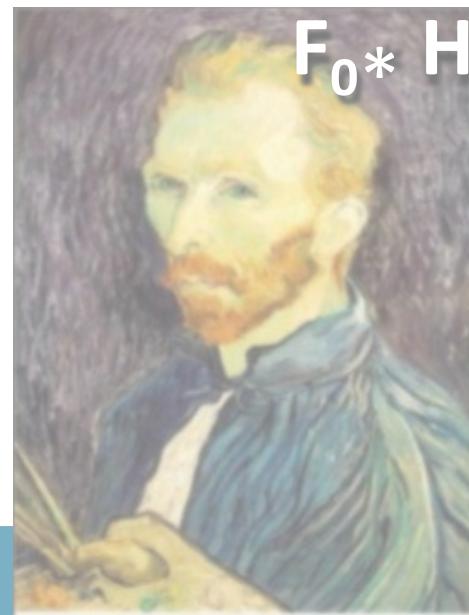
- Solution: filter the image, *then* subsample

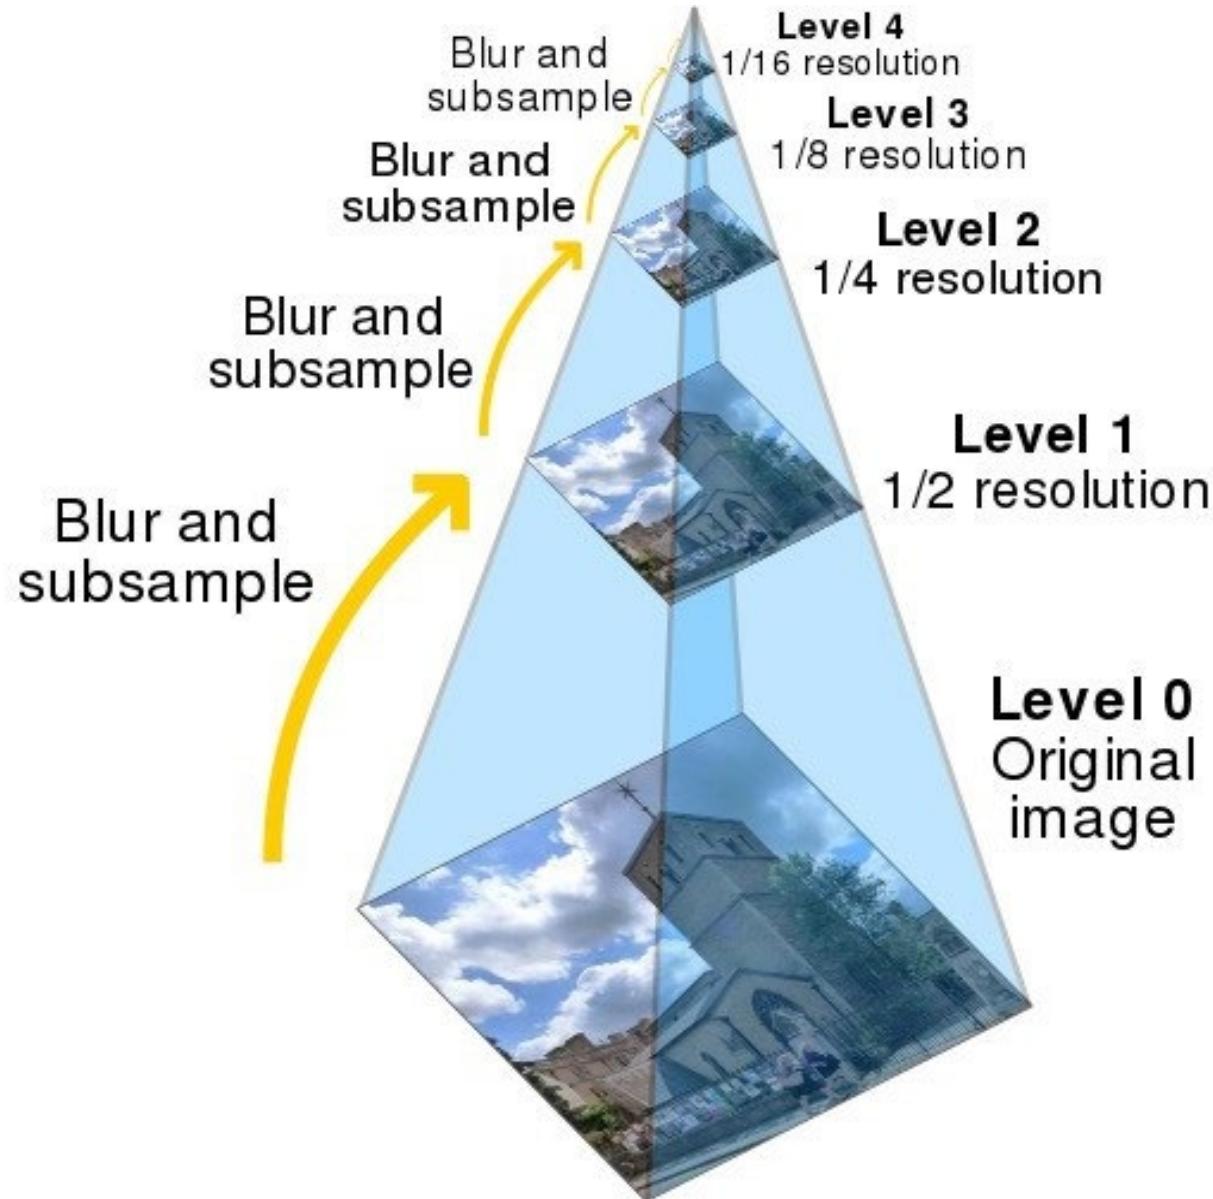


*Gaussian
pyramid*



blur subsample blur subsample ...





Similar to Gaussian Pyramids, there are Laplacian Pyramids.

- We will see them in the following class.

What are they good for?

- Improve Search: search over scale
 - Template matching
 - Classic coarse-to-fine strategy
 - ✓ E.g. find an object at different scales



Upsampling

- This image is too small for this screen:
- How can we make it 10 times as big?
- Simplest approach: repeat each row and column 10 times (“nearest neighbor interpolation”)

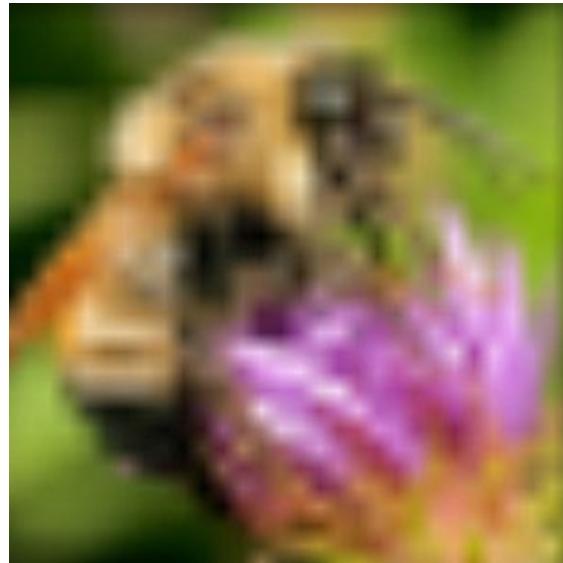


Image interpolation

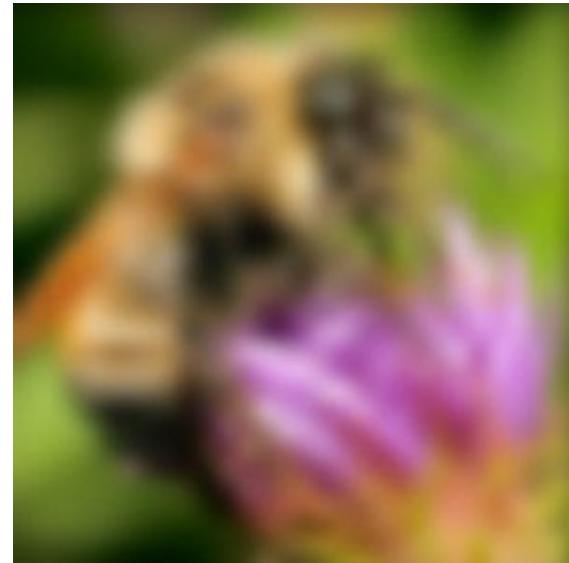
Original image:  x 10



Nearest-neighbor
interpolation

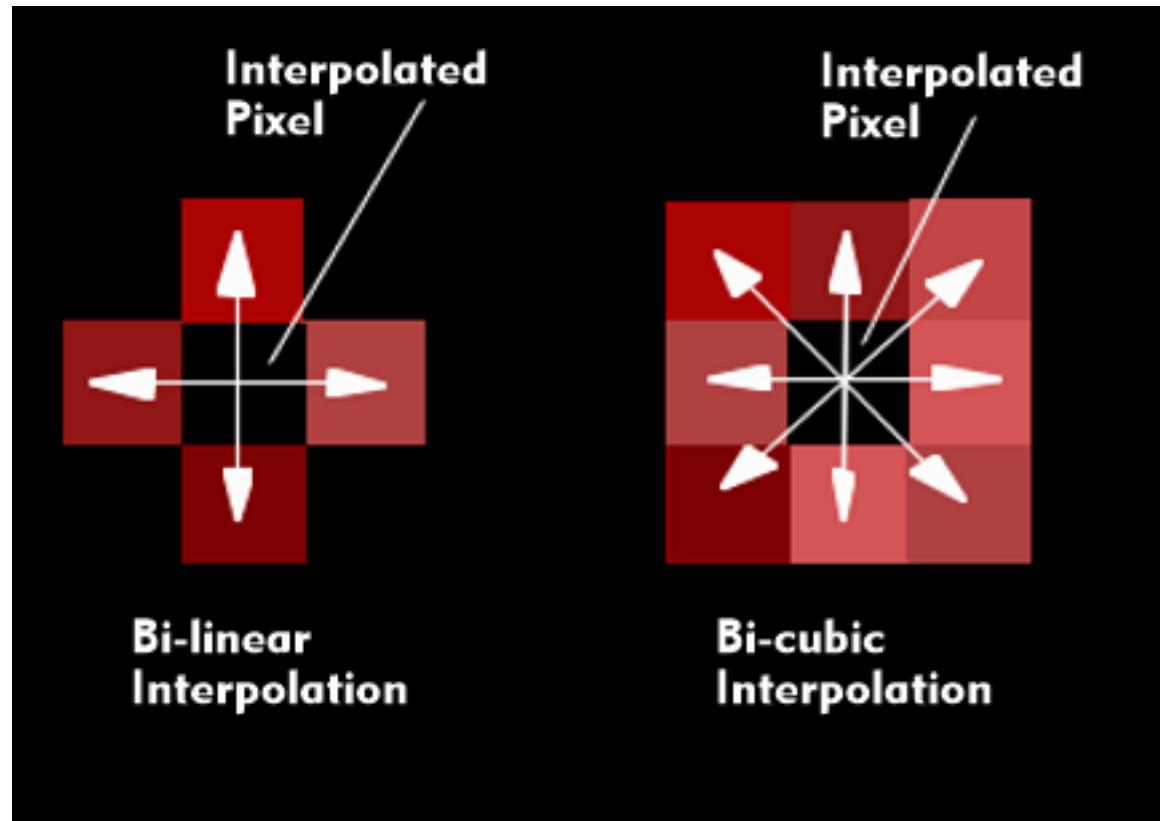


Bilinear interpolation

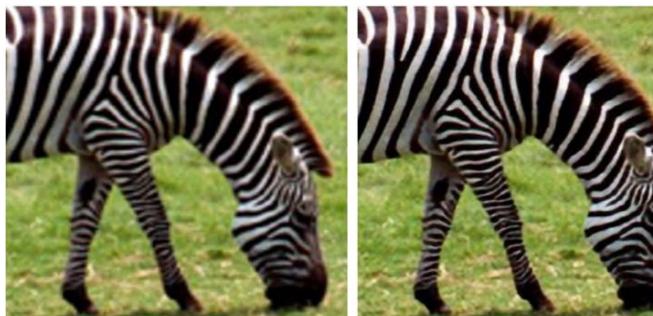


Bicubic interpolation

Bilinear and Bicubic interpolation



Modern methods



(a) Bicubic



(b) SRCNN



(c) A+



(d) RAISR



(e) Bicubic



(f) SRCNN



(g) A+



(h) RAISR

Romano, et al: RAISR: Rapid and Accurate Image Super Resolution 2016, IEEE Transaction on Computational Image

R. Timofte, V. De Smet, and L. Van Gool. A+: Adjusted Anchored Neighborhood Regression for Fast Super-Resolution. In Asian Conference on Computer Vision (ACCV 2014), 1-5 November

2014

Super-resolution with multiple images

- Can do better upsampling if you have multiple images of the scene taken with small (subpixel) shifts
- Some cellphone cameras (like the Google Pixel line) capture a **burst** of photos
- Can we use that burst for upsampling?

Google Pixel 3 Super Res Zoom



Super-resolution Using Hand Motion

Effect of hand tremor as seen in a cropped burst of photos, after global alignment
<https://ai.googleblog.com/2018/10/see-better-and-further-with-super-res.html>



Example photo with and without super res zoom (smart burst align and merge)

Image Transformations & Filtering

- Point Processing
- Linear Filtering
- Sampling & Aliasing
- Image Derivatives
- Edge Detection

Partial derivatives with convolution

Image is a function $f(x, y)$

Remember:

Approximate

:

-1	1
----	---

Another one:

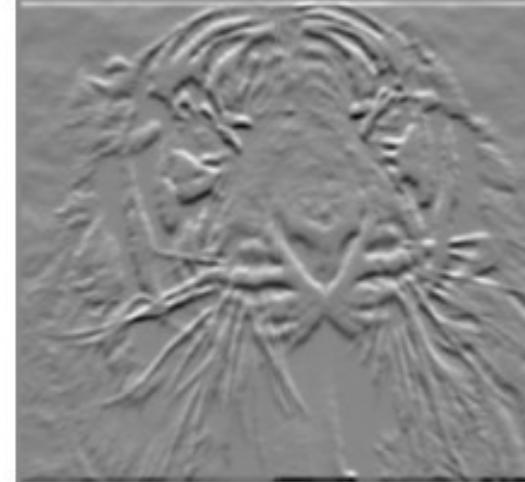
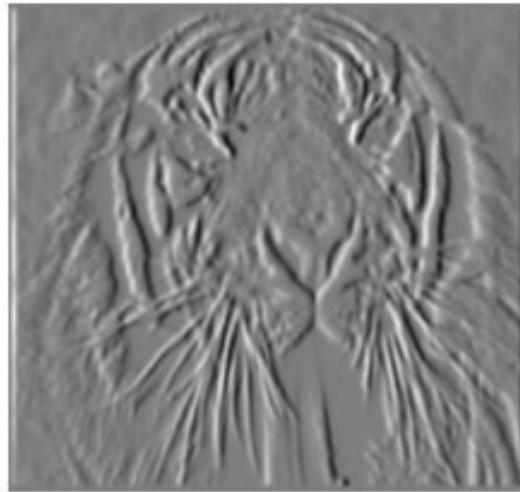
-1	0	1
----	---	---

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\epsilon \rightarrow 0} \frac{f(x + \epsilon, y) - f(x, y)}{\epsilon}$$

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x - 1, y)}{2}$$

Image Gradient



$$\begin{array}{|c|c|} \hline -1 & \quad \\ \hline \end{array}$$

$$\frac{\partial f(x, y)}{\partial x}$$

$$\frac{\partial f(x, y)}{\partial}$$

$$\begin{array}{|c|c|} \hline -1 & \quad \\ \hline \quad & 1 \\ \hline \end{array}$$

$$\begin{array}{|c|c|} \hline 1 & \quad \\ \hline \quad & -1 \\ \hline \end{array}$$

or

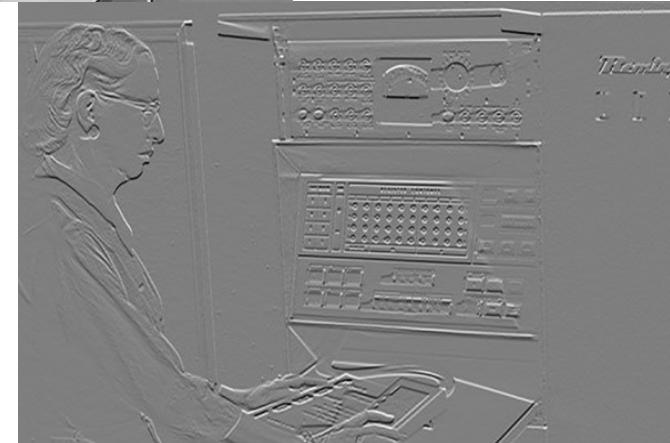
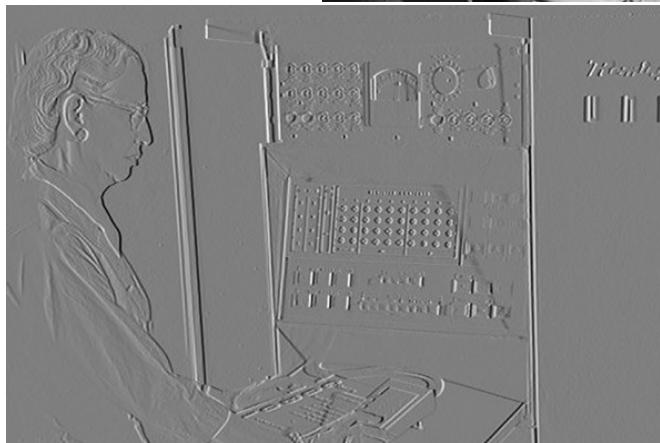
$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$



Partial Derivatives



$$\frac{\partial f(x, y)}{\partial x}$$



$$\frac{\partial f(x, y)}{\partial y}$$

Gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$



Gradient Orientation

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

all the gradients

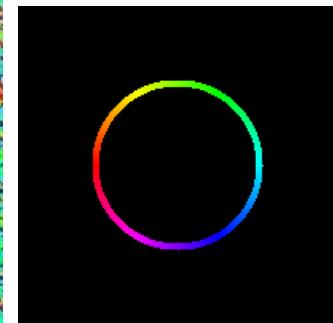
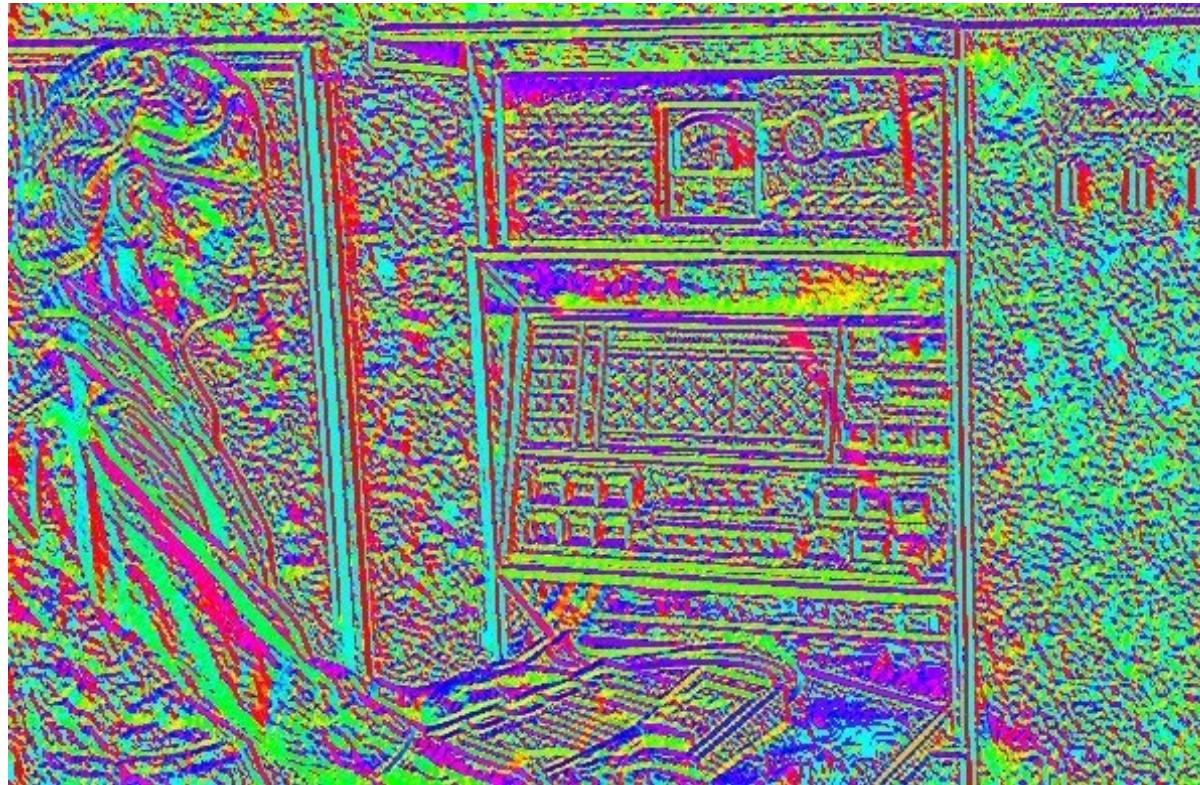
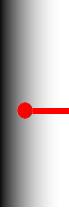
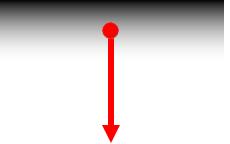


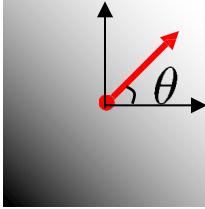
Image gradient

- The gradient of an image:

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

-  $\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$


$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$


$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient points in the direction of most rapid increase in intensity

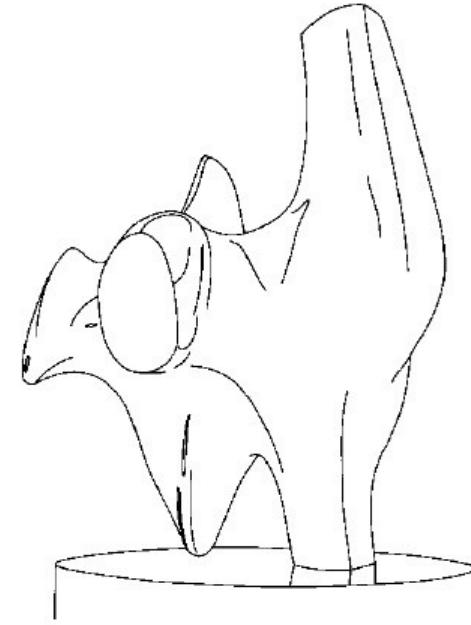
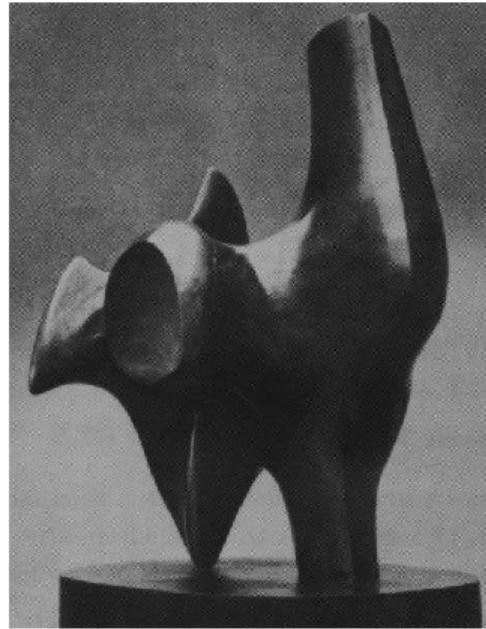
- How does this direction relate to the direction of the edge?

The *edge strength* is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

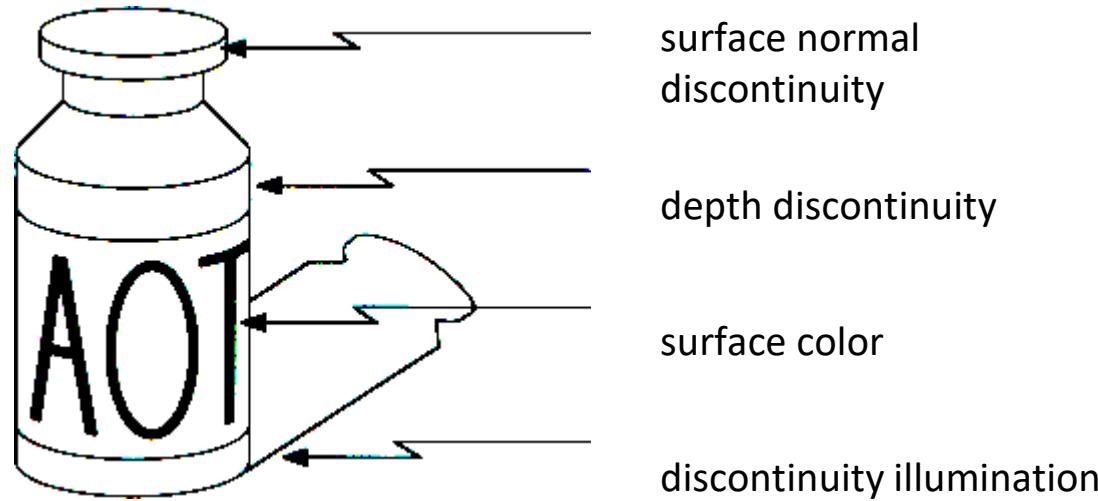
The gradient direction is given by $\theta = \tan^{-1} \left(\frac{\partial f / \partial y}{\partial f / \partial x} \right)$

Edge detection



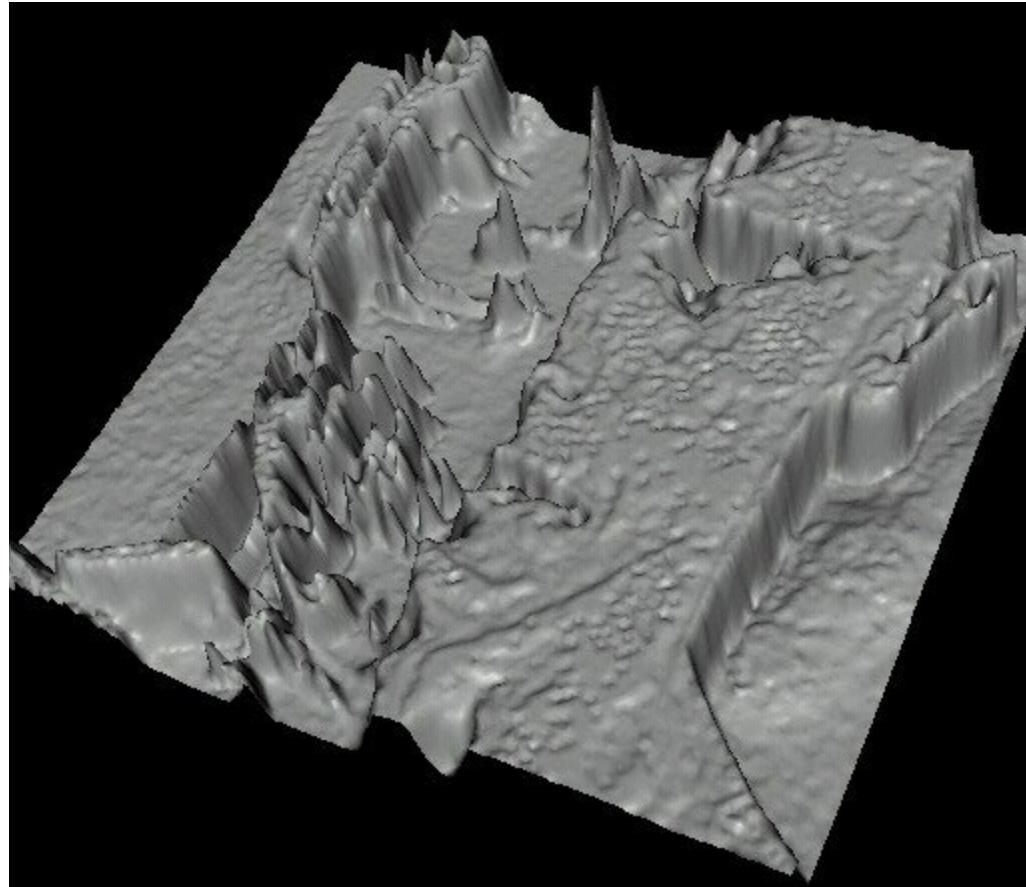
- Convert a 2D image into a set of curves
 - Extracts salient features of the scene
 - More compact than pixels

Origin of edges



- Edges are caused by a variety of discontinuity factors

Images as functions...



- Edges look like steep cliffs

Characterizing edges

- An edge is a place of *rapid change* in the image intensity function

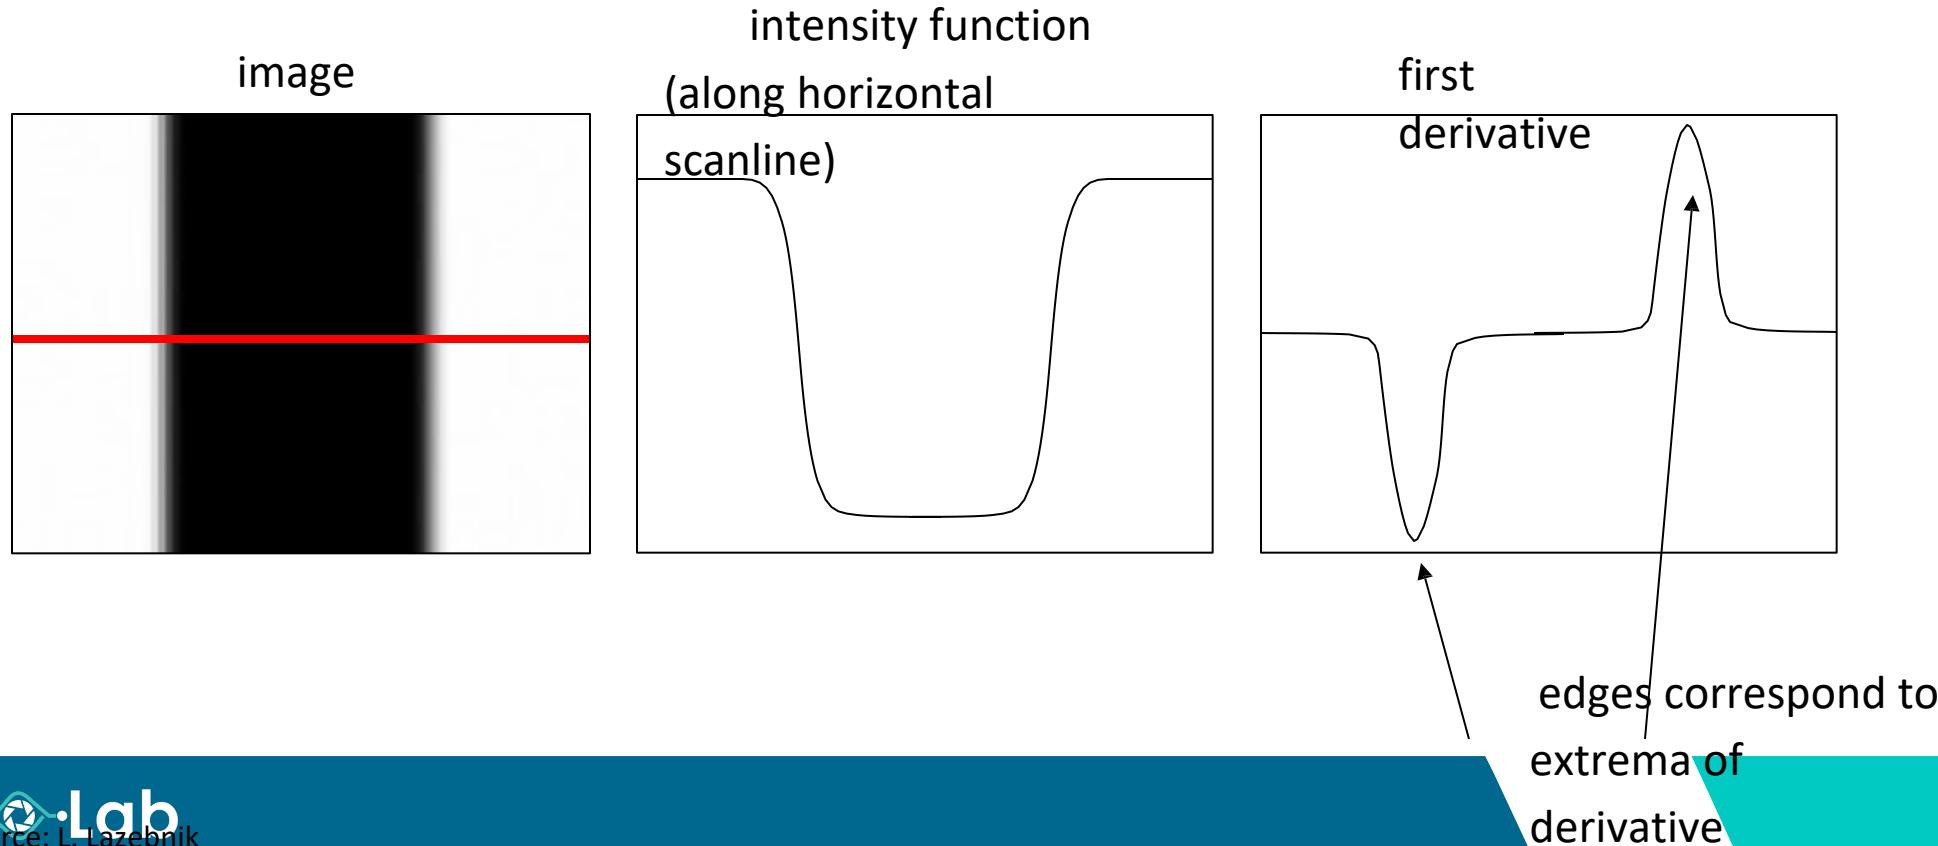
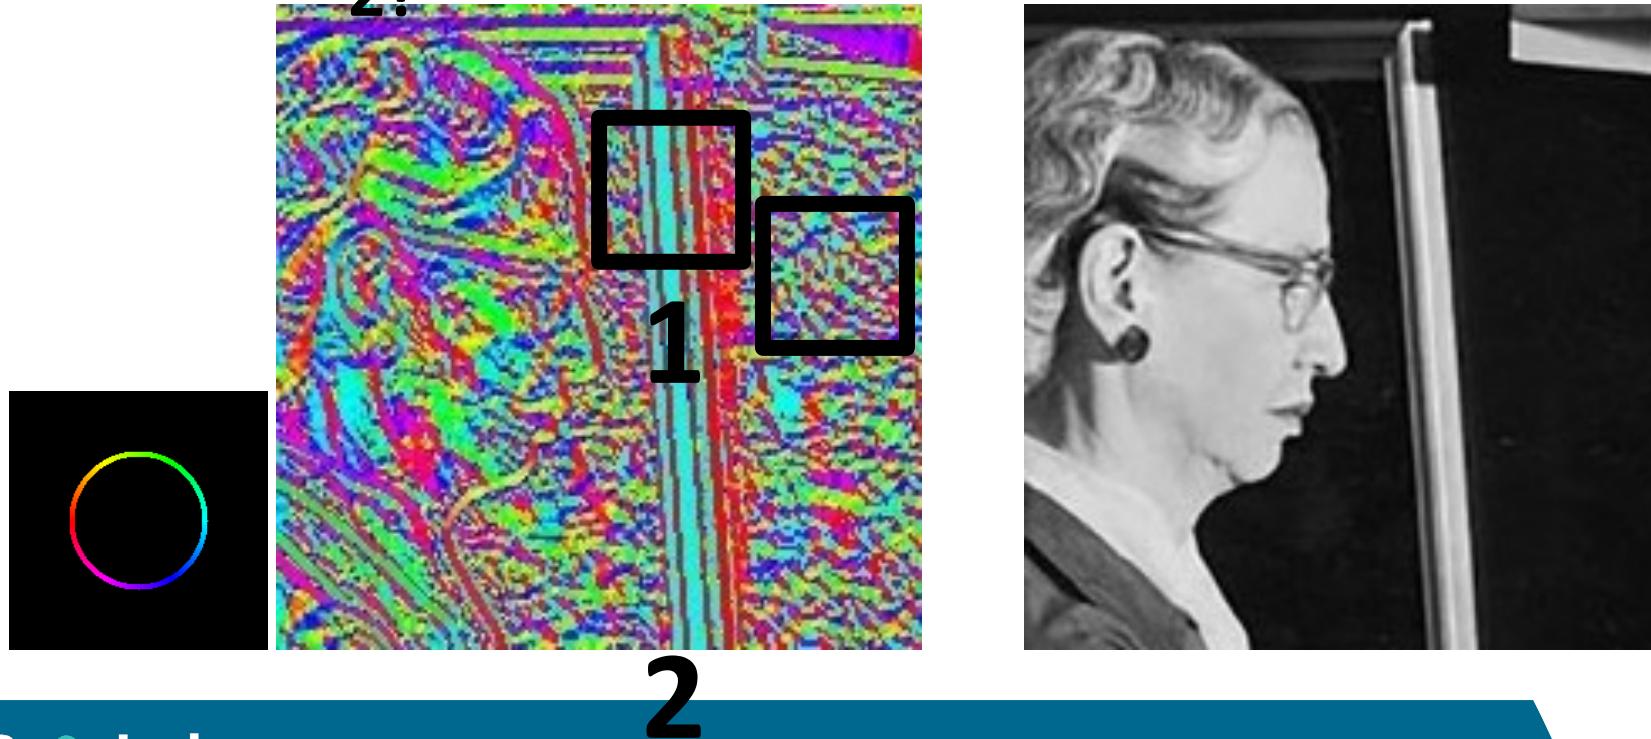


Image Gradient & Edges

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

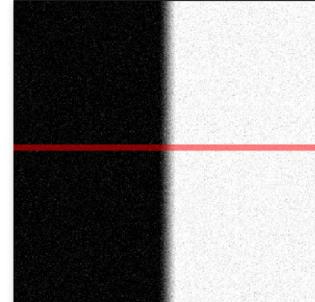
Direction of image gradients

Why is there structure at 1 and not at
2?

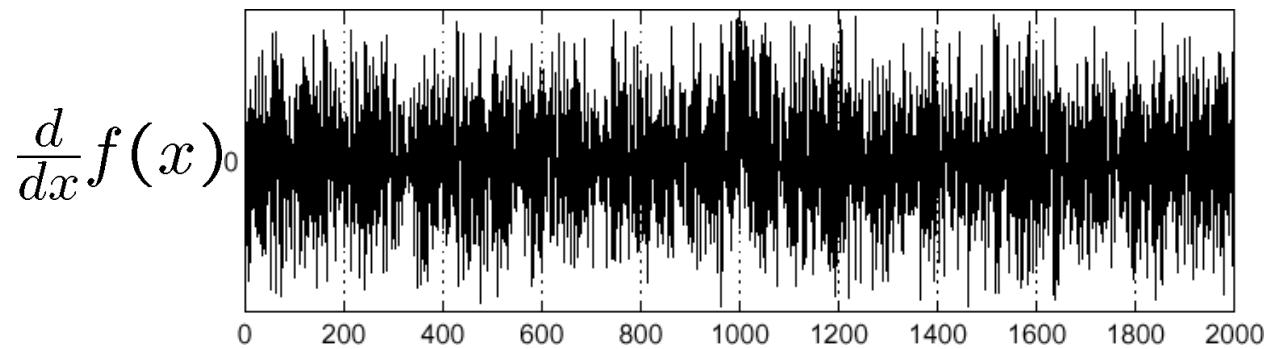
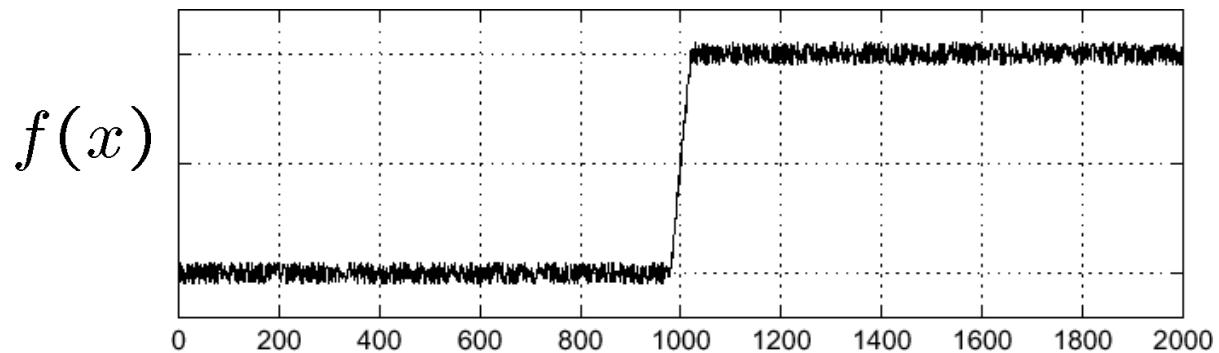


Effects of noise

- Consider a single row or column of the image
 - Plotting intensity as a function of position gives a signal



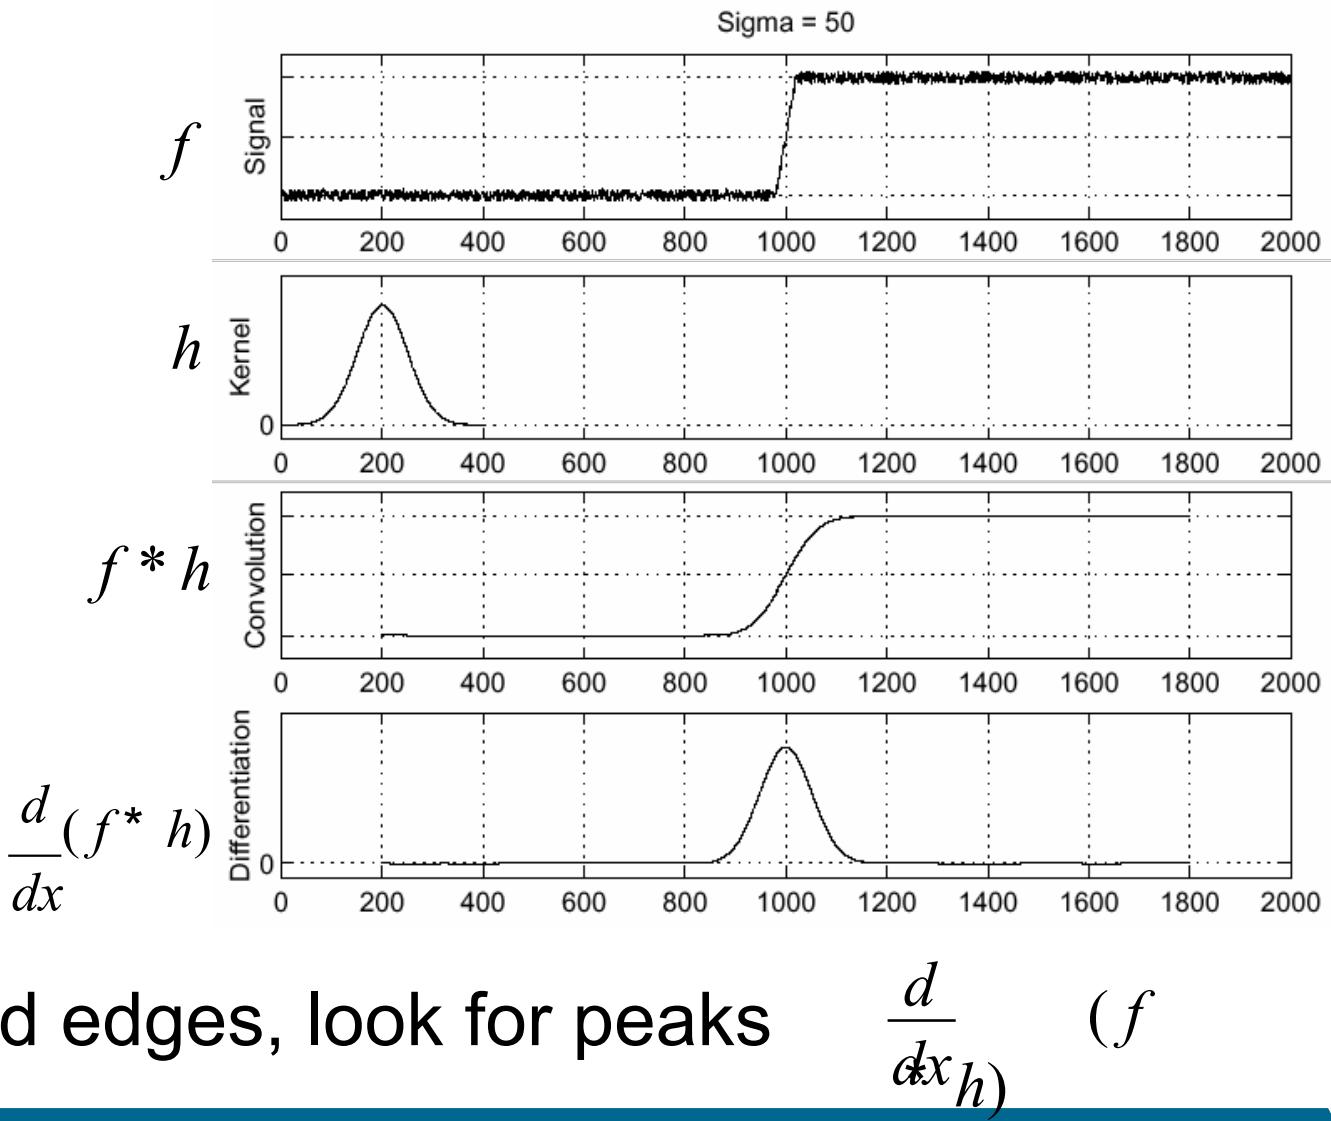
Noisy input image



Where is the edge?

- Issues with derivate filters:
- Sensitive to edges at small spatial scales
 - Also sensitive to noise

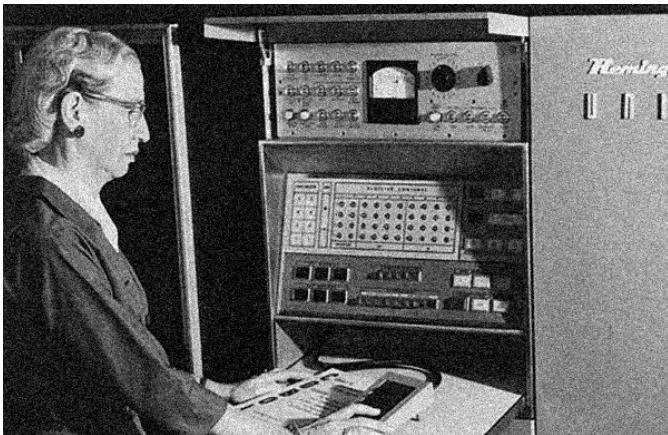
Solution: smooth first



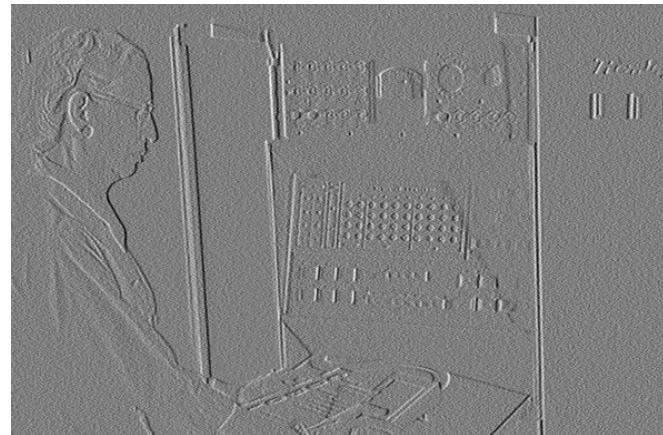
- To find edges, look for peaks in $\frac{d}{dx}(f * h)$

Noise in 2D

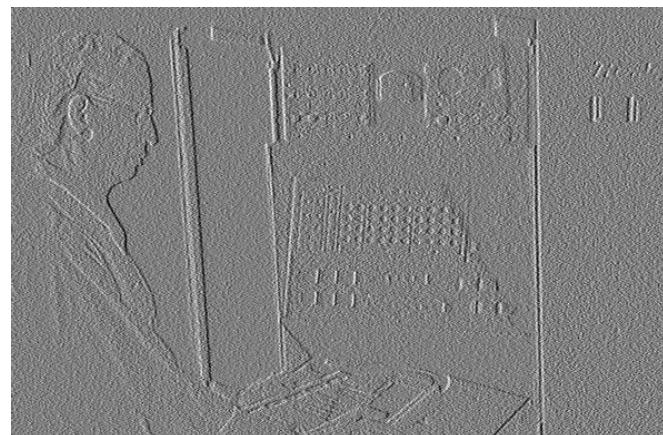
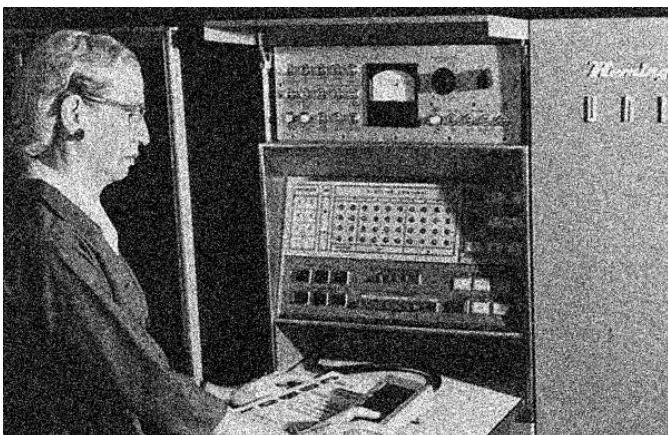
Noisy Input



I_x via $[-1,0,1]$



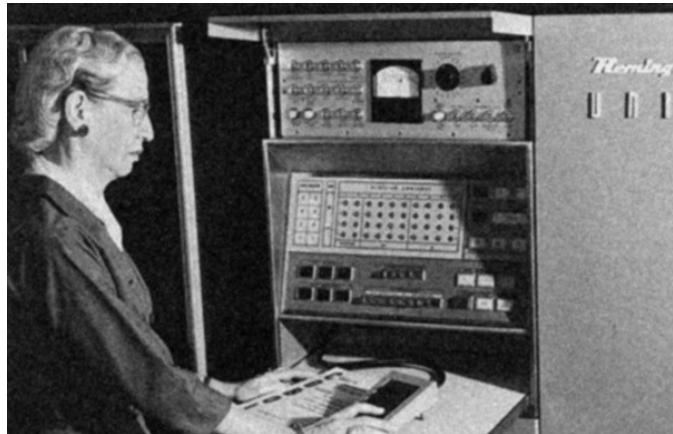
Zoom



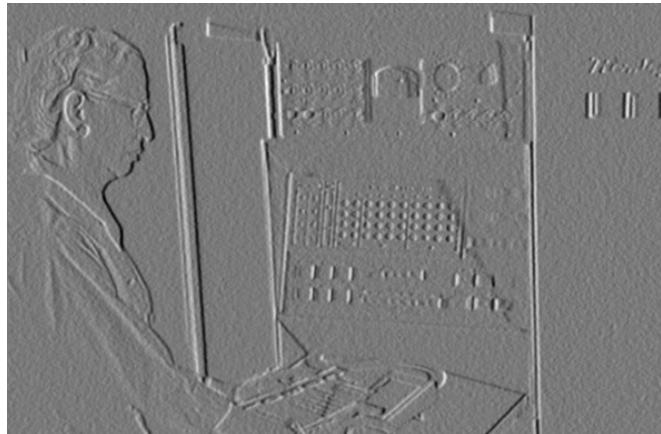
Source: D. Fouhey

Noise + Smoothing

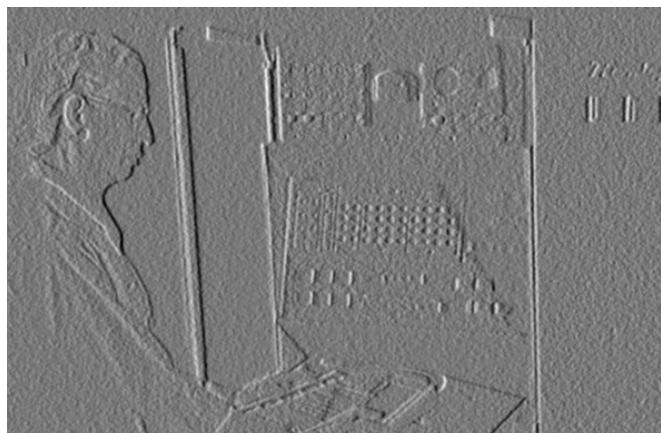
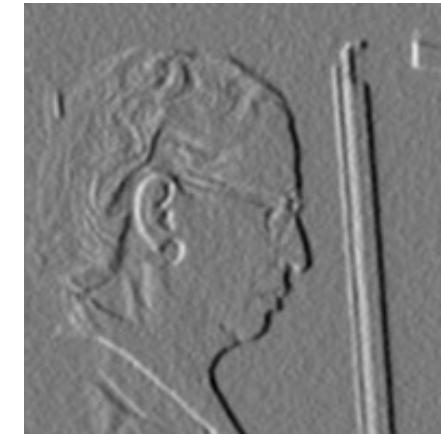
Smoothed Input



$Ix \text{ via } [-1,01]$

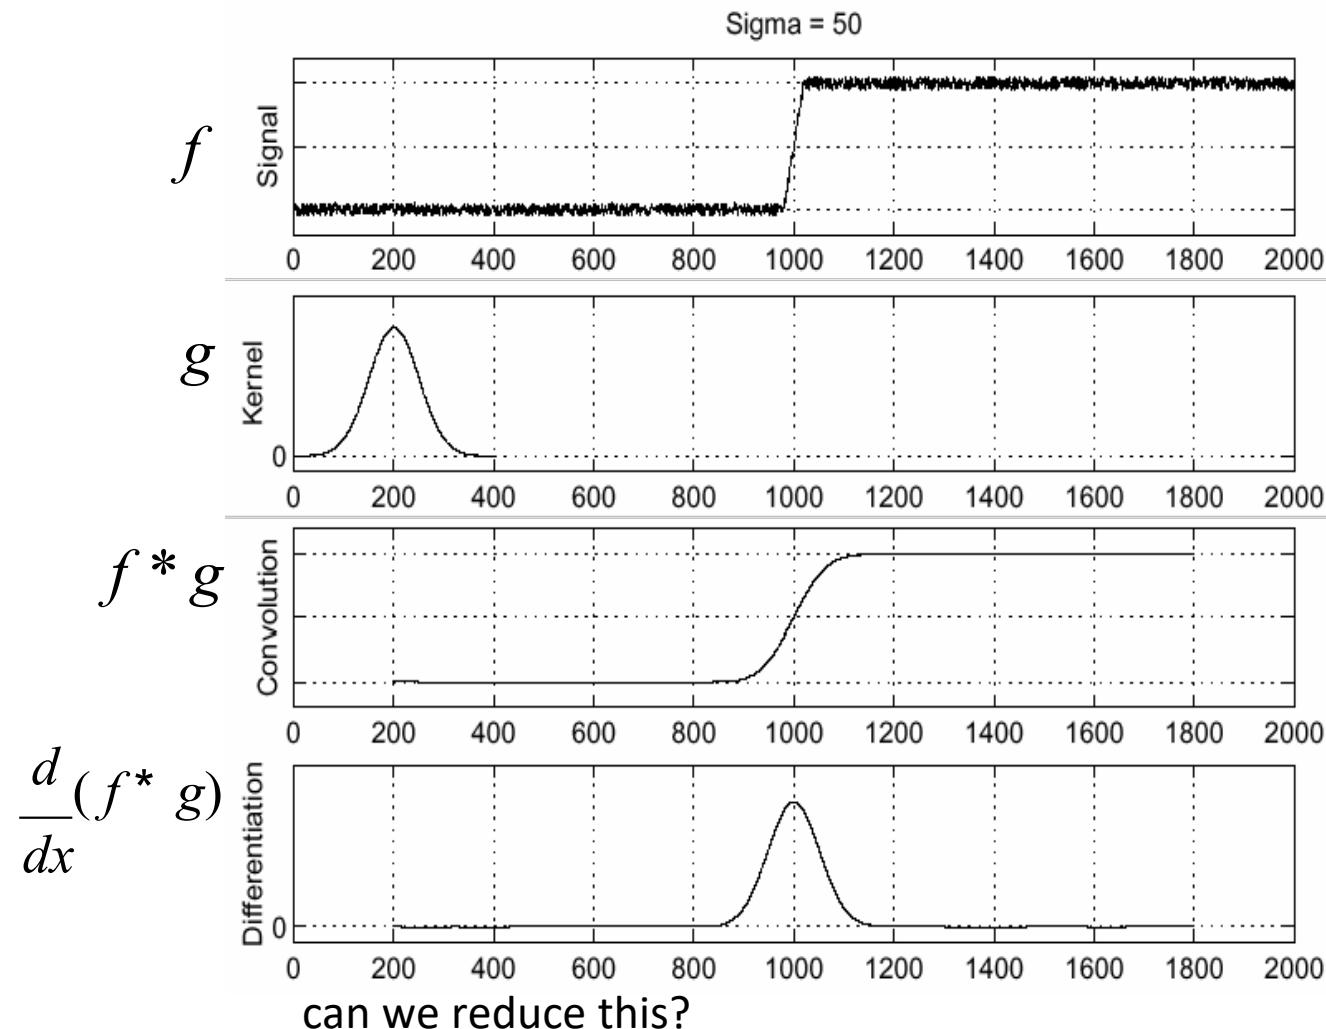


Zoom



Source: D. Fouhey

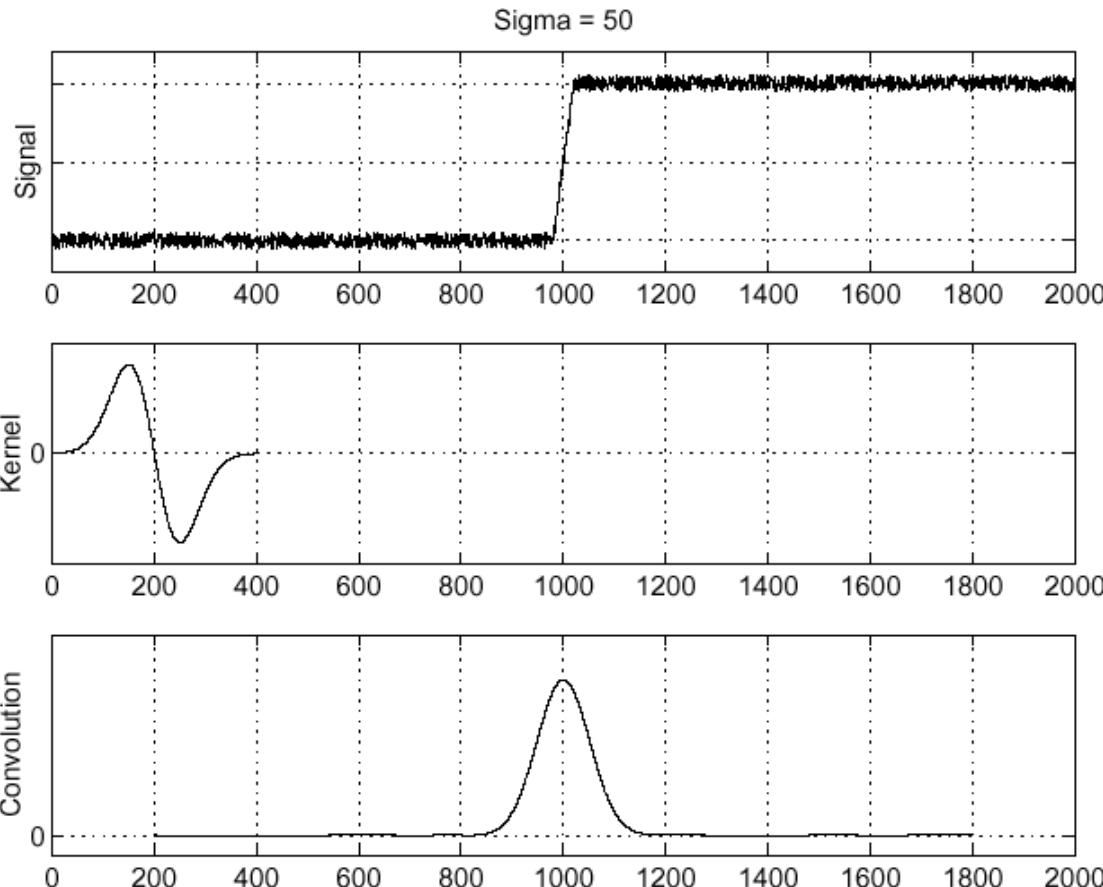
How many steps here?



Derivative theorem of convolution

$$\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$$

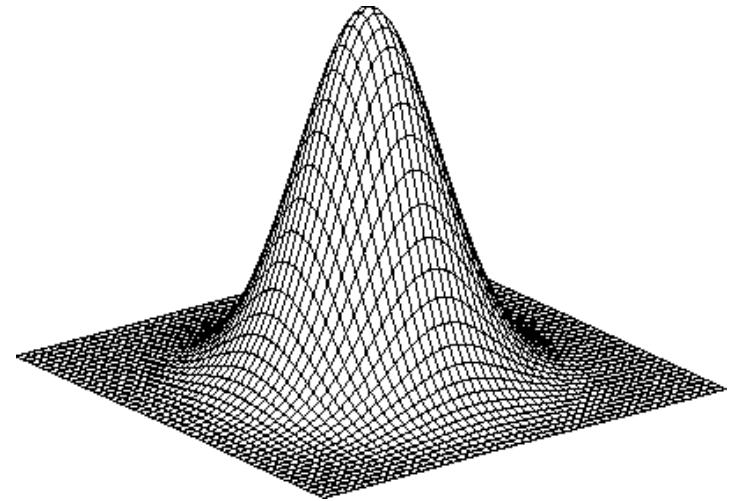
- This saves us one operation:



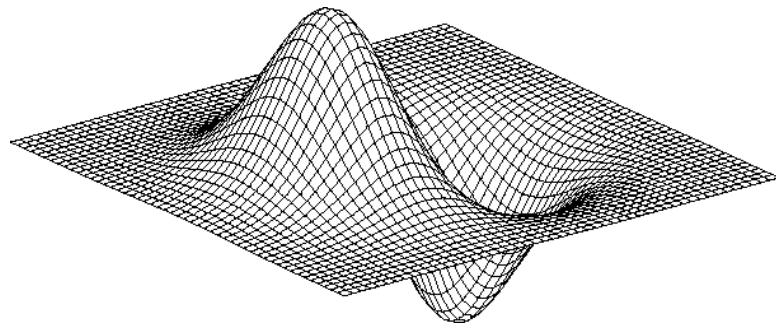
$$\frac{\partial}{\partial x}h$$

$$(\frac{\partial}{\partial x}h) \star f$$

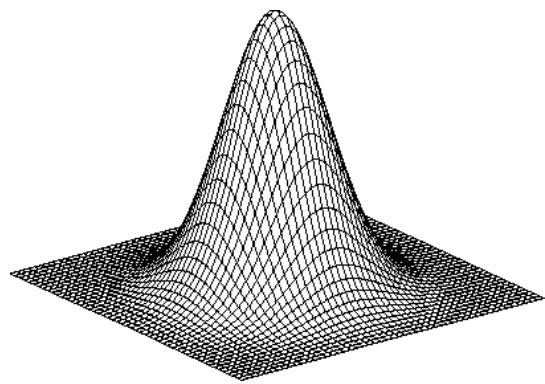
Derivative of Gaussian filter



* [1 -1]=

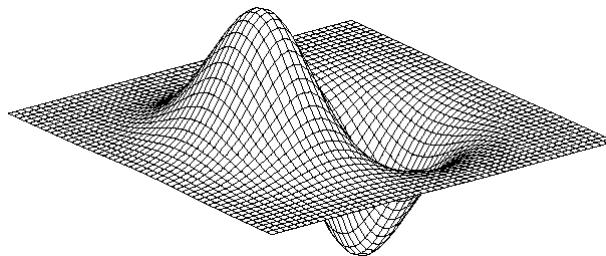


2D edge detection filters



Gaussian

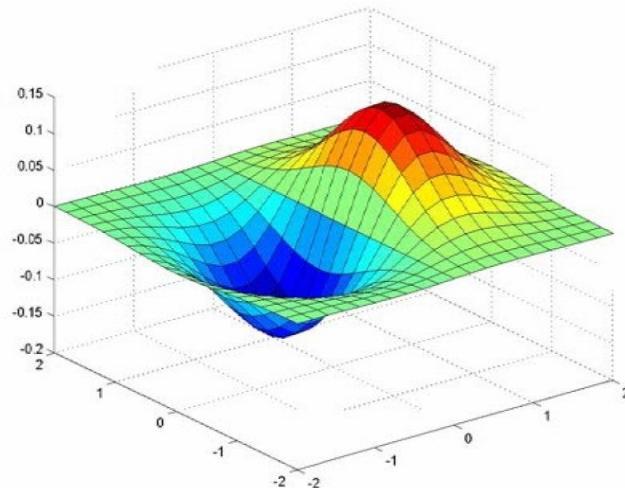
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



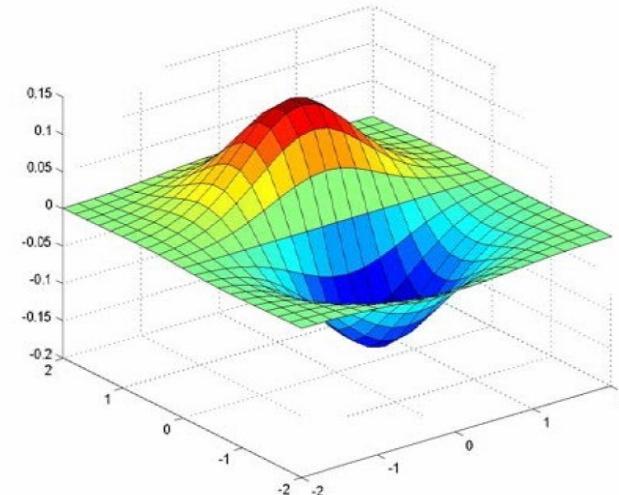
Derivative of Gaussian

$$(x) \quad \frac{\partial}{\partial x} h_\sigma(u, v)$$

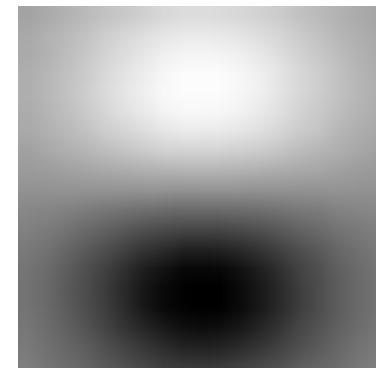
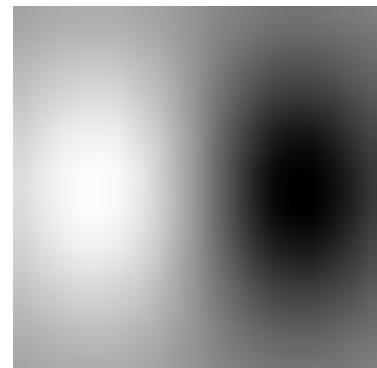
Derivative of Gaussian filter



x-direction



y-direction



Laplace filter

Basically a second derivative filter.

- We can use finite differences to derive it, as with first derivative filter.

first-order
finite difference $f'(x) = \lim_{h \rightarrow 0} \frac{f(x + 0.5h) - f(x - 0.5h)}{h}$ \longrightarrow 1D derivative filter

1	0	-1
---	---	----

second-order
finite difference $f''(x) = \lim_{h \rightarrow 0} \frac{f(x + h) - 2f(x) + f(x - h)}{h^2}$ \longrightarrow Laplace filter
?

Laplace filter

Basically a second derivative filter.

- We can use finite differences to derive it, as with first derivative filter.

first-order
finite difference $f'(x) = \lim_{h \rightarrow 0} \frac{f(x + 0.5h) - f(x - 0.5h)}{h}$ \longrightarrow 1D derivative filter

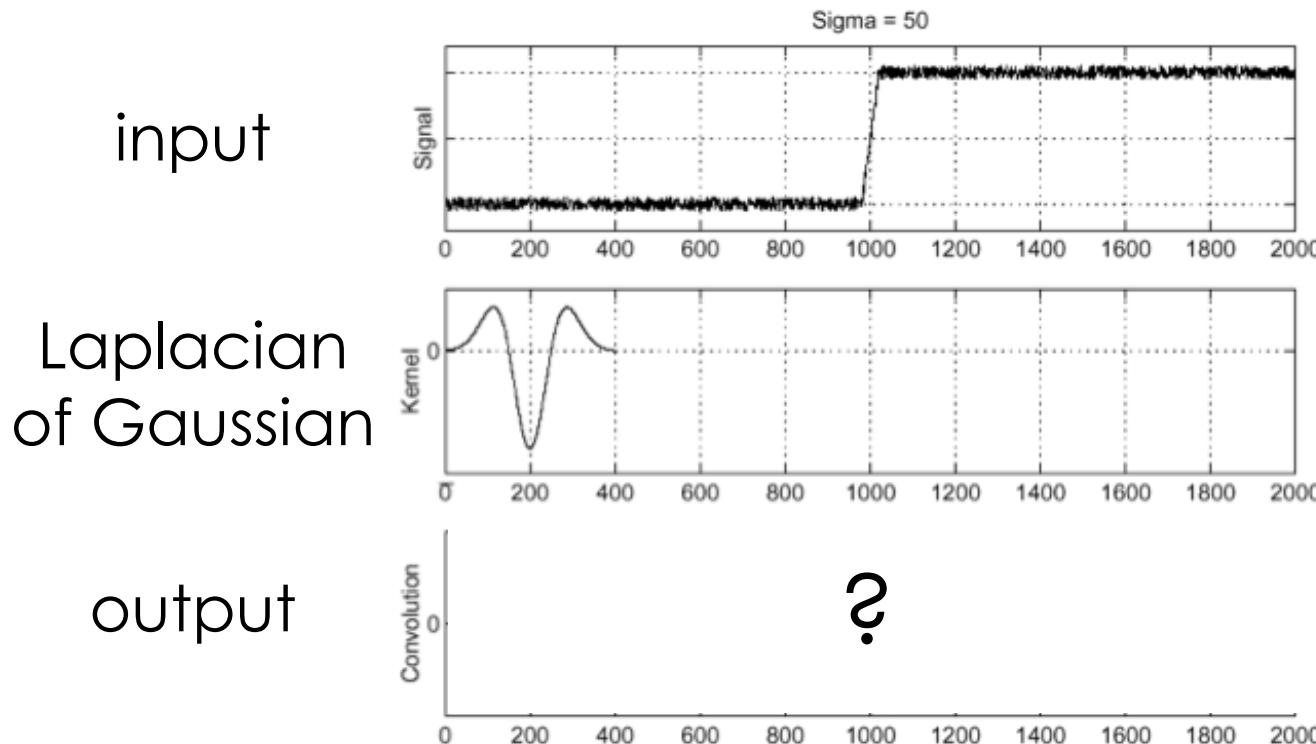
1	0	-1
---	---	----

second-order
finite difference $f''(x) = \lim_{h \rightarrow 0} \frac{f(x + h) - 2f(x) + f(x - h)}{h^2}$ \longrightarrow Laplace filter

1	-2	1
---	----	---

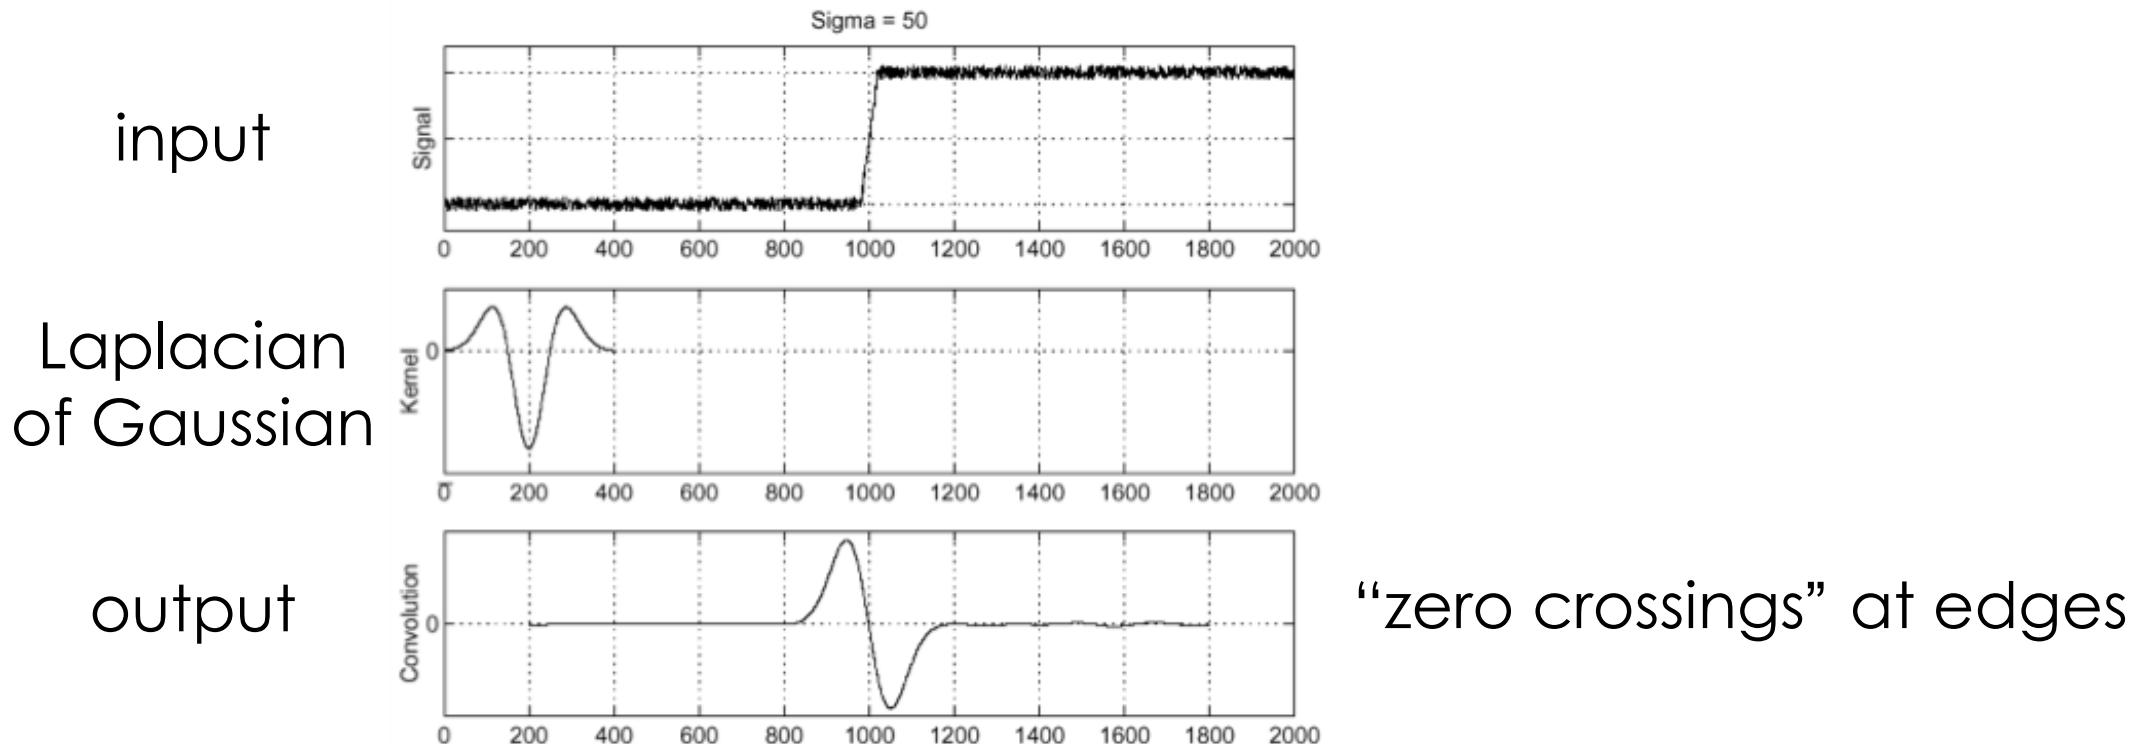
Laplacian of Gaussian (LoG) filter

As with derivative, we can combine Laplace filtering with Gaussian filtering

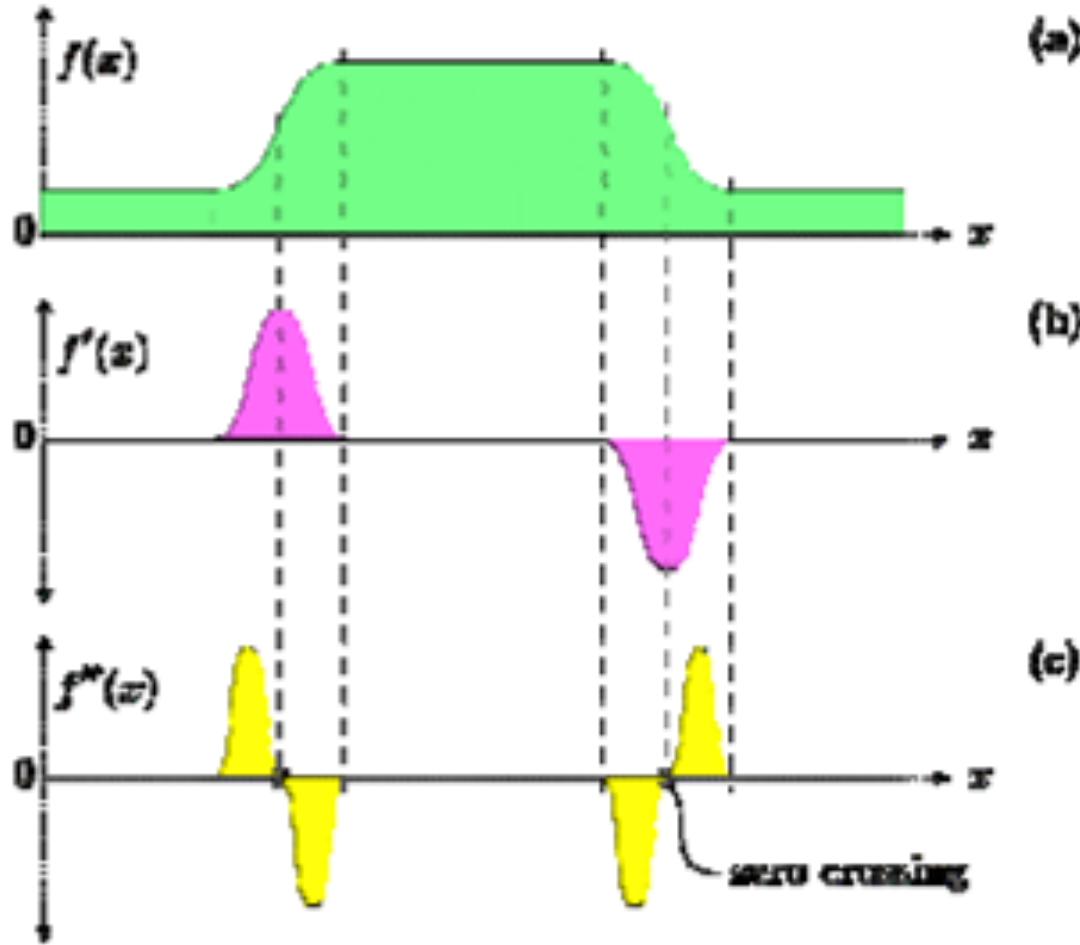


Laplacian of Gaussian (LoG) filter

As with derivative, we can combine Laplace filtering with Gaussian filtering



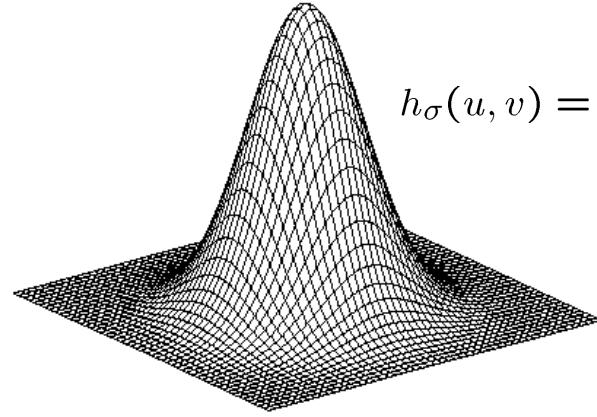
Zero-crossing



First derivative at edge is a maximum
Second derivative at edge is zero

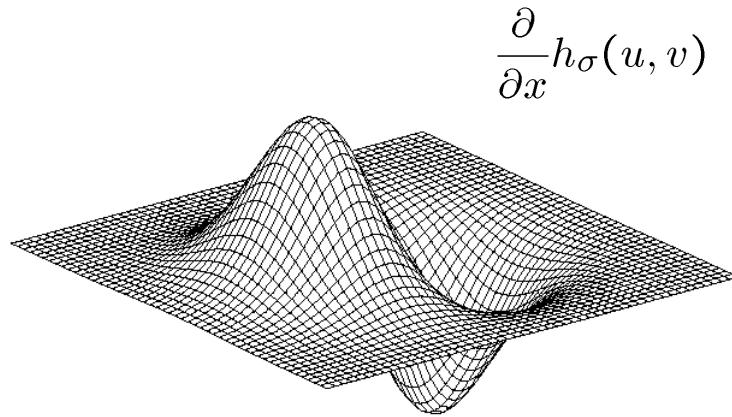
It is more easy to understand when a function is zero than find a maximum

2D Gaussian filters



Gaussian

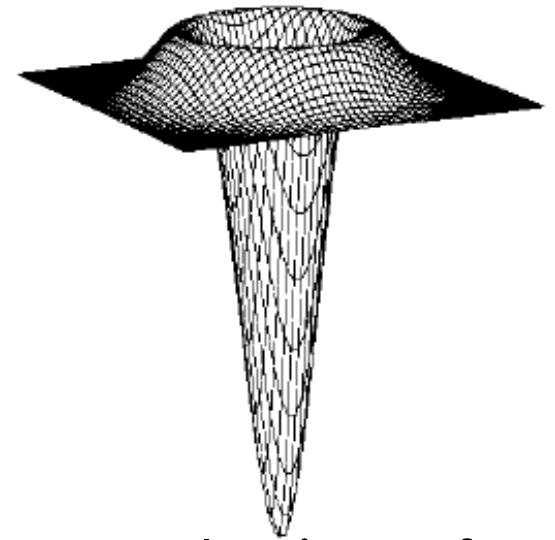
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



Derivative of Gaussian

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

$$\nabla^2 h_\sigma(u, v)$$



Laplacian of
Gaussian

Laplace and LoG filtering examples



Laplacian of Gaussian filtering



Laplace filtering

Laplacian of Gaussian vs Derivative of Gaussian

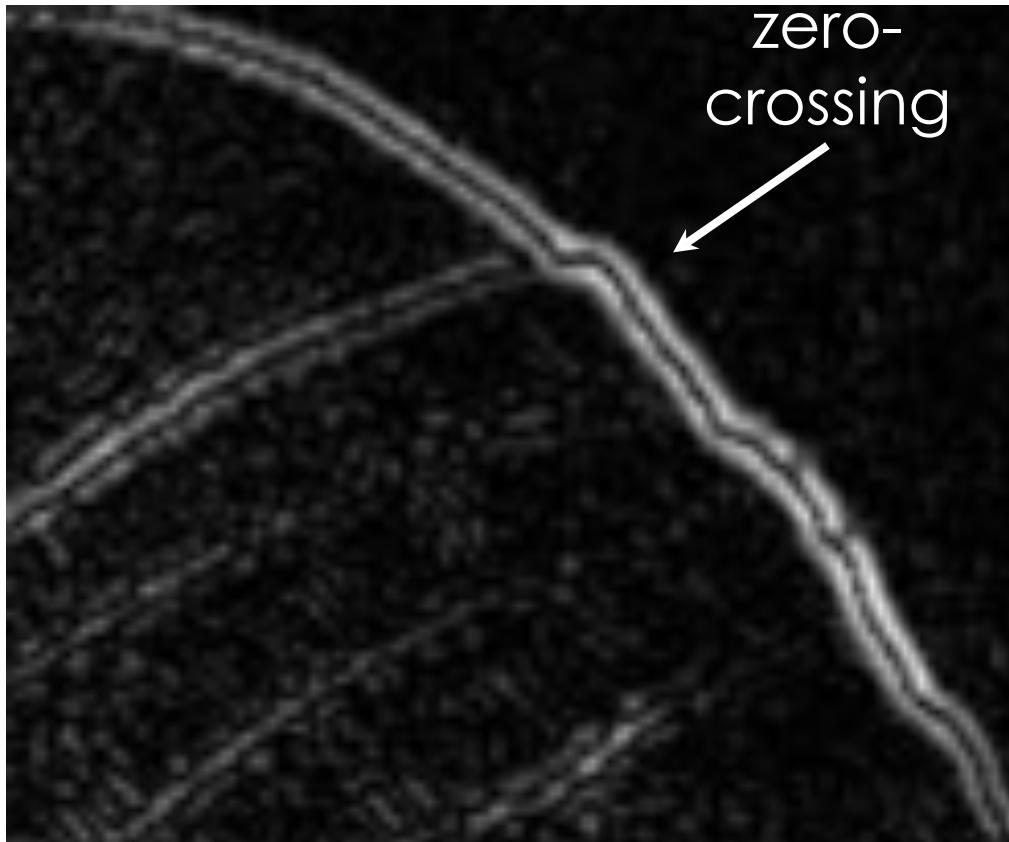


Laplacian of Gaussian filtering

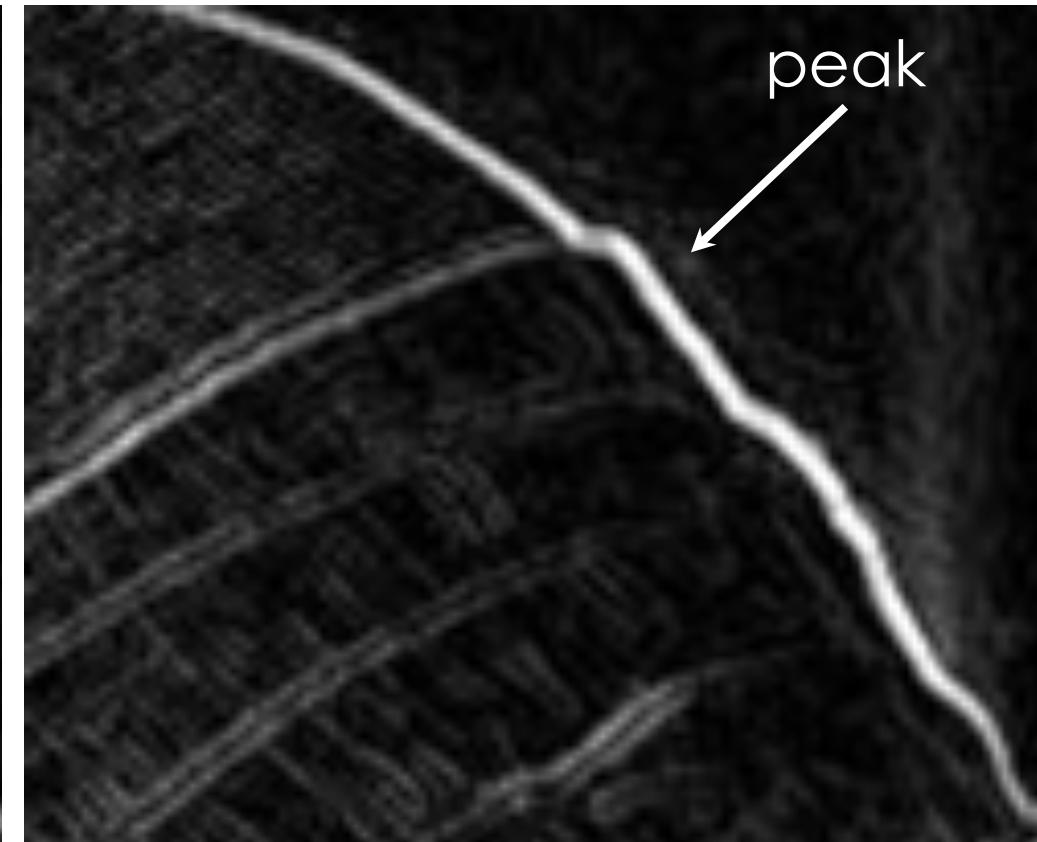


Derivative of Gaussian filtering

Laplacian of Gaussian vs Derivative of Gaussian



Laplacian of Gaussian filtering



Derivative of Gaussian filtering

Zero crossings are more accurate at localizing edges but DoG is computationally more efficient

The Sobel operator

- Common approximation of Derivative of Gaussian

$$\frac{1}{8} \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

s_x

$$\frac{1}{8} \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

s_y

- The standard definition of the Sobel operator omits the $1/8$ term
 - doesn't make a difference for edge detection
 - the $1/8$ term **is** needed to get the right gradient magnitude

Sobel operator: example

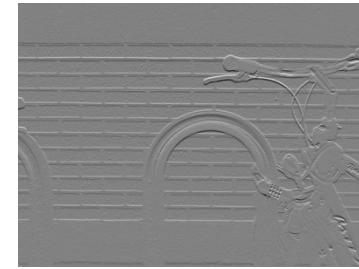
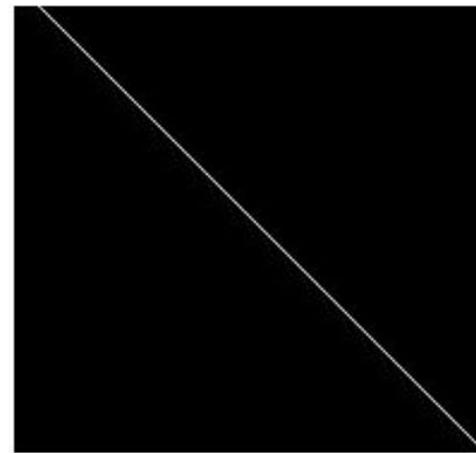




Image with Edge



Edge Location

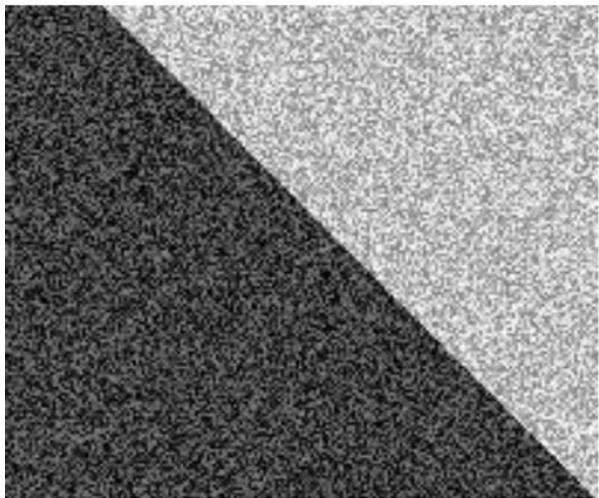
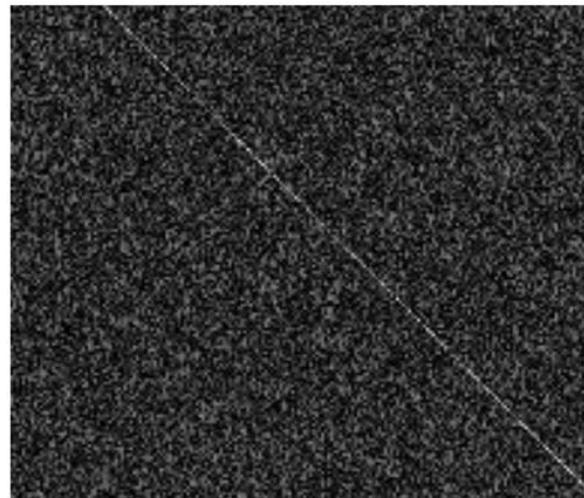
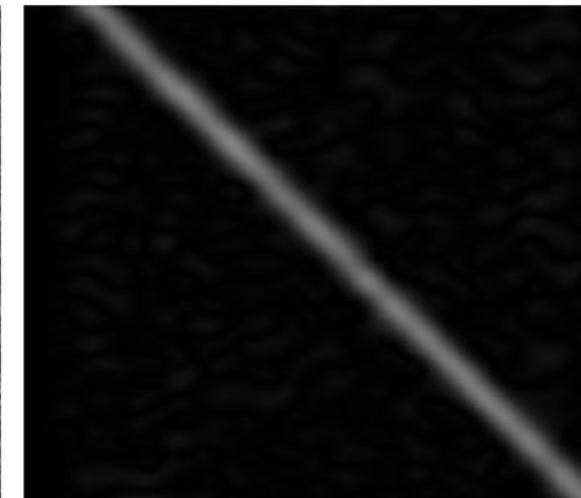


Image + Noise



Derivatives detect
edge *and* noise



Smoothed derivative removes
noise, but blurs edge

The Sobel filter

$$\begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline \end{array}$$

Sobel filter

What
filter is
this?

1D
derivative
filter

The Sobel filter

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 & -1 \end{bmatrix}$$

Sobel filter

Blurring

1D
derivative
filter

Does this filter return large responses on vertical or horizontal lines?

The Sobel filter

Horizontal Sober filter:

$$\begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline \end{array}$$

What does the vertical Sobel filter look like?

The Sobel filter

Horizontal Sober filter:

$$\begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline \end{array}$$

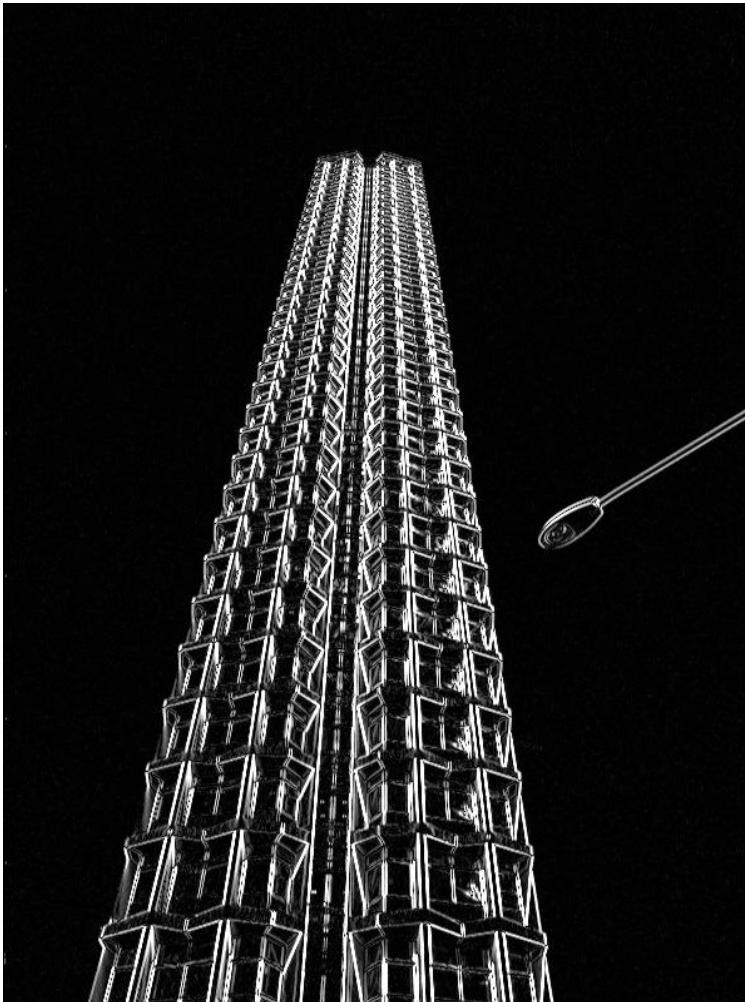
Vertical Sobel filter:

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 0 \\ \hline -1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array}$$

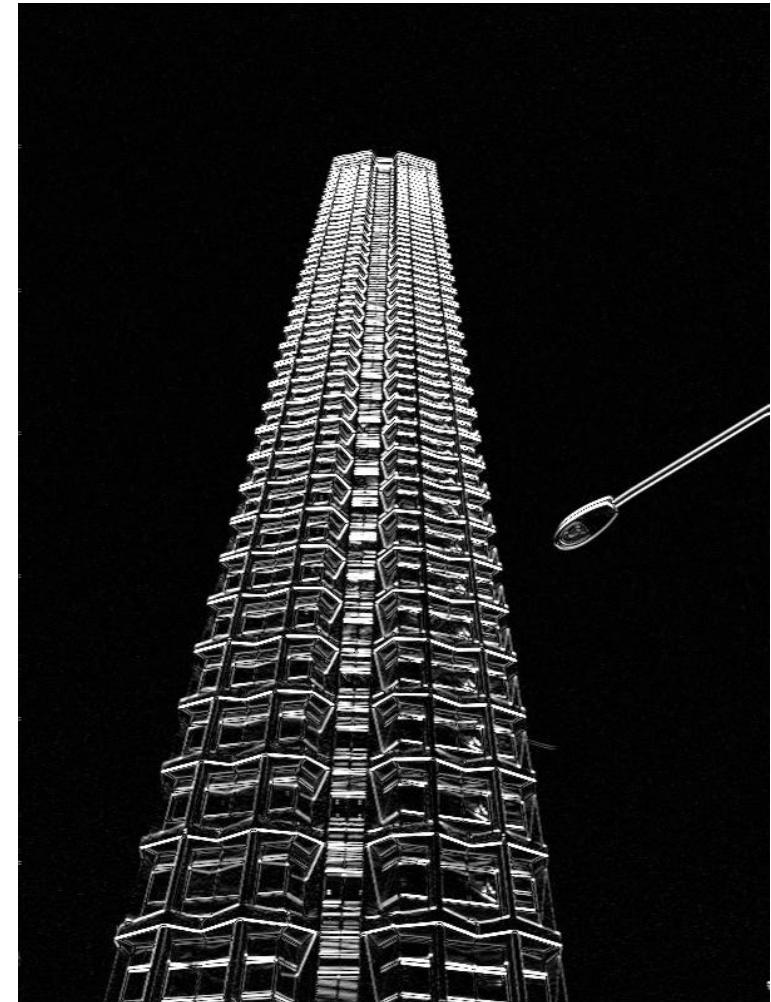
Sobel filter example



original



which Sobel filter?

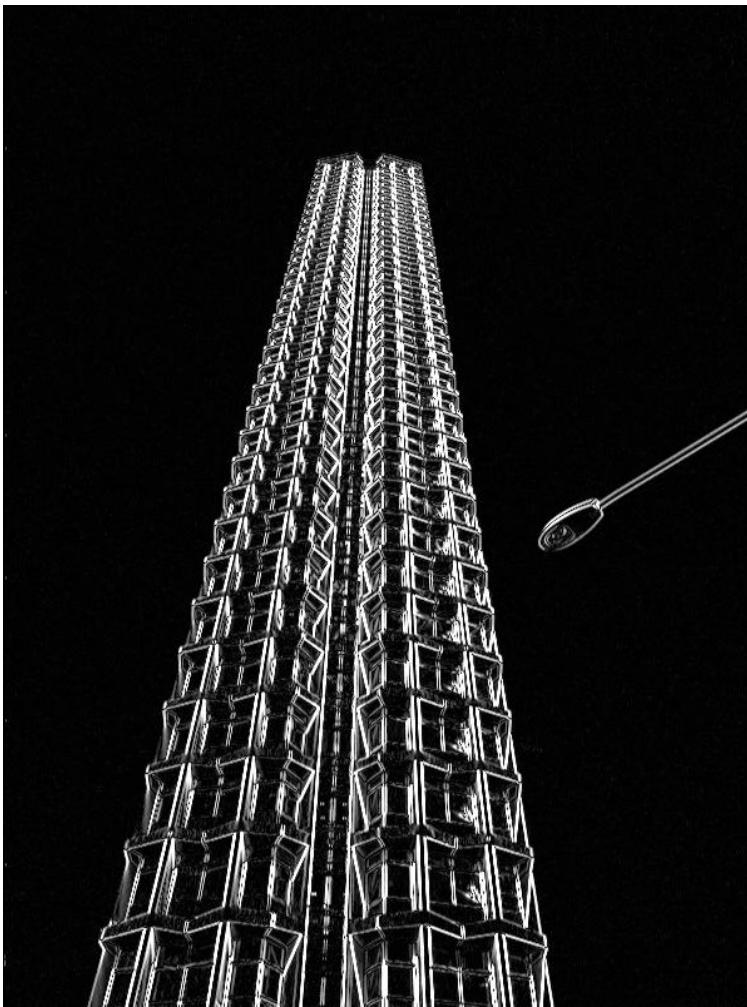


which Sobel filter?

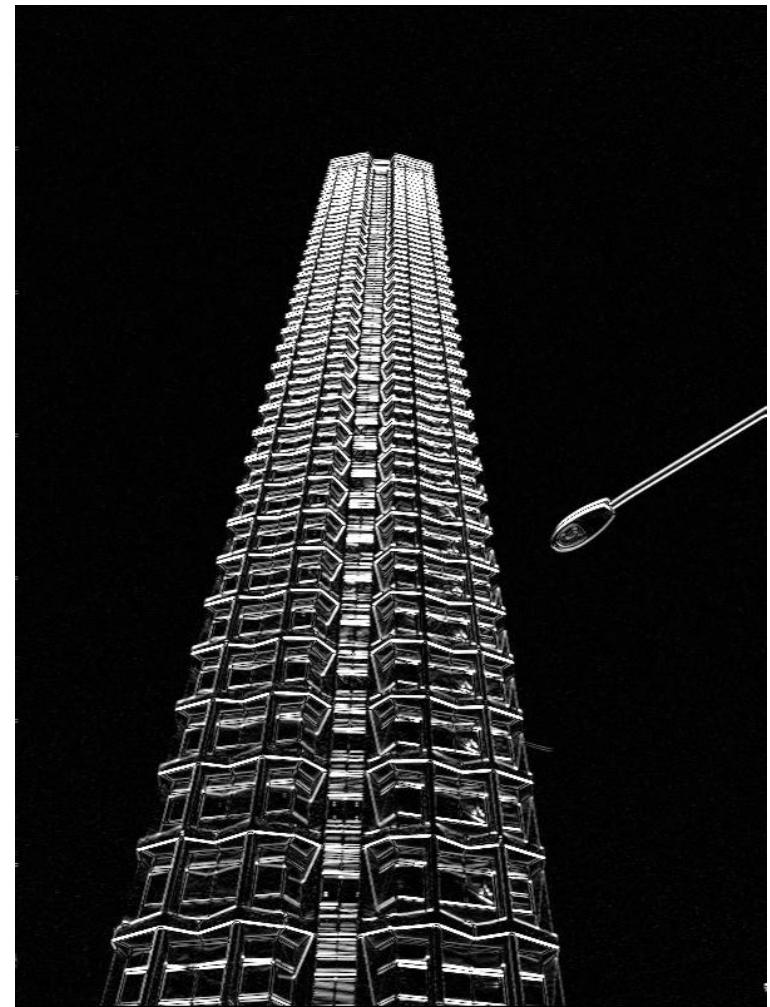
Sobel filter example



original

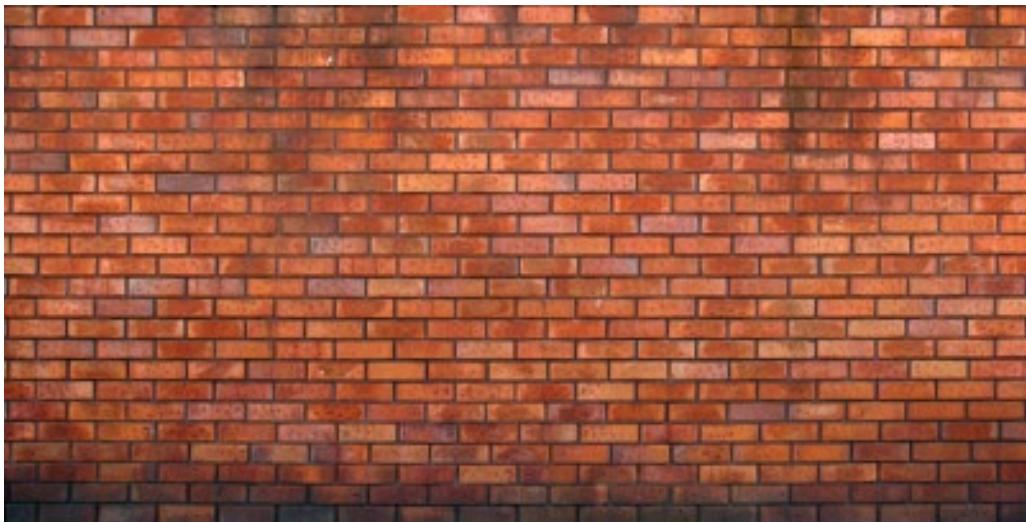


horizontal Sobel filter



vertical Sobel filter

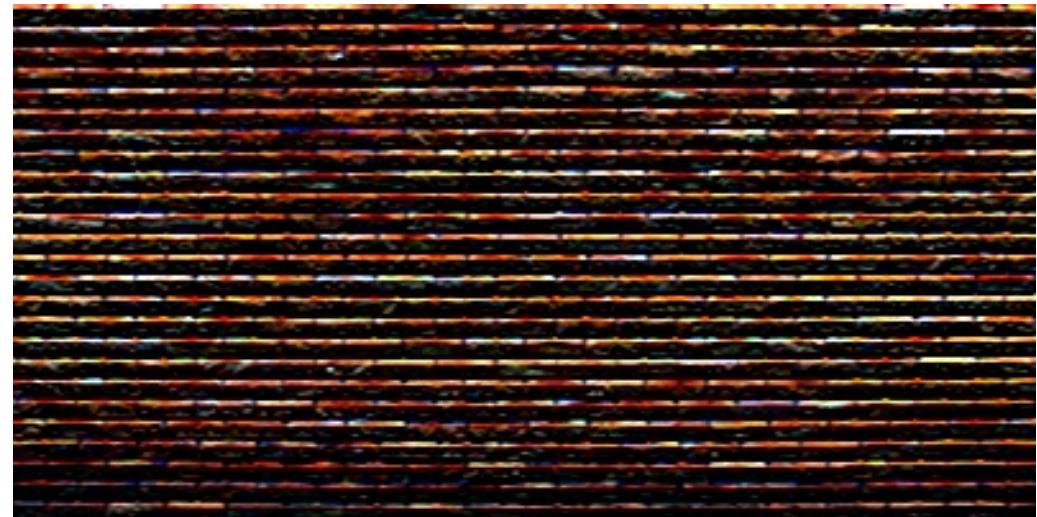
Sobel filter example



original



horizontal Sobel filter



vertical Sobel filter

Several derivative filters

Sobel

1	0	-1
2	0	-2
1	0	-1

1	2	1
0	0	0
-1	-2	-1

Scharr

3	0	-3
10	0	-10
3	0	-3

3	10	3
0	0	0
-3	-10	-3

Prewitt

1	0	-1
1	0	-1
1	0	-1

1	1	1
0	0	0
-1	-1	-1

Roberts

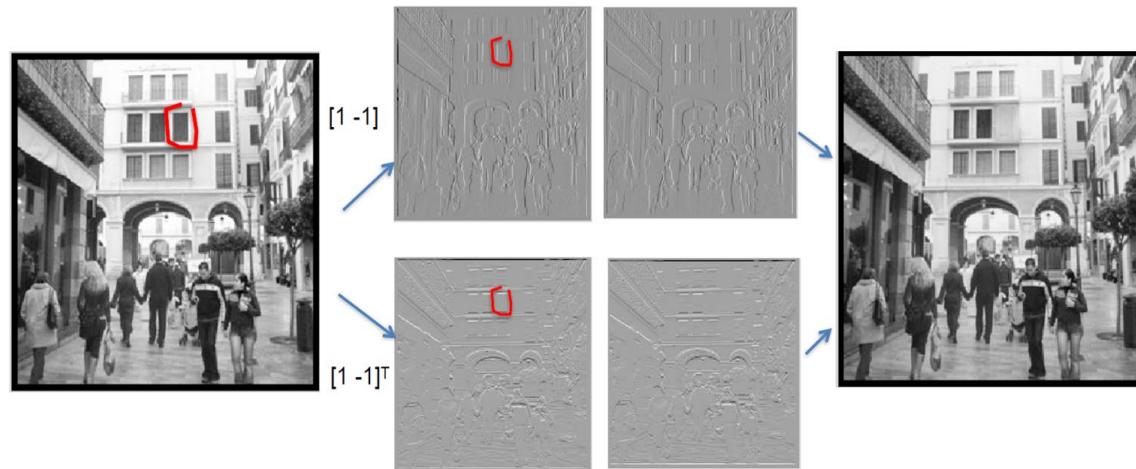
0	1
-1	0

1	0
0	-1

Reconstruction from 2D derivatives

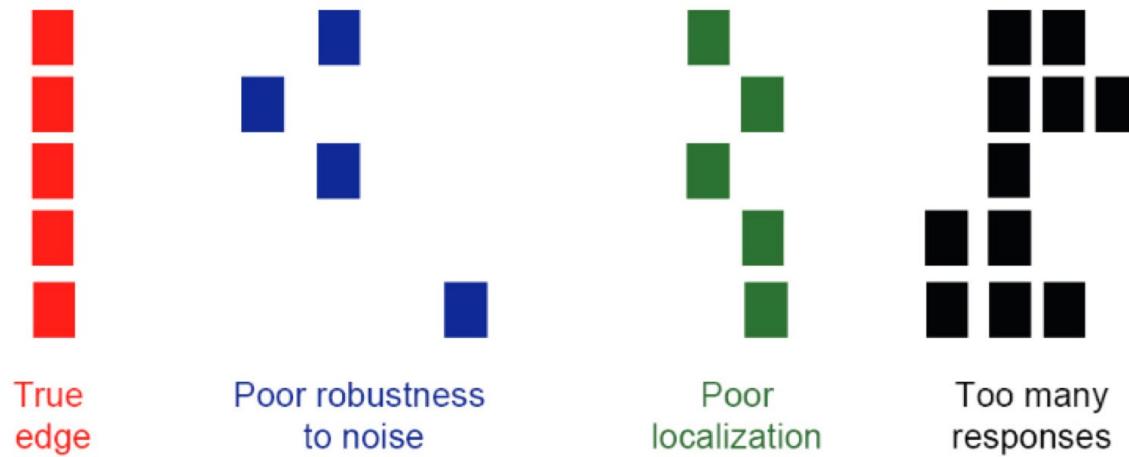
- In 2D, we have multiple derivatives (along n and m) and we compute the pseudo-inverse of the full matrix

Editing the edge image



Edge detector

- The optimal edge detector must be accurate, minimizing the number of false positives and false negatives; have precise localization, pinpointing edges at the positions where they actually occur; and have single response, ensuring that only one edge is found where there only is one edge



Example

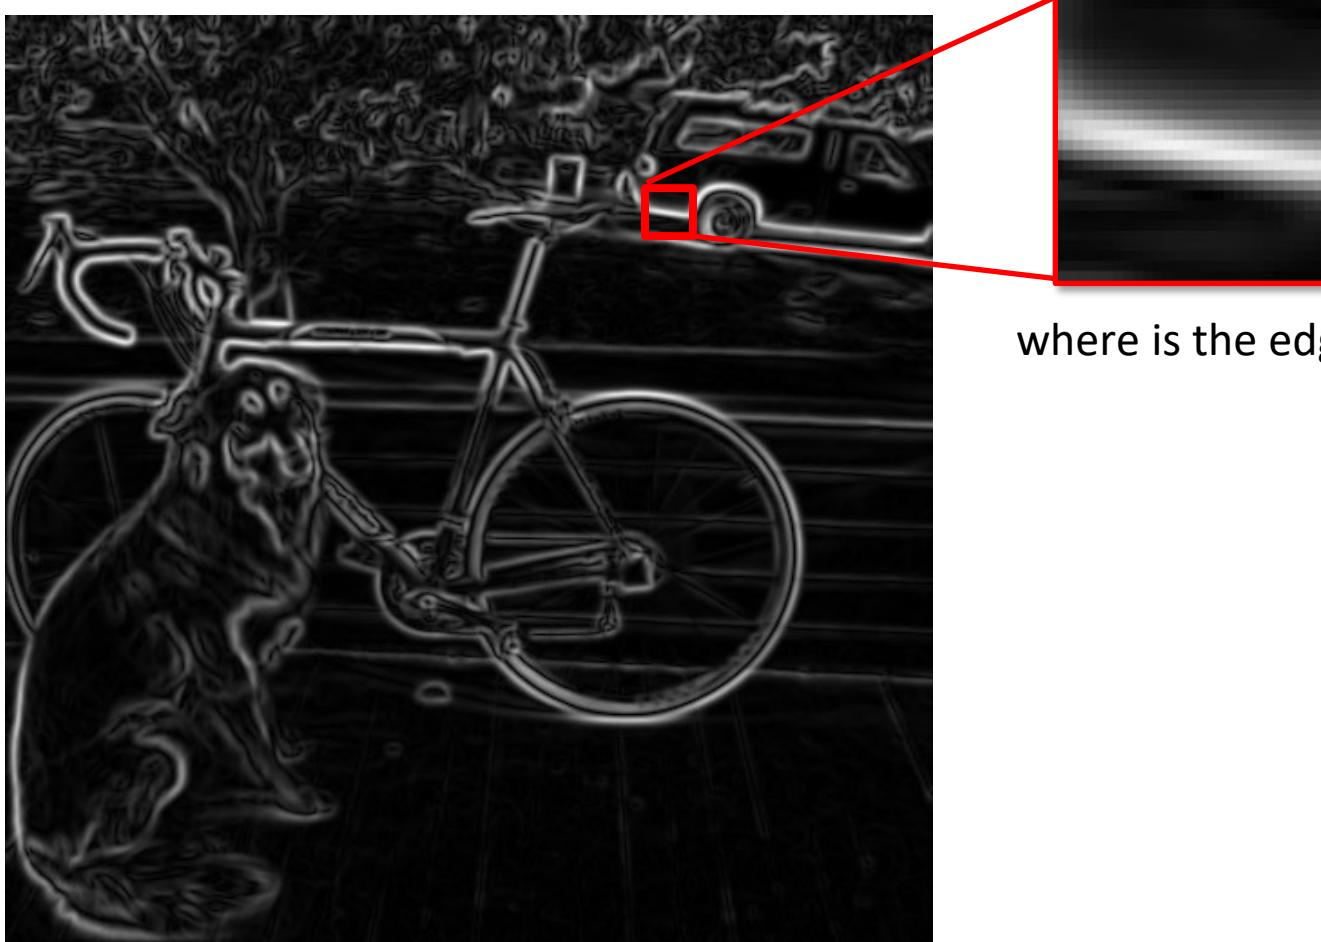


original image

Demo:

[http://bigwww.epfl.ch/demo/ip/
demos/edgeDetector/](http://bigwww.epfl.ch/demo/ip/demos/edgeDetector/)

Finding edges

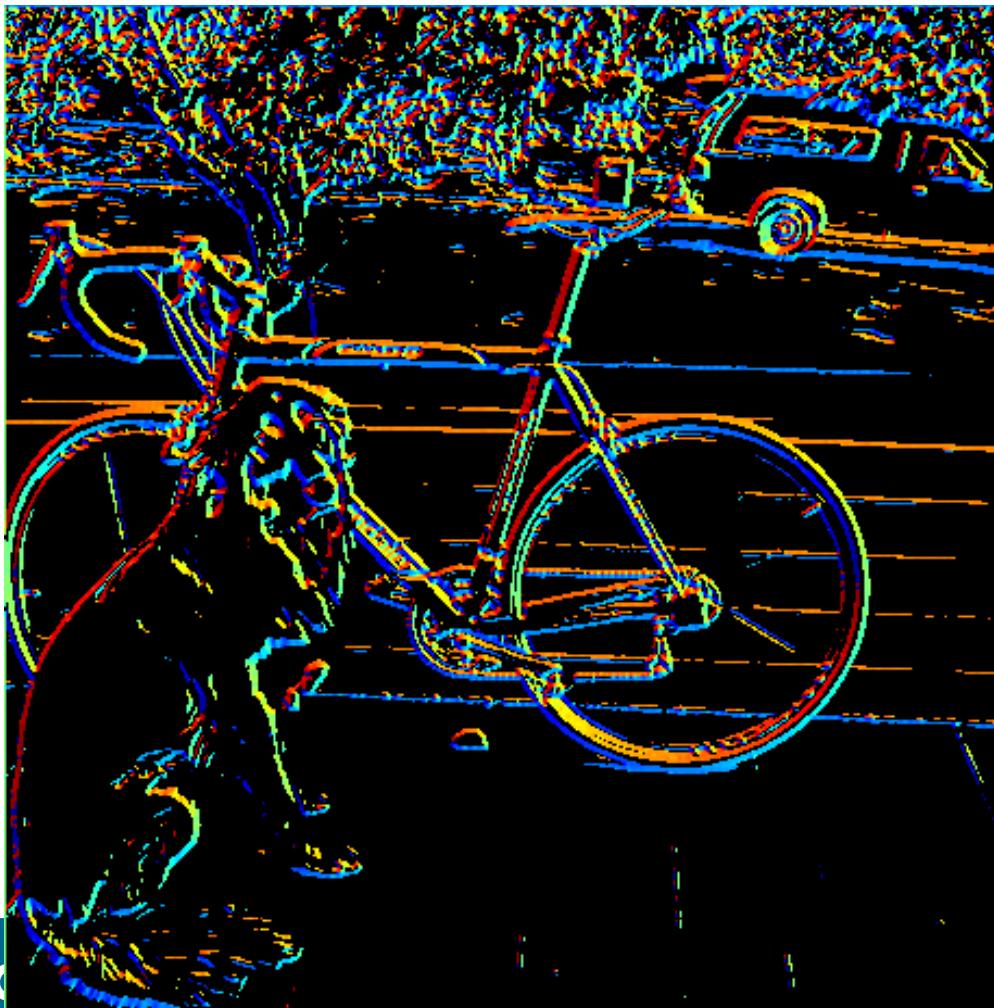


where is the edge?

smoothed gradient magnitude

Get Orientation at Each Pixel

- Get orientation (below, threshold at minimum gradient magnitude)



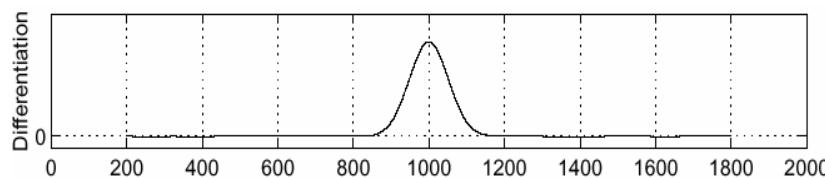
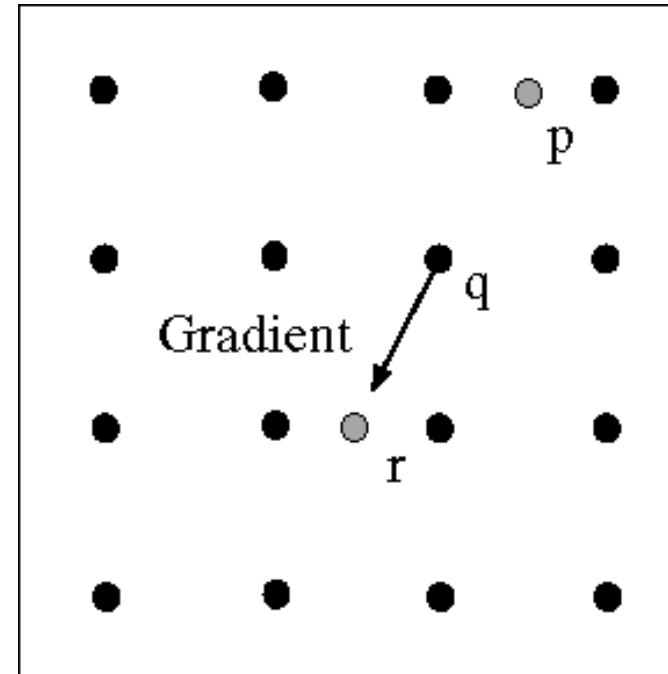
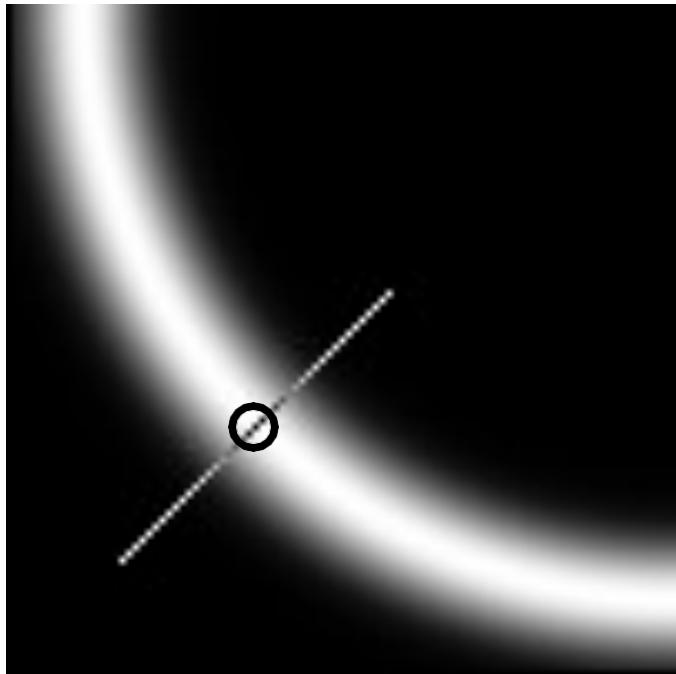
$\theta = \text{atan2}(gy, gx)$

360

Gradient orientation angle

0

Non-maximum suppression



- Check if pixel is local maximum along gradient direction
 - requires *interpolating* pixels p and r

Before Non-max Suppression



After Non-max Suppression



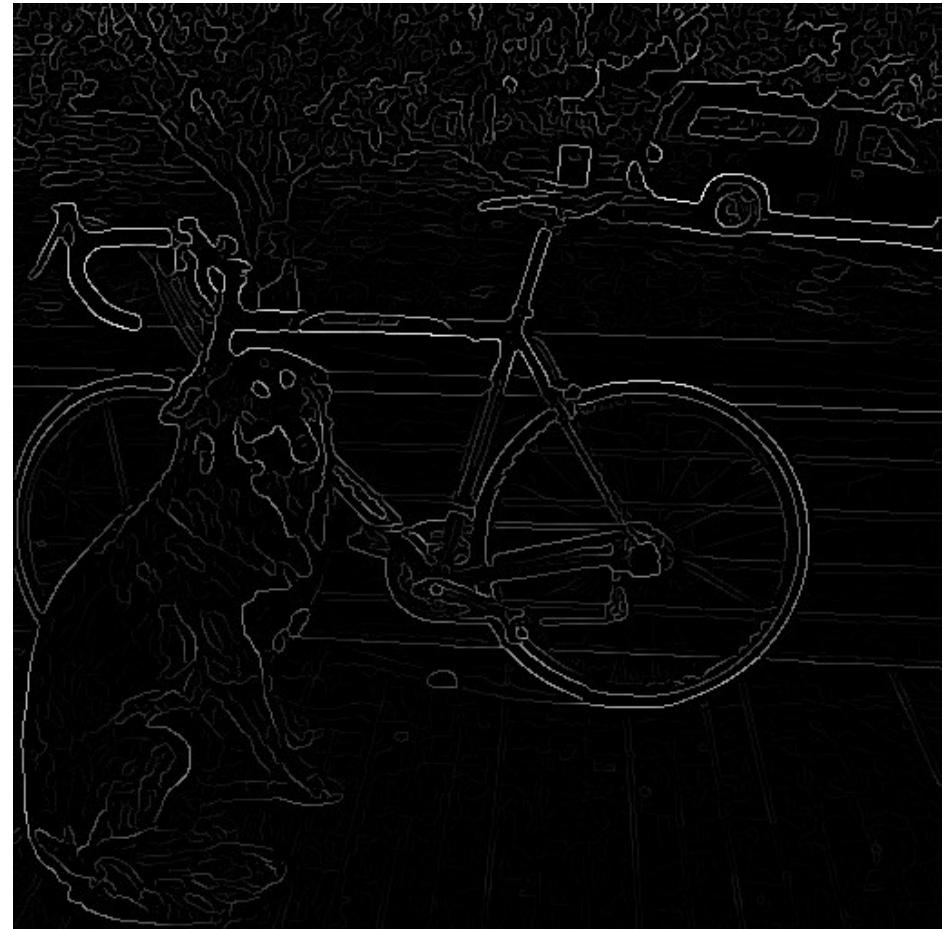
Still noise exists!



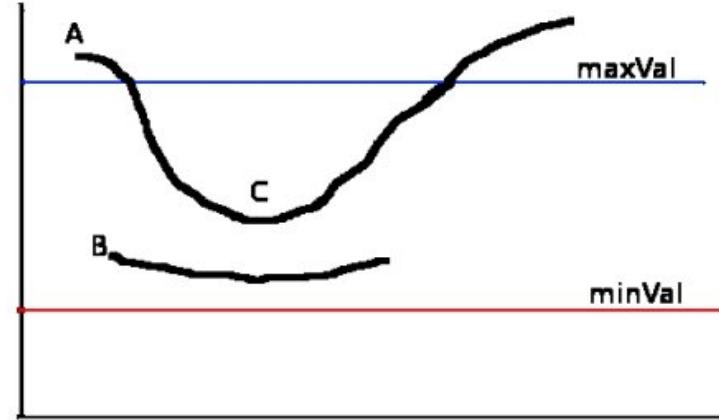
Thresholding edges

Still some noise but we want only strong edges

- 2 thresholds, 3 cases (R is the intensity gradient)
 - $R > T$: strong edge
 - $R < T$ but $R > t$: weak edge
 - $R < t$: no edge
- Strong edges are edges!
- Weak edges are edges iff they connect to strong
 - Look in some neighborhood (usually 8 closest)



Hysteresis thresholding

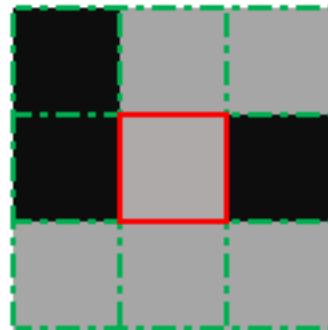


- The edge A is above the maxVal, so considered as "sure-edge".
- Although edge C is below maxVal, it is connected to edge A, so that also considered as valid edge and we get that full curve.
- Edge B, although it is above minVal and is in same region as that of edge C, it is not connected to any "sure-edge", so that is discarded.
- So it is very important that we have to select minVal and maxVal accordingly to get the correct result.

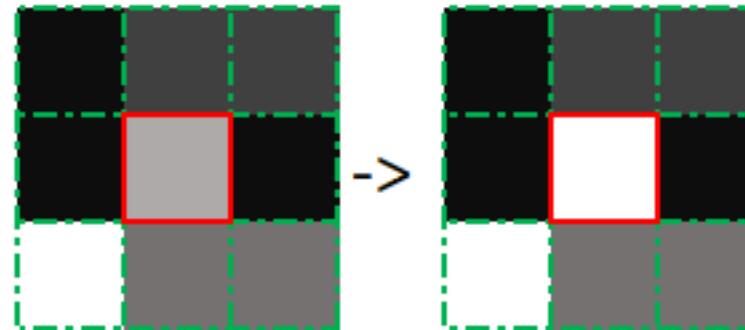
This stage also removes small pixels noises on the assumption that edges are long lines.

Edge Tracking by Hysteresis

Based on the threshold results, the hysteresis consists of transforming weak pixels into strong ones, if and only if at least one of the pixels around the one being processed is a strong one, as described below:



No strong pixels around



One strong pixel around



Canny edge detector



1. Filter image with Derivative of Gaussian

1. Find magnitude and orientation of gradient

1. Non-maximum suppression

1. Linking and thresholding (hysteresis):

- Define two thresholds: low and high
- Use the high threshold to start edge curves and the low threshold to continue them



Canny edge detector

- Our first computer vision pipeline!
- Still a widely used edge detector in computer vision
 - J. Canny, *A Computational Approach To Edge Detection*, IEEE Trans. Pattern Analysis and Machine Intelligence, 8:679-714, 1986.
- Depends on several parameters:
 - high threshold low threshold
 - width of the Gaussian blur σ

Canny edge detector



original

Canny

$\sigma = 1$

Canny

$\sigma = 2$

- The choice of σ ends with desired behavior
 - large σ detects “large-scale”
 - small σ edges detects fine edges

Slide Credits

- CS5670, Introduction to Computer Vision, Cornell Tech, by Noah Snavely.
- CS 194-26/294-26: Intro to Computer Vision and Computational Photography, UC Berkeley, by Alyosha Efros.
- CS 15-463, 663, 862, CMU, by Computational Photography, Ioannis Gkioulekas.

Acknowledgements: some slides and material from Bernt Schiele, Mario Fritz, Michael Black, Bill Freeman, Fei-Fei Li, Justin Johnson, Serena Yeung, R. Szeliski, Ioannis Gkioulekas, Roni Sengupta, Andreas Geiger